

# Smalltalk-80

A. Mével and  
T. Guéguen

Controll  
oller  
el  
roller  
View  
Views  
formation  
nport  
ndow  
r...formation  
isplayBox  
erWidth  
derColor  
deColor  
ndingBox  
elFrame  
belText  
Complemented  
edSubViews  
nimumSize  
aximumSize

ext

ardSystemView label:  
ardSystemView reverseLabel  
ardSystemView displayView  
ardSystemView label  
ardSystemView collapse

playView  
| clippingBox labelDisplayBox |  
clippingBox ← self clippingBox.  
labelDisplayBox ← self labelDisplayBox  
(labelDisplayBox intersect: clippingBox)  
isDisplay.  
labelText isNil  
iffFalse:  
[labelText  
displayOn: Display  
at: (labelDisplayBox center)  
(labelText boundingBox topLeft))  
clippingBox: clippingBox

**Macmillan Computer Science Series**

Consulting Editor

Professor F.H. Sumner, University of Manchester

S.T. Allworth and R.N. Zobel, *Introduction to Real-time Software Design*, second edition

Ian O. Angell and Gareth Griffith, *High-resolution Computer Graphics Using FORTRAN 77*

Ian O. Angell and Gareth Griffith, *High-resolution Computer Graphics Using Pascal*

M.A. Azmoodeh, *Abstract Data Types and Algorithms*

C. Bamford and P. Curran, *Data Structures, Files and Databases*

Philip Barker, *Author Languages for CAL*

A.N. Barrett and A.L. Mackay, *Spatial Structure and the Microcomputer*

R.E. Berry and B.A.E. Meekings, *A Book on C*

G.M. Birtwistle, *Discrete Event Modelling on Simula*

T.B. Boffey, *Graph Theory in Operations Research*

Richard Bornat, *Understanding and Writing Compilers*

Linda E.M. Brackenbury, *Design of VLSI Systems - A Practical Introduction*

J.K. Buckle, *Software Configuration Management*

W.D. Burnham and A.R. Hall, *Prolog Programming and Applications*

J.C. Cluley, *Interfacing to Microprocessors*

J.C. Cluley, *Introduction to Low Level Programming for Microprocessors*

Robert Cole, *Computer Communications*, second edition

Derek Coleman, *A Structured Programming Approach to Data*

Andrew J.T. Colin, *Fundamentals of Computer Science*

Andrew J.T. Colin, *Programming and Problem-solving in Algol 68*

S.M. Deen, *Fundamentals of Data Base Systems*

S.M. Deen, *Principles and Practice of Database Systems*

Tim Denvir, *Introduction to Discrete Mathematics for Software Engineering*

P.M. Dew and K.R. James, *Introduction to Numerical Computation in Pascal*

M.R.M. Dunsmuir and G.J. Davies, *Programming the UNIX System*

K.C.E. Gee, *Introduction to Local Area Computer Networks*

J.B. Gosling, *Design of Arithmetic Units for Digital Computers*

Roger Hutt, *Z80 Assembly Language Programming for Students*

Roland N. Ibbett, *The Architecture of High Performance Computers*

Patrick Jaulent, *The 68000 - Hardware and Software*

J.M. King and J.P. Pardoe, *Program Design Using JSP - A Practical Introduction*

H. Kopetz, Software Reliability  
E.V. Krishnamurthy, Introductory Theory of Computer Science  
V.P. Lane, Security of Computer Based Information Systems  
Graham Lee, From Hardware to Software - an introduction to computers  
A.M. Lister, Fundamentals of Operating Systems, third edition  
G.P. McKeown and V.J. Rayward-Smith, Mathematics for Computing  
Brian Meek, Fortran, PL/1 and the Algols  
A. Mével and T. Guéguen, Smalltalk-80  
Barry Morrell and Peter Whittle, CP/M 80 Programmer's Guide  
Derrick Morris, System Programming Based on the PDP11  
Y. Nishinuma and R. Espesser, UNIX - First contact  
Pim Oets, MS-DOS and PC-DOS - A Practical Guide  
Christian Queinnec, LISP  
Gordon Reece, Microcomputer Modelling by Finite Differences  
W.P. Salman, D. Tisserand and B. Toulout, FORTH  
L.E. Scales, Introduction to Non-linear Optimization  
Peter S. Sell, Expert Systems - A Practical Introduction  
Colin J. Theaker and Graham R. Brookes, A Practical Course on  
Operating Systems  
J-M. Trio, 8086-8088 Architecture and Programming  
M.J. Usher, Information Theory for Information Technologists  
Colin Walls, Programming Dedicated Microprocessors  
B.S. Walker, Understanding Microprocessors  
Peter J.L. Wallis, Portable Programming  
I.R. Wilson and A.M. Addyman, A Practical Introduction to Pascal -  
with BS6192, second edition

#### Non-series

Roy Anderson, Management, Information Systems and Computers  
J.E. Bingham and G.W.P. Davies, A Handbook of Systems Analysis,  
second edition  
J.E. Bingham and G.W.P. Davies, Planning for Data Communications

# **Smalltalk-80**

**A. Mével and T. Guéguel**

Edited by  
**Mario Wolczko,**  
Department of Computer Science,  
University of Manchester



© Editions Eyrolles 1987

Authorised English Language edition of Smalltalk-80, by A. Mével and T. Guéguel, first published 1987 by Editions Eyrolles, 61 boulevard Saint-Germain, 75005 Paris

Translated by M.J. Stewart

© English Language edition, Macmillan Education Ltd, 1987

All rights reserved. No reproduction, copy or transmission of this publication may be made without written permission.

No paragraph of this publication may be reproduced, copied or transmitted save with written permission or in accordance with the provisions of the Copyright Act 1956 (as amended), or under the terms of any licence permitting limited copying issued by the Copyright Licensing Agency, London.

Any person who does any unauthorised act in relation to this publication may be liable to criminal prosecution and civil claims for damages.

First published 1987

Published by

MACMILLAN EDUCATION LTD  
Hounds Mills, Basingstoke, Hampshire RG21 2XS  
and London  
Companies and representatives throughout the world

British Library Cataloguing in Publication Data

Mével, A.

Smalltalk-80. --- (Macmillan computer science series).  
1. Smalltalk -80 (Computer system)  
I. Title II. Guéguel, T.  
005.13'3 QA76.8.S635

ISBN 978-0-333-44514-3  
DOI 10.1007/978-1-349-09653-4

ISBN 978-1-349-09653-4 (eBook)

# Contents

<b>Preface</b>	xi
<b>1 Introduction</b>	1
1.1 Presentation	1
1.2 Smalltalk and its growth	1
<b>2 Principles of the Language</b>	6
2.1 Objects and messages	6
2.2 Classes and instances	7
2.3 Methods	8
<b>3 Classes and Instances</b>	9
3.1 Declaration of variables	9
3.1.1 Instance variables	9
3.1.2 Class variables	9
3.1.3 Global variables	10
3.1.4 Pool variables	10
3.2 Classes and subclasses	10
3.2.1 General remarks	10
3.2.2 Definition of a class	11
3.3 Metaclasses	11
<b>4 Syntax</b>	13
4.1 Literals	13
4.1.1 Numbers	13
4.1.2 Characters	14
4.1.3 Character strings	14
4.1.4 Symbols	14
4.1.5 Arrays	14
4.2 Variables	15
4.3 Messages	15
4.3.1 Composition of a message	15
4.3.2 Value returned by a message	16
4.3.3 Message priority	17
4.3.4 Cascaded messages	18
4.4 Blocks	18
4.4.1 Description of blocks	18
4.4.2 Blocks with arguments	18
4.4.3 Use of blocks in control structures	19

<b>5 Methods</b>	21
5.1 Messages and methods	21
5.1.1 Methods search	21
5.1.2 Object returned by a method	22
5.1.3 Temporary variables	22
5.2 Pseudo-variables	23
5.2.1 Pseudo-variable self	23
5.2.2 Super pseudo-variable	24
5.3 Primitive methods	25
<b>6 Some Elements of the User Interface</b>	26
6.1 General description	26
6.2 Main menu	27
6.3 Workspaces	27
6.4 Browsers	28
<b>7 The Class Object</b>	31
7.1 Identification of an object	31
7.2 Copies of objects	32
7.2.1 Non-duplicated instance variables	32
7.2.2 Duplicated instance variables	32
7.2.3 Copy message	33
7.3 Comparisons of objects	33
7.3.1 Equivalence of two objects	33
7.3.2 Equality of two objects	33
7.3.3 Comparison with nil	34
7.4 Indexed objects	34
7.5 Representation of objects	34
7.6 Control of errors	35
7.7 Control of messages	35
7.8 Dependencies between objects	36
7.9 Primitive messages	38
<b>8 Description of Architecture of Classes</b>	40
8.1 Class Behavior	41
8.1.1 Dictionary of methods	41
8.1.2 Instances	43
8.1.3 Manipulation of the class hierarchy	45
8.2 The classes ClassDescription and Class	47
8.2.1 Access to the description of a class	48
8.2.2 Category of classes and messages	48
8.2.3 Copy of messages	49
8.2.4 Compilation of methods	49
8.2.5 Access to variable names	50
8.2.6 Saving a class on a file	50
8.3 The class Metaclass	50
8.4 Multiple inheritance	50
<b>9 The Magnitude Classes</b>	53
9.1 The class Magnitude	53
9.2 The class Date	54

## Contents

vii

9.2.1 Creation of instances	54
9.2.2 Information about the calendar	54
9.2.3 Arithmetic	55
9.2.4 Conversion	55
9.3 The class Time	56
9.3.1 Creation of instances	56
9.3.2 Information about the time	56
9.3.3 Conversions	56
9.3.4 Arithmetic	56
9.4 The class Character	57
9.4.1 Access to the instances	57
9.4.2 Testing the instances	57
<b>10 The Numeric Classes</b>	59
10.1 The class Number	59
10.1.1 Arithmetic	59
10.1.2 Mathematics	60
10.1.3 Tests	61
10.1.4 Truncation and rounding	62
10.1.5 Conversions	62
10.1.6 Nature of the object returned by arithmetic operations	63
10.2 The class Float	64
10.3 The class Fraction	64
10.4 The class Integer	64
10.4.1 Enumeration	64
10.4.2 Some additional arithmetic functions	65
10.4.3 Bit manipulation	65
10.4.4 Changing base	66
<b>11 The Class Collection</b>	67
11.1 Creation of instances	67
11.2 Manipulating the elements of a collection	68
11.2.1 Adding elements to a collection	68
11.2.2 Removing elements from a collection	68
11.3 Tests on a collection	69
11.4 Enumeration of a collection	69
11.5 Conversion of a collection	70
11.6 The subclasses of the class Collection	71
11.6.1 The class Bag	71
11.6.2 The class Set	71
11.6.3 The class Dictionary	71
11.6.4 The class IdentityDictionary	74
11.6.5 The class SequenceableCollection	74
The class OrderedCollection	77
The class SortedCollection	78
The class LinkedList	79
The class Interval	80
The class ArrayedCollection	81

<b>12 The Class Stream</b>	84
12.1 The class PositionableStream	85
12.1.1 The class ReadStream	87
12.1.2 The class WriteStream	87
12.1.3 The class ReadWriteStream	88
12.2 The class ExternalStream	88
12.2.1 The class FileStream	89
12.2.2 The class Directory	89
<b>13 Processes</b>	91
13.1 The class Process	91
13.2 The class ProcessorScheduler	92
13.3 The class Delay	93
13.4 Semaphores	94
<b>14 The Classes Point and Rectangle</b>	96
14.1 The class Point	96
14.1.1 Creating a point	96
14.1.2 Messages supported by the class Point	96
14.2 Coordinate system within Smalltalk environment	98
14.3 The class Rectangle	99
14.3.1 Creating a rectangle	100
14.3.2 Instance methods	100
<b>15 The Graphics Classes</b>	105
15.1 The class Bitmap	105
15.2 The class DisplayObject	105
15.2.1 The class DisplayText	106
15.2.2 The class DisplayMedium	107
15.2.3 The class Form	110
15.2.4 The class Cursor	112
15.2.5 The class DisplayScreen	112
15.2.6 The class Path	113
15.3 The class BitBlt	114
15.4 The class Pen	115
<b>16 Basic Elements of the Interface</b>	118
16.1 Mouse	118
16.2 Windows	118
16.3 Pop-up menus	119
16.4 Standard windows	119
16.5 The system menu	120
<b>17 Text Editing Windows</b>	122
17.1 Scrollbars	122
17.2 How to select text	124
17.3 Edit commands	127
17.4 Workspaces	127
17.5 System workspace	128

<b>18 Browsers</b>	129
18.1 Description	129
18.2 The categories of class	130
18.3 The classes	131
18.4 Protocols	134
18.5 Methods	135
18.6 The code window	136
18.6.1 Editing a class	137
18.6.2 Editing a method	137
<b>19 Other Windows</b>	139
19.1 Inspectors	139
19.1.1 Description	139
19.1.2 Other inspectors	140
19.2 Debuggers	142
19.2.1 List of methods	143
19.2.2 The code window	143
19.2.3 Inspector of the receiver	144
19.2.4 Inspector of local variables	144
19.3 Views of external files	144
19.4 Graphics editors	147
19.4.1 The Form Editor	147
19.4.2 The Bit Editor	148
<b>20 Management of User Events</b>	150
20.1 The class InputState	150
20.2 The class InputSensor	150
<b>21 Model-View-Controller System</b>	152
21.1 The class View	152
21.1.1 The class WindowingTransformation	153
21.1.2 Instance variables of the class View	154
21.1.3 Messages understood by instances of View	156
21.2 The class Controller	158
21.2.1 Instance variables of Controller	158
21.2.2 Messages understood by instances of Controller	159
21.3 The class ControlManager	160
21.4 The class MouseMenuController	160
21.5 The classes StandardSystemView and StandardSystemController	162
21.6 Text-scrolling windows	163
21.6.1 Text-editing windows	164
21.6.2 Lists	165
21.7 Example	167
<b>Index</b>	175

# Preface

With the advent of affordable personal computing power, the Smalltalk-80 system is rapidly gaining in popularity. The Smalltalk-80 system can provide the personal computer user with a glimpse into the future of the programming process over the next decade.

This book is an introduction to programming in the Smalltalk-80 system. It describes the major elements of the system, namely the programming language, class structure and user interface, and serves as a concise reference for the most important classes in the system. Also included are many examples of Smalltalk-80 usage, and the source to a complete application. The description is based on the version 2 Virtual Image from Xerox, probably the most popular version of Smalltalk in use today.

The Smalltalk-80 system is large, and one cannot hope to describe it all in a book of this size. For a more comprehensive source, the reader is referred to *Smalltalk-80: The language and its Implementation*, by A. Goldberg and D. Robson, Addison-Wesley, 1983, and *Smalltalk-80: The Interactive Programming Environment*, by A. Goldberg, Addison-Wesley, 1984.

# 1 Introduction

## 1.1 Presentation

Smalltalk is one of several things: a programming language, an operating system, a programming environment, and a design and programming methodology. Many people consider Smalltalk to be a tool for knowledge representation, thus making it a serious competitor to Lisp and Prolog. It has moved on from its position as internal support system at the Xerox Palo Alto Research Center, which it held for ten years, to being used for the industrial production of software. The application areas in which it excels are application prototyping, interactive graphics systems and simulation. Although there are still some problems with speed of execution (mainly stemming from the interpretive nature of the system), spectacular results have been achieved in cost savings with certain types of software, especially in applications containing a high graphics or interactive component. The principal reason for this is that the Smalltalk environment takes the maximum advantage of the reuse of software components.

## 1.2 Smalltalk and its growth

Smalltalk is available on an increasing number of modern machines. It was first developed at Palo Alto in the beginning of the 1970s at the instigation of Alan Kay, and has evolved through three important stages:

- Smalltalk-72 was inspired by the Flex machine and the languages Lisp and Plasma
- Smalltalk-76 formalised the class, object and inheritance concepts inspired by Simula-67
- Smalltalk-80 is the latest version of the language, which was released from the Palo Alto laboratories.

Highlighted below are some of the stages in the design and development of Smalltalk at the PARC laboratories leading up to its general publication.

October 1972 building of first Smalltalk interpreter  
December 1972 definition of Smalltalk-72 according to the following basic principles

- control with automatic de-allocation
- interaction by messages
- behaviour of objects defined by classes

**1974 Smalltalk-74**

- introduction of message dictionaries
- graphic control by BitBlt

**1976 Smalltalk-76**

- unification of classes and contexts as objects
- inheritance by class hierarchy

**1978 Smalltalk-78**

- portability achieved with NoteTaker
- improvement of performance
- experimental implementation on an 8-bit microcomputer (Z80 or 6502) TinyTalk

**Summer 1979** first publications on Smalltalk produced. Invitations to several companies to carry out experimental implementations and participate in the evaluation of the product: Apple, DEC, Tektronix, Hewlett Packard

**August 1981** publication of an issue of 'Byte' devoted to Smalltalk

**September 1983** first Smalltalk licences granted by Xerox

**Summer 1984** appearance of first implementations of Smalltalk on different machines (Sun, SM90, etc)

**Early 1985** Tektronix 4404 workstation made available

**Spring 1985** 'Methods' (a Smalltalk clone) available for the IBM PC (Digitalk Inc) and Smalltalk for the Apple Lisa (Macintosh XL)

**Autumn 1985** Smalltalk available on the Apple Macintosh 512K (Apple) and on the IBM PC/AT (SoftSmarts Inc). More powerful Tektronix 4405 and 4406 models announced.

The essential characteristic of Smalltalk is its economy of design. In fact, this is based on the four concepts of object, message, class and inheritance.

Other mechanisms, such as the block or the dictionary are seen by the programmer as constructs developed from these basic concepts. This effort towards unification can be pushed even further, but it currently has no equivalent in other languages and systems. What it does do is to facilitate rapid mastery of the programming tools available.

From the outset of his research work, Alan Kay established one of the principles that was later to be followed in the different versions of Smalltalk. That was the principle of reaction. At any time, an object is active and must be capable of presenting itself to the user in a tangible form. This property greatly facilitates the creation and development of models.

One important difference between the architecture of a traditional product and that of a product based on Smalltalk is the absence of any frontier between the application and the system, which is thus completely open (this system is called the virtual image in Xerox terminology). To take an example, in a traditional

approach, it is possible to find an application that implements a table control by hash coding while the compiler also uses the same algorithm. With Smalltalk, it would be a design error not to reuse something that is already present in the system.

Programming with Smalltalk is incremental by nature and therefore lends itself very well to the progressive development of models (by stepwise refinement). This has two important consequences:

- There is an intentional fusion of the roles of programmer and user, because the same metaphors are used. This means that the user of a program will be able to move progressively from simple utilisation to modification and extension of the program with the same type of interaction.

- Of course, it will become possible to mix intimately the different activities of programming and debugging.

It would be possible to expand on other characteristics of Smalltalk that improve programming efficiency. For example, one can cite the abandonment of the concept of mode (in the context of insertion mode, query mode, etc). One can also think of the principle of strong modularity which requires that no part of the system depends upon details of the internal functioning of other parts, this property being induced by the architecture of communicating objects. However, it is above all the importance of the graphics interface and its good integration with the system and the language that one associates with Smalltalk.

The graphics interface consists of a bit-mapped screen and a three-button mouse that permits zones of pixels to be designated on the screen. These devices allow

- working in an environment consisting of several windows
- mixing text and graphics
- maintaining a dialogue very easily, thanks to pop-up menus which disappear as soon as the button that calls them is released
- selecting points, areas and text, all with great ease.

It should be appreciated that the graphics screen and the mouse pointing device have nothing to do with 'add-ons' to an existing system. They are thoroughly integrated into the programming environment which exploits them to the full.

Many research projects are currently evaluating the value of Smalltalk-80 in industrial software production. It is still too soon to pass overall judgement on this product, but some characteristics can be underlined here and now.

The fashion for object oriented languages goes beyond the phenomenon of Smalltalk. A large number of tools that make explicit reference to the same concepts are available or currently appearing. Without striving to be exhaustive, the following can be mentioned: Formes (P. Cointe), Kool (P. Albert), Ulysses (Y. Autret), Mering (J. Ferber), Ceyx (J.M. Hullot), X-Lisp (D. Betz), R-Lisp (J.L. Roos), Ross (D. McArthur), Loops (Bobrow), Glisp (Novak), Act1 and Act2 (Liebermann and Theriault), Flavors (Cannon and Moon), LRO (Ch. Roche), C++ (B. Stroustrup), Objective-C (Cox and Love), Clascal (Apple), Methods (Digitalk Inc), ObjectivePascal (N. Wirth and Apple), Oblogis (P.Y. Gloess), Neon (Kryia System Inc), Eiffel (B. Meyer), etc. All of these, mentioned at random, have quite diverse fields of application, but they all, to varying degrees, share certain aspects of Smalltalk. We can therefore say that the language conceived of by Alan Kay has had an important influence on the evolution of software production tools. The question is to establish whether this influence will be direct (effectively using Smalltalk), or indirect (borrowing ideas). Simula-67, which Smalltalk draws on considerably, has had a great influence on current programming languages, but its influence has been mainly indirect.

- The key to the question that has just been put on the future of the industrial use of Smalltalk depends very much on the choice that educational institutions make on the use of this language for their courses. The current availability of Smalltalk would indicate important growth potential from 1986 onwards. The present barrier to the wider industrial spread of Smalltalk is not the cost of the necessary workstations, but the lack of staff trained to use it.

- The extreme flexibility of Smalltalk allows the programmer to redefine the majority of his system functions with no restrictions. This can pose problems in a shared programming environment. Solutions exist, but are not yet integrated. However, for use in prototyping, this is not a serious handicap and the economies achieved promise increasing use of this tool.

There are currently numerous development projects based on Smalltalk. The most serious test of a widespread use of this product will be the existence of an exchange supply of software components written in Smalltalk.

The aim of this book is not to expound the theory of object languages but rather to offer a practical guide to programming in Smalltalk. There can be no question of developing here all the aspects of Smalltalk, because in its most recent versions it contains more than 230

classes and 4500 methods. Nevertheless, we have tried not to neglect any of the important points of Smalltalk; in particular, we have devoted several chapters to the user interface and more precisely to the MVC (Model, View, Controller) interface system that forms the basis of every Smalltalk interface.

## 2 Principles of the Language

### 2.1 Objects and messages

All the entities used in Smalltalk-80 are objects. For example, an object can represent a number, a character, a drawing, a list, a program, an editor, etc.

An object has two main characteristics:

- it allows data belonging to it to be stored and accessed
- it can reply to a certain number of messages.

For example, the object 3.14 will respond to the message `sin` by returning the value 0, which is written:

```
3.14 sin
```

The only means of manipulating an object is to send it a message. Messages therefore constitute the interface between objects and the external world. The way in which an object responds to a message is private to it and does not have to be known externally.

Writing an application in Smalltalk consists of determining the objects necessary for the description of the problem and in defining the possible operations for each object.

For example, let us take a computer-aided drawing application. This must allow us to draw easily all simple geometric shapes. We will restrict ourselves to points, straight lines and circles. We can see already that just stating the problem indicates some of the objects that will have to be created; this produces a first breakdown of the problem into objects:

- drawing
- point
- line
- circle

Drawings will be defined by a list of points, a list of straight lines and a list of circles.

Points will be defined by two coordinates x and y.

Lines will be defined by two points.

Circles will be defined by a point and a radius.

The following operations are defined on drawings:

- display (displays the drawing on the screen)
- points (returns the list of points in the drawing)
- lines (returns the list of lines in the drawing)

- circles (returns the list of circles in the drawing)
- addPoint: aPoint (adds a point to the drawing)
- addLine: aLine (adds a line to the drawing)
- addCircle: aCircle (adds a circle to the drawing)

The following operations are defined on points:

- display (displays the point on the drawing)
- xy (returns the coordinates of the point)
- x: aNumber (initialises the x-coordinate of the point)
- y: aNumber (initialises the y-coordinate of the point)

The following operations are defined on lines:

- display (displays the line on the drawing)
- origin (returns the startPoint of the line)
- corner (returns the the endPoint of the line)
- origin: aPoint (initialises the startPoint of the line)
- corner: aPoint (initialises the endPoint of the line)

The following operations are defined on circles:

- display (displays the circle on the drawing)
- centre (returns the centre point of the circle)
- radius (returns the radius of the circle)
- centre: aPoint (initialises the centre of the circle)
- radius: aNumber (initialises the radius of the circle).

This example demonstrates the modular aspect of Smalltalk. When the basic objects of an application and the operations on these objects are defined, the programmer can use them without knowing the detail of their construction; he can also reuse them in contexts that are quite different from those for which they were constructed.

## 2.2 Classes and instances

In the previous example, a drawing is made up of several points, several lines and several circles. All the points must interpret the messages sent to them in the same way.

This leads us to group together objects of the same kind into classes (that is, those objects that represent

the same sort of entity and interpret messages in the same way).

Thus, the class of points whose elements will be particular points is defined.

Objects contained in a class are called instances of this class. The class describes the form in which information about its instances are stored; it also describes the way in which its instances will reply to the messages that they receive. In a class the instances are represented by a certain number of variables called instance variables. For the class of points the instance variables are the coordinates x and y; for the class of lines the instance variables are the origin and the corner.

### **2.3 Methods**

The way in which an object responds to a message is described in a method. All instances of a class use the same method to respond to a message. Each class contains a list of methods that allows its instances to respond to the messages that are sent to them.

# 3 Classes and Instances

## 3.1 Declaration of variables

In the Smalltalk environment we encounter several types of variable that differ in their scope. These different types are:

- instance variables
- class variables
- global variables
- pool variables

### 3.1.1 Instance variables

Instance variables are variables that belong to an object and which allow it to be differentiated from other instances of its class. We can distinguish two types of instance variable, those that are indexed and those that are named.

#### Named instance variables

These are instance variables that are identified by a name. For the instances of the class **Point** there are two named instance variables, **x** and **y**.

When the name of an instance variable appears in an expression, this name refers to the value of the corresponding instance.

When a new instance is created, it contains instance variables specified by its class; the default value of these instance variables is the undefined object called **nil**.

#### Indexed instance variables

These are instance variables that are not accessed by name but by a numeric index. In contrast to objects that possess only named instance variables, those that possess indexed instance variables may have a varying number of instance variables.

In Smalltalk-80 there is for example a class with indexed instance variables which allows arrays to be represented; for each instance the number of variables corresponds to the size of the table.

### 3.1.2 Class variables

Class variables are variables that are shared by all the

instances of a class. For example, in the case of a computer-assisted drawing application, the outline of a circle is made up of several line segments; the number of these segments is the same for all circles, and one can therefore define it as a class variable.

### **3.1.3 Global variables**

Global variables are variables that are shared by all objects. They are stored in a dictionary called Smalltalk. For example, all the classes are referenced by global variables whose name is the access key to the dictionary called Smalltalk.


### **3.1.4 Pool variables**

Pool variables are variables that are accessible by the instances of several classes. In order to define a set of variables shared by several classes, it is necessary to define a dictionary that is common to these classes, this dictionary being itself a global variable.

## **3.2 Classes and subclasses**

### **3.2.1 General remarks**

Each object belongs to one and only one class, as can be represented thus:



It can, however, be useful to share some elements of the description between several classes or to describe one class by means of another.(See also section 5.2.2.)

This is what happens in Smalltalk, where each class is described in terms of another class called its superclass. The instances of the new class are identical to those of its superclass, except for the additions

made explicit in the new class. The new class thus defined is a subclass of its superclass.

One subclass is itself a class and can therefore also have one or more subclasses. The set of classes thus takes on a tree structure whose root is the class called **Object**. **Object** is the only class that does not possess a superclass.

The instances of each subclass inherit from all the instance variables and from all the class variables, as well as from all the methods of the superclass.

### 3.2.2 Definition of a class

To give a complete definition of a class, we need the following elements:

- its name
- its superclass
- its new class variables
- its new pool variables
- its new instance variables
- the list of its new methods.

The name of the class is obligatory and must be different from the name of any other existing class.

The superclass is obligatory and must correspond to a class already in existence.

Variables are optional, but if they exist their name must be different from any name already defined in the set of its superclasses.

New methods may be added, or methods already defined in the set of its superclasses may be redefined.

### 3.3 Metaclasses

As we have said earlier, each Smalltalk entity is an object. Consequently, a class is itself an object. We have also seen how each object belongs to one and only one class. A class is therefore itself an instance of another class that we call its metaclass. In earlier versions of Smalltalk all the classes were instances of one single metaclass. For reasons of flexibility (especially when creating new instances of a class), each class is the sole instance of a metaclass.

The metaclasses are classes and therefore contain methods that allow their instances (that is, the classes) to respond to the messages that they receive. As a result, these methods will be called class methods, while the other methods will be called instance methods. When a class is created, a new metaclass is automatically created. Among the messages sent to classes we find for example the message `new`, which allows a new instance of the class that received this message to be created.

In contrast to other classes, metaclasses have no name; neither are they metaclass instances. They are all instances of the same class called **Metaclass**.

One can access the metaclass of a class by sending it the message `class`. So, if we call `Point` the class of points, we can access its metaclass with the expression:

`Point class`

The metaclasses are classes and therefore have a superclass called **Class**.

To summarise, we can regard metaclasses in two ways:

- as instances of the class **Metaclass**,
- as a subclass of the the class **Class**.

**Metaclass** and **Class** are two subclasses of the class **ClassDescription**.

# 4 Syntax

An expression is a sequence of characters which, when evaluated, always returns an object. In Smalltalk-80 there are four lexical types, as follows:

- literals that represent certain predefined objects like numbers, characters, character strings;
- variable names that represent variables that can be accessed in the context of the expression;
- messages;
- blocks that represent an expression whose evaluation is postponed.

## 4.1 Literals

### 4.1.1 Numbers

Numbers are objects that represent numerical values and respond to mathematical messages.

A number may be positive or negative; if it is negative it will be preceded by the minus sign -.

A number may be a decimal or an integer; if it is decimal it includes the decimal point in its sequence of figures, for example:

```
4  
4.56  
-7  
-7.234  
432
```

When the base used is not base 10, the number begins with its base expressed in base 10, followed by the character 'r' and ending with the expression of the number in that base. When the base is greater than 10, the first letters of the alphabet are used to represent numbers greater than 9, for example:

Smalltalk expression	Base 10 equivalent
2r10	2
2r1010	10
2rl.1	1.5
8rl1	9
8rl23	83
16rA1	161
16r-A1	-161

A number may be expressed in scientific notation; in this case the list of figures is followed by the character 'e' and the exponent is expressed in base 10. The number in front of the exponent is multiplied by the base to the power of the exponent, for example:

Scientific notation	Base 10
2.549e3	2459
-2.459e-3	-0.002459
2rl0e3	16
2rl0e-3	0.25

#### 4.1.2 Characters

Characters represent the elements of the alphabet. They are indicated by the '\$' followed by a character, for example:

```
$a
$A
$$
$9
```

#### 4.1.3 Character strings

A string of characters represents a sequence of characters. This is indicated by placing the list of characters between quote marks. If the string itself contains an apostrophe, this must be doubled in order to avoid errors in interpretation, for example:

```
'this is a string'
'it''s a string too'
```

#### 4.1.4 Symbols

Symbols are strings of characters that are used to represent the names of objects in the system. They are written by prefixing them with the hash '#', for example:

```
#name1
#ThisIsASymbol
```

#### 4.1.5 Arrays

An array is a structured object whose elements may be accessed by their numeric index. Each element of the array is a literal. It is indicated by the hash and the list of its elements is enclosed by brackets. If an array or a symbol is included among the elements of the array, there is no need for a hash in front of these elements, for example:

```
 #(1)
 #(2 7 9)
 #(l 'one' one (l one 'one is a string'))
```

#### 4.2 Variables

A variable allows an object to be referenced; all variables except indexed variables have a name. A variable name may be used in an expression to designate an object. From a given object one may access objects referenced by variables whose name is known by this object.

Because the range of instance variables of an object is limited to the object itself, they will be inaccessible from any other object, and can therefore only be used by methods defined for this object.

A variable name is a sequence of alphanumeric characters and must begin with a letter, for example:

```
v1  
nameOfVariable  
NameOfVariable
```

By convention, variables whose range is limited to the object (instance variables and temporary variables) begin with a lower case letter, and variables that are accessible by several objects (global variables, pool variables, class variables) begin with an upper case letter.

To assign a value to a variables, the special character '`<-`' is used, preceded by the variable name followed by an expression whose evaluation will return an object that will then be referenced by the variable, for example:

```
variable1 <- 1  
variable2 <- 'variable 2 references a string of  
characters'  
variable3 <- Point class  
variable4 <- variable5 <- 123
```

#### 4.3 Messages

Messages allow interactions between objects, for example:

```
2 sqrt  
returns the square root of 2,  
2 + 8  
returns the sum of 2 and 8,  
variable1 + 4  
returns the sum of 4 and the number referenced by vari-  
able 1,  
variable1 > variable4  
compares the objects referenced by the two variables,  
10 raisedTo: 2  
returns 10 raised to the power of 2.
```

##### 4.3.1 Composition of a message

An expression requires a receiver to appear, followed by

a selector and possibly arguments. The receiver and the arguments are described by expressions, while the selector is described by a literal, for example:

```
2 sqrt
the receiver is 2, the selector is sqrt,
2 + 8
the receiver is 2, the selector is +, the argument is 8,
variable1 + 4
the receiver is the object referenced by variable1, the selector is +, the argument is 4,
variable1 > variable4
the receiver is the object referenced by variable1, the selector is >, the argument is the object referenced by variable4,
10 raisedTo: 2
the receiver is 10, the selector is raisedTo:, the argument is 2.
```

There are three types of messages, characterised by the number of arguments:

#### **Unary messages**

These are messages that have no argument, for example:

```
2 sqrt
```

#### **Keyword messages**

This is the most common form of message. The selector is made up of one or more keywords preceding each argument. Each keyword terminates with ':', for example:

```
10 raisedTo: 2
the selector is raisedTo:, the unique keyword is raisedTo:, the argument is 2,
list at:2 put: 1
the selector is at:put:. the keywords are at: and put:, the arguments are 2 and 1.
```

#### **Binary messages**

In the case of a single argument, it is possible not to use a keyword for the selector but to replace it with a selector made up of one or two non-alphanumeric characters; in this case the message is said to be binary, for example:

```
2 + 8
variable1 >= variable4
```

#### **4.3.2 Value returned by a message**

In response to the message the receiver returns an object which becomes the value of the expression. In the following example

```
2 + 3 * 4
```

the first message sent to object 2 has as its selector `+`, and as its argument 3; the object returned in response to this message is object 5. The expression then becomes

```
5 * 4
```

and the result returned by the first expression is the object 20.

In the case where the expression includes an assignment, the object returned is the object referenced by the assignment variable. For example, the expression

```
variable1 <- 2 + 3 * 4
```

returns the object 20.

#### 4.3.3 Message priority

Messages are taken in order from the left of the expression, with the following priorities (in decreasing order of precedence):

- unary message
- binary message
- keywords message

In the example

```
10 raisedTo: 2 + 3 sqrt
```

the highest priority message is the unary message `sqrt`; the value returned will be 1.73205. The second message will be the binary message `+`, the returned value being 3.73205. The last message will be the keyword message `raisedTo:`, with the value returned being 10 raised to the power of 3.73205, that is 5395.74.

In order to alter the order in which messages are sent, parentheses may be used. The highest priority message is the one enclosed between parentheses. For example, in the expression

```
2 + (3 * 4)
```

the first message sent will be `*`, which will go to the object 3 with argument 4; the result will be 12, which will become the argument for the selector `+` of the second message. The final object returned will thus be 14.

When several keywords appear in a message, they are grouped under a single selector. For example, in the expression

```
10 raisedTo: 2 raisedTo: 3
```

the selector will be `raisedTo: raisedTo:`, which is not a correct selector. Parentheses must therefore be used, as follows:

```
(10 raisedTo: 2) raisedTo: 3
```

or

```
10 raisedTo: (2 raisedTo: 3)
```

#### 4.3.4 Cascaded messages

It is possible to send several messages to the same object by separating these messages with semi-colons. In the following example

```
list at: 1 put: 'one'  
list at: 2 put: 'two'  
list at: 3 put: 'three'
```

could be written

```
list at: 1 put: 'one' ; at: 2 put: 'two' ; at: 3 put:  
'three'
```

The use of cascaded messages is not obligatory but it does make the writing of expressions more concise.

### 4.4 Blocks

#### 4.4.1 Description of blocks

Blocks are Smalltalk objects used in control structures. A block is a sequence of expressions separated by full stops; it begins with the character '[' and ends with the character ']'. Evaluation of a block must be requested explicitly; this allows it to be used in control structures. An example of a block is:

```
[variable1 <- 1. variable2 <- 2]
```

If there is a point between the last expression and the final bracket, it will be ignored.

A block is an object that can execute expressions placed within brackets if it receives the message value. For example

```
[variable1 <- 1. variable2 <- 2] value
```

and

```
variable1 <- 1. variable2 <- 2
```

will produce the same result.

Evaluation of a block returns the value returned by the last expression in the block. When a block contains no expression the value returned is nil.

#### 4.4.2 Blocks with arguments

It is possible to pass one or more arguments to a block before its execution. In a block the arguments are specified at the beginning of the block by identifiers beginning with the colon ::. Arguments are separated from expressions by the bar character |.

For example, the block

```
[:i | variable1 <- i]
```

is a block with an argument that is i. The block

```
[:i :j | variable1 <- i. variable2 <- j]
```

is a block with two arguments, i and j.

The arguments of a block are variables local to that block.

To evaluate a block with an argument, you must send it the message value: with an argument that will be passed to the block. In the event of two arguments, the message to be used is value:value::.

Use value:value: value: and value:value:value:value: for three or four arguments; use valueWithArgs: anArray for more.

#### **4.4.3 Use of blocks in control structures**

When we refer to classic programming languages like Pascal or C, we find that they contain three types of control structure: the loop (for, do), the test (if then else) and the conditional loop (while, until). These structures are also found in Smalltalk-80 and make use of blocks. Smalltalk-80 does not have an equivalent to the multi-branch (case) statement.

#### **Loops**

Starting with a list of elements it is possible to execute a block for each element in the list, with each element becoming the argument of the block. This is done by sending the message do: to a list of elements. For example

```
listIntegers do: [:eachInteger | sum <- sum +
    eachInteger]
```

will have the effect of calculating the sum of all the integers contained in the list.

#### **Tests**

Boolean objects are returned from messages like >, <, >=, <= etc. These Boolean objects, true and false, respond to the messages ifTrue:, ifFalse:, ifTrue:ifFalse: and ifFalse:ifTrue:. The arguments of these messages are blocks that will be evaluated or not, depending on whether the receiver is true or false. In the example

```
anInteger < 0 ifTrue: [anInteger <- 0 - anInteger]  
the sign of anInteger is changed if it is negative.
```

In the example

```
anInteger < 0 ifTrue: [anInteger <- 0] ifFalse:  
[anInteger <- anInteger sqrt]
```

anInteger is set to 0 if it is negative, otherwise its square root is assigned to it. To clarify the expression, it can be formatted differently and may for example be written as follows:

```
anInteger < 0 ifTrue: [anInteger <- 0]  
ifFalse: [anInteger <- anInteger sqrt]
```

#### **Conditional loops**

A conditional loop operates for as long as a condition is true or holds. This condition is evaluated at each

iteration of the loop. That can for example be achieved by sending the message `whileTrue:` or the message `whileFalse:` to a block with another block as argument. The example below allows the factorial of a number to be calculated iteratively

```
fact <- 1.  
[aNumber > 1] whileTrue:  
[fact <- fact * aNumber.aNumber <- aNumber - 1]
```

So long as `aNumber` is greater than 1, the value returned by the evaluation of the first block is true and therefore the second block is evaluated. As soon as `aNumber` is equal to 1, the evaluation of the first block returns false and the iteration stops. The result will be referenced by the variable `fact`. This example can be handled with the message `whileFalse:`

```
fact <- 1.  
[aNumber <= 1] whileFalse:  
[fact <- fact * aNumber.aNumber <- aNumber - 1]
```

Note that these control structures are constructed from the basic language features, namely, blocks methods and messages, and require no extra features in the language.

For example, the `ifTrue: ifFalse:` control structure is built from two methods, one in class `True` (whose sole instance is `true`),

```
ifTrue: aBlock ifFalse: aSecondBlock  
    ↑aBlock value
```

and another in class `False` (whose sole instance is `false`),

```
ifTrue: aBlock ifFalse: aSecondBlock  
    ↑aSecondBlock value
```

# 5 Methods

## 5.1 Messages and methods

A method describes how an object is going to reply to a message. The name of a method is the same as the name of the message to which the method must allow a response to be made. A method consists of its name followed by a list of expressions separated by full stops. For example, the method of the class Point which returns the x-coordinate of that point will be written

```
x                         message name  
↑x                         return of instance variable x
```

The method of class Point which alters the x-coordinate of the point will be written

```
x: aValue                  message name  
x <- aValue                assigns to x the value aValue
```

The method of class Point which displays the point on the drawing will be written

```
display                   message name  
|square|                  definition of a variable  
                           local to method  
square <- Form new.      square becomes a new  
                           instance of Form  
square extent: 2@2        size of square is fixed as 2  
                           x 2 pixels  
square black.            square is black  
square displayAt: x@y    display square at position  
(x,y)
```

**Form** is a predefined Smalltalk class.

**new** is a method of class which creates a new instance.

**extent:**, **black** and **displayAt:** are three instance methods of the class **Form**.

The arguments found in the method name are variables local to the method and cannot be altered. For example

```
x:aValue  
aValue <- aValue + 1.  
x <- aValue
```

produces a syntax error.

### 5.1.1 Methods search

When an object receives a message, it looks in its class

for a method having the same selector as the message received. If it does not find it, it continues its search in the superclass of its class; this search is continued through the sequence of its superclasses until the correct method is found.

If on reaching the root class, Object, the method has not been found, there is an error. The receiver is then sent the message `doesNotUnderstand:`, whose argument is the message whose selector has not been found.

The method which replies to the message `doesNotUnderstand:` is set up in the class Object and is therefore known to all objects. It indicates to the programmer that there is an error.

If the search is successful, each expression of the method found is evaluated. When all expressions have been evaluated, the receiver returns to the sender an object which is the response to the message received.

### 5.1.2 Object returned by a method

When a method terminates the receiver returns an object to the sender. The default object is the receiver itself. When it is required to return another object, this is indicated in the method by prefixing the object to be returned with the character '`^`'.

As soon as this character is found, the object produced by the expression that follows is returned to the sender and the method terminates. For example, the following method of class Point

```
xPositive
  x > 0 ifTrue: [↑true]
    ifFalse: [↑false]
```

will return true if the x-coordinate of the point is positive, and false otherwise. This method may be written more simply as

```
xPositive
  ↑x > 0
```

### 5.1.3 Temporary variables

As we saw in the example dealing with the display method in class Point, a variable, `square`, was introduced. Temporary variables like this may be declared in a method, but they are accessible only during execution of the method.

To declare temporary variables their name must be placed between two '`|`' characters.

Declaration of these variables must occur between the method name and the list of expressions. Every time that a method is executed these variables are set to nil. For

example, the method

### **invertXandY**

```
|temporaryVariable|
temporaryVariable <- x.
x <- y
y <- temporaryVariable
```

allows the coordinates of a point to be inverted using the temporary variable called temporaryVariable.

## **5.2 Pseudo-variables**

A pseudo-variable is a variable whose value cannot be altered by assignment.

Some pseudo-variables are constants; that is, they always refer to the same object. So far we have met three, namely

- nil the default value of a variable
- true the object returned by a true condition
- false the object returned by an false condition.

Other pseudo-variables depend on the context in which they are used. Such is the case with the pseudo-variables, self and super.

### **5.2.1 Pseudo-variable self**

The pseudo-variable **self** always refers to the receiver of the message itself. Every method has access to the self pseudo-variable. For example, the method

#### **abs**

```
self < 0 ifTrue: [↑0 - self]
```


returns the absolute value of the receiver.

self also allows recursion to be handled, as in the following example:

#### **factorial**

```
self = 0 ifTrue: [↑1].
↑self*(self - 1) factorial
```

When the expression 3 factorial is evaluated, self will represent 3 at the first pass. The method will request the response to the message (self - 1) factorial, namely 2 factorial, and so on until 0 factorial. This may be shown diagrammatically as follows:



When a message is sent to self, the search for the method corresponding to the message will begin in the

class of the object referenced by self. This class may well be a subclass of the class containing the first method. This is a consequence of the strategy used to look for a method.


### 5.2.2 Super pseudo-variable

There is a second pseudo-variable, called super, that is also accessible by all methods.

This pseudo-variable again refers to the receiver, but its strategy for searching for a method is different.

When a message is sent to super, the method search begins not in the receiver's class, but in the superclass of the class that the message was sent from.

For example, say we have a hierarchy of three classes, as shown in the following figure:



in which a b c are three respective instances of classes A B C. Say that in class A the method is

**name**  
↑ 'A'

in class B the method is

**name**  
↑ 'B'

and in class C the method is

**name**  
↑ 'C'

Say that in class B the methods are

**myName**  
↑ self name

and

**superName**  
↑ super name

The different results of the messages myName and superName sent to objects b and c will be

Expression	Result
c myName	'C'
c superName	'A'
b myName	'B'
b superName	'A'

One of the most frequent cases of the use of the super pseudo-variable is when redefining a method in a subclass. For example:

```
initialise
    super initialise.
    self init
```

This method calls the method with the same name that is located in a superclass of the class in which it is embedded.

### 5.3 Primitive methods

We have just seen that in order to respond to a message an object executes a method; execution of this method results in the evaluation of one or more expressions and therefore also to the sending of messages to other objects. To avoid an infinite regress of message sending, there must be methods that do not send any messages, but perform basic operations. These methods are not embedded in the virtual image but in the virtual machine; they are written in machine code and are referred to as primitive methods. For example, the mathematical operations + \* / - are primitive methods.

There are about a hundred primitive methods; the number varies according to the implementation.

When a method is a primitive method, it begins with an expression of the form


```
<primitive #n>
```

where n is the number of the primitive.

If the primitive cannot be executed correctly, the expressions that follow it in the method are executed; this allows errors to be controlled.

# 6 Some Elements of the User Interface

Part of this book is devoted to the user interface in Smalltalk. However, at this point we want to give the reader some elements so that he or she may test the examples that have already been presented and those that are to follow in the book. It is worth repeating once again that learning Smalltalk is best done in a practical context; this means that the support of a machine, while not absolutely necessary, is highly desirable for tackling the remainder of this book.



## 6.1 General description

When compared with a traditional machine, the Smalltalk interface has three special features:


- a bitmap screen
- a three-button mouse
- multiple windows and pop-up menus.

A window is an area of the screen corresponding to the user interface of a particular application.

A pop-up menu is a menu that appears on the screen when the user presses one of the mouse buttons and which disappears when the button is released to execute the selected function.

### 6.2 Main menu

Depending on the window types, the pop-up menus vary. When the cursor is outside any window, the user can make the menu shown below appear by pressing the middle button of the mouse.



We shall briefly describe the main functions of this menu.


- restore display  
redisplays the complete screen,
- file list  
opens a window accessing the file system,
- browser  
opens a window viewing the classes and methods,
- workspace  
opens a window in which expressions can be evaluated,
- save  
saves the virtual image in a file,
- quit  
exits the Smalltalk system.

### 6.3 Workspaces

A workspace is a window in which it is possible to edit text and evaluate expressions.

In this type of window the left-hand button of the mouse is used to designate a position in the text or select a part of this text, showing it in inverse video.

The middle button on the mouse is used to edit the text or evaluate expressions.



The main functions of the middle button menu are:

again

which repeats the last text replacement,

undo

cancels the last editing function,

copy

copies the selected text into a buffer memory,

cut

removes the selected portion of text, placing a copy in the buffer memory

paste

inserts the text held in the buffer at the position marked by the cursor or replaces the selected section of text with the contents of the buffer,

do it

evaluates the selected portion of text,

print it

evaluates the selected portion of text and displays the result of the evaluation.

#### 6.4 Browsers

Browsers are windows used to browse through the classes and methods in the system.

For greater clarity of presentation, classes and methods are grouped into categories. These categories are arbitrary and have no effect on the system architecture.

A browser type window contains five sub-windows. The four at the top are windows that can access lists; the one at the bottom is a text editing window. The four top windows contain respectively:

- the class categories - classes belonging to the same category are classes which deal with similar problems
- classes
- the methods (or protocols) categories, methods that

System Browser				
Collections-Support	-----	copying	copy:from:in:rule:	
Graphics-Primitives	Cursor	displaying	magnifyBy:	
Graphics-Display Object	DisplayMedium	display box access	nextLifeGeneration	
'Graphics-Paths	DisplayObject	pattern	reflect:	
Graphics-Views	DisplayScreen	bordering	rotate2:	
Graphics-Editors	DisplayText	coloring	rotateBy:	
Graphics-Support	Form	image manipulation	shapeFill:interiorPoint:	
Kernel-Objects	InfiniteForm	printing	shrinkBy:	
Kernel-Classes	instance	fileIn/Out	spread:from:by:spacing:	
Kernel-Methods	class	editing	wrapAround:	
<b>copy: destRectangle from: sourcePt in: sourceForm rule: rule</b>				
"Make up a BitBit table and copy the bits"				
(BitBit				
destForm: self				
sourceForm: sourceForm				
halftoneForm: nil				
combinationRule: rule				
destOrigin: destRectangle origin				
sourceOrigin: sourcePt				
extent: destRectangle extent				
clipRect: (0@0 extent: width@height)) copyBits				
"				
[Sensor redButtonPressed] whileFalse:				
[Display copy: (30@30 extent: 300@300) from: Sensor cursorPoint in: Display rule: Form				

have a close functionality are grouped into the same category


- methods.

The bottom window serves mainly for editing classes and methods. When a class is selected and no method category is chosen, the bottom window contains the definition of the class. The figure below shows a browser containing the definition of the class Point.

System Browser				
Collections-Sequ	-----	accessing		
Collections-Text	Pen	comparing		
Collections-Array	Point	arithmetic		
Collections-Strea	Quadrangle	truncation and round		
Collections-Supp	Rectangle	c polar coordinates		
Graphics-Primitiv	instance	point functions		
Graphics-Display	class	converting		
Graphics-Paths		coercing		
Graphics-Views				
Object subclass: #Point				
instanceVariableNames: 'x y '				
classVariableNames: ''				
poolDictionaries: ''				
category: 'Graphics-Primitives'				

Having edited the definition of an existing class, it is possible to alter the characteristics of this class or create a new class. In order to save these characteristics, you need to use the **accept** function of the middle button mouse menu.

When a protocol is chosen, the user can create a new method or alter an existing method in the bottom window. To link the method into the system the accept function must again be used; this has the effect of compiling the text of the method. The figure below shows a browser viewing the text of the method `<=` in class Point.



The two types of windows that we have just described are sufficient for testing most of the examples in this book. The other principal types of windows are:

- Debuggers - for examining the state of computation during the evaluation of an expression
- FileLists - which allow access to the external files
- Transcripts - there is generally one single window of this type in which the system displays information relating to the progress of the Smalltalk session.

# 7 The Class Object

The class **Object** is the superclass of all other classes and consequently the methods that are defined in it are accessible by all Smalltalk-80 objects. Of course, these methods can be redefined in any class.

## 7.1 Identification of an object

The characteristics of an object depend on the class to which it belongs. Certain messages allow access to these characteristics. We have already met one of these:

class

which when sent to an object returns the class of that object.

For example,

2 class

returns **SmallInteger**,

\$A class

returns **Character**.

There are others:

isKindOf: aClass

allows one to find out if the receiver is an instance of the argument aClass or of one of its subclasses.

For example,

2 isKindOf: SmallInteger

returns true,

2 isKindOf: Number

returns true, because **SmallInteger** is a subclass of Number.

isMemberOf: aClass

allows one to find out if the receiver is a direct instance of aClass.

For example,

2 isMemberOf: SmallInteger

returns true,

2 isMemberOf: Number

returns false.

respondsTo: aSelector

allows one to find out if the receiver understands the message. For example,

2 respondsTo: #+

returns true,

2 respondsTo: #at:

returns false.

The selectors of messages are instances of class **Symbols** and are therefore prefixed with the hash.


## 7.2 Copies of objects

It is possible to copy an object, that is, create a new object having the same characteristics as the original.

There are two ways of making a copy, depending on whether the instance variables of the object are duplicated or not.

### 7.2.1 Non-duplicated instance variables

Instance variables are shared between the original and its copy, as in the following diagram.




If the original or its copy alters one of its instance variables, this alteration is also made for the other object. The message to make this copy is

```

shallowCopy
For example,
a <- 'ABC'
b <- a shallowCopy
assigns a copy of a to b.
  
```

### 7.2.2 Duplicated instance variables

Instance variables belong to the copy; there is no longer any dependence between the original and its copy. We have the following arrangement:



The copy is made in a recursive way at the level of the instance variables; that is, the objects referenced by the instance variables can themselves have their own instance variables. These instance variables are also duplicated up to the point of having objects without instance variables.

The message to make this copy is  
deepCopy

For example, if f1 is an instance of class **Form** (that is, a bitmap - see 15.2.3), then after evaluating the following expressions (the reverse message changes black to white and vice versa):

```
f2 <- f1 deepCopy.  
f3 <- f1 shallowCopy.  
f1 reverse
```

f2 will contain a copy of the original bitmap, whereas f1 and f3 will both refer to the original bitmap, which has been reversed.

### 7.2.3 Copy message

The class **Object** contains the definition of the method that will respond to the message

```
copy
```

This method calls `shallowCopy` directly. In the subclasses of **Object** it is possible to redefine this method according to the nature of the objects to be copied.

Some variables can thus be duplicated while others cannot.

## 7.3 Comparisons of objects

Two types of comparisons may be made. The equivalence and the equality of two objects can be tested.

### 7.3.1 Equivalence of two objects

Two expressions are equivalent if they refer to the same object. The message for testing equivalence is

```
==anObject
```

For example:

```
variable1 <- anObject.  
variable2 <- anObject.  
variable1==variable2
```

returns true, because variable 1 and variable 2 refer to the same object. We can also test if two objects are not equivalent with the message

```
~~anObject
```

### 7.3.2 Equality of two objects

Whereas equivalence is defined in class **Object**, and should not be redefined, equality should be defined appropriately for each class.

For example, two instances of class **Point** are considered to be equal if their x- and y-values are equal, but are not equivalent unless they are the same object:

```
variable1 <- 3@4  
variable2 <- 3@4  
variable1 = variable2 returns true  
but variable1==variable2 returns false.
```

### 7.3.3 Comparison with nil

There are two messages to compare an object with the nil object. The message

isNil

returns true if the receiver is nil. The message  
notNil

returns true if the receiver is not nil.

### 7.4 Indexed objects

We have seen how an object can have named instance variables and/or indexed instance variables. For objects with indexed instance variables, there are several messages that allow access to the objects referenced by these variables. These are

at: anIndex

which returns the object referenced by the instance variable that has the index anIndex.

at: anIndex put: anObject

allows the object anObject to be assigned to the instance variable of index anIndex.

basicAt: AnIndex

is equivalent to at:, but must not be redefined in the subclasses of Object.

basicAt: anIndex put: anObject

is equivalent to at:put:, but must not be redefined in the subclasses of Object.

size

returns the number of indexed instance variables in the receiver.

basicSize

is equivalent to size, but must not be redefined in the subclasses of Object.

For example,

8 size	returns 0
--------	-----------

'ABC' size	returns 3
------------	-----------

#(1 2 3 (2 4)) size	returns 4
---------------------	-----------

(#(1 2 3 (2 4)) at: 4) size	returns 2
-----------------------------	-----------

### 7.5 Representation of objects

Smalltalk-80 represents objects internally in a binary form which is not readable by users. Therefore, it also provides methods to generate representations of objects in a human-readable form. There are two possible forms of external representation:

- a first external representation in summary form,
- a second representation which is an expression that, if evaluated, will recreate the object.

There are two methods in class Object that create such representations. These are

printString

which returns a character string giving a summary representation of the receiver, and

storeString

which returns a character string giving a syntactic representation of the receiver.

For example,

#(2 3) printString

will return the character string '(2 3)',

#(2 3) storeString

will return the character string '#(2 3)'.

(Set with: l with: 2) printString

will return the character string 'Set(l 2)',

(Set with: l with: 2) storeString

will return the character string '((Set new) add: 1; add: 2; yourself)' which allows an identical copy of the receiver to be reconstructed.

## 7.6 Control of errors

The search for a method begins at the level of the class of the receiver; it continues through the hierarchy of superclasses. If the method is not found, there is an error; this is conveyed by the message doesNotUnderstand: with as its argument the message that caused the error. This method signals the error to the programmer; it is defined within the class Object and is therefore accessible to all objects. This method can be redefined for a particular class, for example, to make an automatic correction.

There are a number of other methods defined within class Object that are intended to control errors. These methods are:

error:aString

indicates to the user that there is an error and displays the string aString,

primitiveFailed

indicates to the user that a primitive method has failed,

shouldNotImplement

indicates to the user that the method that should respond exists but should not be used,

subclassResponsibility

indicates to the user that the method must be redefined in the class of the receiver.

## 7.7 Control of messages

It is sometimes useful to choose which message to send to an object by computing the selector within an expression.

One cannot simply write:

receiver variableSelector

because the receiver will try to execute the method variableSelector.

There are, however, five methods that are available to deal with this problem:

```
perform:aSymbol
sends to the receiver the unary message aSymbol,
perform:aSymbol with:anObject
sends to the receiver the keyword or binary message aSymbol with argument anObject,
perform:aSymbol with:anObject with:anOtherObject
sends to the receiver the keyword message aSymbol with arguments anObject and anOtherObject,
perform:aSymbol with: anObject1 with: anObject2 with:
anObject3
sends to the receiver the keyword message aSymbol with arguments anObject1, anObject2 and anObject3,
perform:aSelector withArguments: anArray
sends to the receiver the message aSelector with the arguments contained in anArray. There will be an error if the number of elements in the array does not correspond to the number of arguments expected by the selector.
```

For all the messages, there will be an error if the number of arguments does not match the number expected by the selector. For example,

```
sort: aSelector
|test|
self size
to: 1
by: -1
do: [:j | 1 to: j - 1 do:
[:k|
    test <- (self at:k)
    perform: aSelector
    with: (self at: k + 1).
test ifFalse:
    [self swap: k with: k + 1]]]
```

is a method for sorting an array, whose argument aSelector is the sort criterion. This method must be defined in the class SequenceableCollection.

```
#(3 7 0 1) sort: # <
will return the array #(0 1 3 7)
#(3 7 0 1) sort: # >
will return the array #(7 3 1 0).
```

## 7.8 Dependencies between objects

In the Smalltalk-80 system it is possible to define dependencies between objects in such a way that, when an object is dependent on a second one, any alteration of the second is signalled to the first.

The class Object has defined within it methods that allow dependencies between objects to be created explicitly. These methods are:

addDependent:anObject

anObject is added to the list of objects depending on the receiver,

removeDependent:anObject

anObject is removed from the list of objects depending on the receiver,

dependents

returns an instance of **OrderedCollection** containing the objects depending on the receiver,

release

removes all the dependencies on the receiver, leaving it with no dependency on any other object,

changed

the receiver signals to all its dependent objects that it has changed by sending them the message update,

changed:aParameter

the receiver signals to all its dependent objects that it has changed by sending them the message update:aParameter, whose argument aParameter indicates the nature of the change,

update:aParameter

is the method that must respond to the message sent by changed or changed:. This method is used to update the receiver,

broadcast:aMessage

sends the unary message aMessage to all the objects depending on the receiver,

broadcast:aMessage with:anObject

sends the keyword message aMessage with the argument anObject to all the objects depending on the receiver.

We will take as an example the channel-selecting keys of a TV set. Only one of these keys can be selected at once. One key is characterised by its state (selected or not). The class Key of the keys can thus be defined with the following methods

class name	Key
superclass	Object
name of instance variables	selected
class methods	

**new**

  ↑super new initialise

instance methods

**initialise**

  selected <- false

**select**

  selected ifFalse: [selected <- true. self changed]

```
update: aKey
  aKey == self iffFalse: [selected <- false]
```

The class Keyboard is also defined

```
class name           Keyboard
superclass          Object
instance variables name   keys
class methods
```

```
new: numberOfRowsKeys
  ↑super new createKeys: numberOfRowsKeys

instance methods
select: aNumberOfKey
  (keys at: aNumberOfKey) select
createKeys: numberOfRowsKeys
  keys <- Array new: numberOfRowsKeys.
  1 to: numberOfRowsKeys do:
    [:index | keys at: index put: Key new].
  keys do:
    [:aKey |
      keys do:
        [:aDependentKey |
          aKey == aDependentKey iffFalse:
            [aKey addDependent: aDependentKey]]]
```

The method `createKeys:` allows a keyboard to be initialised with a certain number of keys. The created keys are neither selected nor deselected at the start. A key is dependent on all the other keys. Note the pseudo-variable `super` in the method `new:` which avoids recursive calls of this method.

The example can be made to run with:

```
keyboard <- Keyboard new: 5.
keyboard select: 1.
keyboard select: 4.
```

Selection of key 1 of the keyboard will cause the message `update:` to be sent to all the other keys, which will deselect them whatever their status.

### 7.9 Primitive messages

There are a certain number of primitive methods described in the class `Object`. These methods are:

```
become:anObject
allows the pointers of the receiver and of the argument
anObject to be exchanged. All the variables that refer
to the receiver will refer to the argument and vice
versa. There will be an error if one of the two objects
is an instance of the class SmallInteger,
```

```
instVarAt: anIndex
```

returns the named instance variable of rank anIndex. The numbering of named instance variables corresponds to their order of definition,


instVarAt: anIndex put: anObject  
assigns the object anObject to the named instance variable of rank anIndex,

nextInstance  
returns the instance following the receiver in the list of all instances of the class of the receiver. It returns nil if all the instances have been returned.

## 8 Description of Architecture of Classes

We know that each object is an instance of a class. As a class is itself an object, it is the instance of a class that we call its metaclass. To access the metaclass of a class we send it the message `class`.


The hierarchy that appears at the level of the metaclasses is the same as that which appears at the level of the classes; that is, if a class A is a subclass of a class B, the metaclass of A will be a subclass of the metaclass of B, as in the diagram below.



Each metaclass has only one instance and thus there are exactly as many metaclasses as classes. The metaclasses are also objects and thus belong to a class that is called **Metaclass**.

Metaclass is a class and thus has a metaclass which is an instance of Metaclass. Metaclass and its metaclass are both instances of each other.

Object is the only class that does not have a superclass; on the other hand its metaclass has a superclass called **Class**. The diagram below illustrates this architecture.



### 8.1 Class Behavior

The class **Behavior** defines the elementary methods necessary to objects that have instances. The instances of Behavior are objects representing the classes. As instance variables they have a dictionary of methods, a pointer to a superclass and a list of subclasses.

#### 8.1.1 Dictionary of methods

The methods described in a class are held in a dictionary called the dictionary of class methods. The keys to the dictionary are the selectors, and the objects assigned to these keys are methods in compiled form.

#### Creation

The methods that allow the method dictionary to be manipulated are:

```
methodDictionary: aDictionary  
assigns a dictionary to the method dictionary of the receiver;  
addSelector: aSelector withMethod: aMethod  
adds a method to the method dictionary of the receiver,  
with selector aSelector and for compiled method aMethod;  
removeSelector: aSelector  
removes from the method dictionary of the receiver that  
method with selector aSelector.
```

#### Compilation

The methods that allow one or more methods to be compiled are:

```
compile: code  
compiles the text of the expressions contained in the  
argument code. The argument code is a string of characters.  
An error is produced if the argument code cannot be  
compiled,  
compile: code notifying: requestor  
compiles the text of the expressions contained in the  
argument code. If the argument code cannot be compiled,  
an error is signalled to the requestor object by a message,  
recompile: aSelector  
compiles the method in the method dictionary whose key  
is aSelector;  
decompile: aSelector  
decompiles the method in the method dictionary whose key  
is aSelector. It returns a string of characters represent-  
ing the decompiled form of the compiled method;  
compileAll  
compiles all the methods in the method dictionary;  
compileAllSubclasses  
compiles all the methods contained in the method dictio-  
naries of the subclasses of the receiver.
```

### Access to selectors or methods

The methods that allow access to the selectors or methods in the method dictionary are:

selectors

returns an instance of Set containing the selectors in the method dictionary of the receiver,

allSelectors

returns an instance of Set containing all the selectors in the method dictionary of the receiver and of its superclasses, that is, all the selectors assigned to the messages that are understandable by an instance of the receiver,

compiledMethodAt: aSelector

returns the compiled method assigned to the selector aSelector contained in the method dictionary of the receiver,

sourceCodeAt: aSelector

returns a string of characters corresponding to the method assigned to the selector aSelector in the method dictionary of the receiver,

sourceMethodAt: aSelector

returns an instance of **Text** corresponding to the method assigned to the selector aSelector in the method dictionary of the receiver. In this book the name of the method is shown in bold type (as are class names).

Some examples of the use of these methods might be:

Number selectors

returns Set (to:do: even roundTo: storeOn: .... log: tan rounded)

Number allSelectors

returns Set (raisedTo: asPoint between:and: negated .... min:max: truncated)

Number sourceCodeAt: #abs

returns:

`abs

"Answer a Number that is the absolute value (positive magnitude) of the receiver"

self <0 ifTrue: [↑self negated]

iffFalse: [↑self]'

### Testing

The contents of the methods dictionary can be tested in order to find out the methods that are accessible by the instances of the receiver or the methods that possess certain characteristics. The methods that allow these tests to be carried out are:

hasMethods

returns true if the methods dictionary of the receiver is not empty,

includesSelector: aSelector

returns true if the instances of the receiver can res-

pond to the message with selector aSelector, that is, if the selector is a key in the method dictionary of the receiver or of its superclasses,

whichClassIncludesSelector: aSelector  
returns the first class contained in the sequence of superclasses of the receiver whose method dictionary contains the key aSelector; it returns nil if no class is found,

whichSelectorsAccess: instanceVariableName  
returns an instance of Set containing the selectors of the methods in the method dictionary of the receiver that access the instance variable called instanceVariableName,

whichSelectorsReferTo: anObject  
returns an instance of Set containing the selectors of methods in the method dictionary of the receiver that access the object anObject.

scopeHas: aName ifTrue: aBlock  
executes the block aBlock if the variable called aName is accessible by the receiver.

Here are some examples of the use of these methods:

Number hasMethods  
returns true because the method dictionary of the class Number is not empty;

Number includesSelector: # ==  
returns false because this method is not in the method dictionary of the class Number but in the method dictionary of the class Object;

Number canUnderstand: # ==  
returns true because this method is contained in the method dictionary of a superclass of the class Number;

Number whichClassIncludesSelector: # ==  
returns class Object.

### 8.1.2 Instances

The instances of a class are created by sending the message new or new: to it. These methods are defined in class Behavior, but can be redefined in any subclass. If a class redefines the method new or new:, the method will necessarily contain the message super new in order to avoid a recursive call of itself. It can be useful to redefine these methods if, for example, you want to initialise instance variables with values other than nil.

#### Creation

The four methods defined in class Behavior that allow instances to be created are:

new  
returns a new instance of the receiver. This instance

```
has no indexed instance variables;
basicNew
is identical to new, but this method must not be
redefined;
new: anInteger
returns a new instance of the receiver with a number of
indexed instance variables given by anInteger;
basicNew: anInteger
is identical to new:, but this method must not be
redefined.
```

Here are some examples of the use of these methods:

```
matrix1 <- Array new.
matrix2 <- Array new: 3.
```

matrix1 is an instance of the class Array which has no indexed instance variables. The message new is equivalent to the message new: 0 for those classes that have indexed instance variables.

### Accessing

There are different methods that allow one or more instances of a class to be accessed. These methods are:

```
allInstances
returns an instance of OrderedCollection containing all
the direct instances of the receiver (without the instances
of its subclasses).
someInstance
returns a direct instance of the receiver,
instanceCount
returns the number of direct instances of the receiver,
instVarNames
returns an instance of Array containing the names of the
named instance variables,
subclassInstVarNames
returns an instance of Set containing the names of the
named instance variables of the subclasses of the receiver,
allInstVarNames
returns an instance of Array containing the names of the
named instance variables of the receiver and of its
superclasses,
classVarNames
returns an instance of Set containing the names of the
class variables of the receiver,
allClassVarNames
returns an instance of Set containing the names of the
class variables of the receiver and of all its super-
classes,
sharedPools
returns an instance of Set containing the names of the
dictionaries that contains the shared variables,
allSharedPools
```

returns an instance of Set containing the names of the dictionaries that contain the shared variables defined in the receiver and its superclasses.

#### **Enumeration**

There are certain methods that allow an operation to be carried out on each instance. These methods are:

allInstancesDo: aBlock

evaluates the block aBlock for each instance of the receiver,

allSubInstancesDo: aBlock

evaluates the block aBlock for all the subclass instances of the receiver.

#### **Testing**

Other methods allow the instance variables of the receiver to be tested

isPointers

returns true if the instance variables of the receiver are stored in the form of pointers,

isBits

returns true if the instance variables of the receiver are stored in the form of bits,

isBytes

returns true if the instance variables of the receiver are stored in the form of bytes,

isWords

returns true if the instance variables of the receiver are stored in the form of words,

isFixed

returns true if the instances of the receiver have no indexed instance variables,

isVariable

returns true if the instances of the receiver have indexed instance variables,

instSize

returns the number of named instance variables of the receiver.

Some examples are:

Number isFixed

returns true because the instances of Number do not have indexed instance variables,

Array instSize

returns 0 because the instances of Array do not have named instance variables.

#### **8.1.3 Manipulation of the class hierarchy**

##### **Creation**

Methods defined in class Behavior allow alteration of

the class hierarchy by setting a class's superclass and subclasses. These methods are:

```

superclass: aClass
fixes the superclass of the receiver at aClass,
addSubclass: aClass
adds the subclass aClass to the receiver,
removeSubclass: aClass
removes the subclass aClass from the receiver.
```

### **Accessing**

There are methods that allow the hierarchy of the receiver to be accessed. These are:

```

subclasses
returns an instance of Set containing the immediate subclasses of the receiver,
allSubclasses
returns an instance of OrderedCollection containing all the subclasses of the receiver, that is, its immediate subclasses and the subclasses of its subclasses, and so on,
```

```

withAllSubclasses
returns an instance of OrderedCollection containing the receiver and all of its subclasses,
```

```

superclass
returns the superclass of the receiver,
allSuperclasses
returns an instance of OrderedCollection containing the superclass of the receiver followed by the other superclasses in hierarchical order; the last element is always Object.
```

For example,

```

Number subclasses
returns Set (Float Fraction Integer),
Number allSubclasses
returns OrderedCollection (Fraction Integer Float SmallInteger LargePositiveInteger LargeNegativeInteger),
Integer superclass
returns the object representing the class Number,
Integer allSuperclasses
returns OrderedCollection (Number Magnitude Object).
```

### **Testing**

There are certain methods that allow the hierarchy of the receiver to be tested. These methods are:

```

inheritsFrom: aClass
returns true if the class aClass belongs to the sequence of superclasses of the receiver,
kindOfSubclass
returns a string describing the type of the receiver.
There are four types of class in Smalltalk-80:
- classes which have no indexed instance variables
```

where the named instance variables represent pointers to other objects; the returned string will be 'subclass:'

- those classes which have indexed instance variables that can represent
- pointers; the returned string will be 'variableSubclass:'
- bytes; the returned string will be 'variableByteSubclass:'
- words; the returned string will be 'variableWordSubclass:'

For example:

```

Integer inheritsFrom: Number
returns true because Number is the superclass of Integer,
Integer kindOfSubclass
returns 'subclass:',
Float kindOfSubclass
returns 'variableWordSubclass:'.
```

### **Enumeration**

There are certain methods that allow operations to be carried out on a part of the class hierarchy related to the receiver. These methods are:

allSubclassesDo: aBlock  
evaluates the block aBlock for all the subclasses of the receiver, that is, for its immediate subclasses and the subclasses of its subclasses, and so on

allSuperclassesDo: aBlock  
evaluates the block aBlock for all the superclasses of the receiver,

selectSubclasses: aBlock  
evaluates the block aBlock for all the subclasses of the receiver and groups in an instance of Set all the classes for which evaluation returns true. This instance of Set is returned.

selectSuperclasses: aBlock  
evaluates the block aBlock for all the superclasses of the receiver and groups in an instance of Set all the classes for which evaluation returns true. This instance of Set is returned.

### **8.2 The classes ClassDescription and Class**

The description of a class is not totally described by the class Behavior. Behavior does not allow the names of instance variables and class variables to be stored; nor does it allow a name to be given to the class or a comment describing the class to be stored.

The class **ClassDescription**, which is a subclass of the class Behavior, and the class **Class**, which is a subclass of the class **ClassDescription**, are the classes

that allow a name to be given to a class, a comment to be assigned to it and the names of named instance variables to be stored.

### 8.2.1 Access to the description of a class

The additional methods that allow access to the description of a class are:

```
    name
returns a string representing the name of the receiver,
    comment
returns a string giving the comment assigned to the receiver,
    comment: aString
assigns the string aString to the commentary associated with the receiver,
    addInstVarName: aString
adds to the list of named instance variables of the receiver the named instance variable whose name is aString,
    removeInstVarName: aString
removes from the list of named instance variables of the receiver the named instance variable whose name is aString. It returns an error if aString is not the name of a named instance variable.
```

### 8.2.2 Category of classes and messages

To clarify the description of the class hierarchy, the classes are grouped into categories. This grouping is arbitrary and simply allows classes that relate to the same problem to be grouped together; for example, the category Numeric-Numbers groups classes Number, Integer, Float, Fraction, SmallInteger, LargePositiveInteger, LargeNegativeInteger, Random. This grouping has no effect on the hierarchy of the classes.

Just as classes can be grouped into categories, so can the methods of classes.

The methods that allow access to the category of a class or a method are:

```
    category
returns the category of the receiver,
    category: aString
assigns to the category of the receiver the string aString,
    removeCategory: aString
removes all the methods from the methods dictionary of the receiver whose category is aString,
    whichCategoryIncludesSelector: aSelector
returns the category of the method whose selector is aSelector; it returns nil if the method is not found in the method dictionary of the receiver.
```

### 8.2.3 Copy of messages

Some methods are provided that copy one or more methods from another class. These methods are:

copy: aSelector from: aClass

copies the method of selector aSelector described in the class aClass and places it in the method dictionary of the receiver,

copy: aSelector from: aClass classified: aCategory  
copies the method of selector aSelector described in the class aClass and places it in the method dictionary of the receiver under category aCategory,

copyAll: arrayOfSelectors from: aClass  
copies the methods whose selectors are contained in the arrayOfSelectors described in the class aClass and places them in the method dictionary of the receiver,

copyAll: arrayOfSelectors from: aClass classified:  
aCategory  
copies the methods whose selectors are contained in the arrayOfSelectors described in the class aClass and places them in the method dictionary of the receiver under the category aCategory,

copyAllCategoriesFrom: aClass  
copies all the methods described in the class aClass and places them in the method dictionary of the receiver,

copyCategory: aCategory from: aClass  
copies all the methods described in the class aClass under the category aCategory and places them in the method dictionary of the receiver,

copyCategory: aCategory from: aClass classified:  
aNewCategory  
copies all the methods described in the class aClass under the category aCategory and places them in the method dictionary of the receiver under the category aNewCategory.

### 8.2.4 Compilation of methods

There are two particular methods that allow source code to be compiled and the result placed under a category. These two methods are:

compile: aSourceCode classified: aCategory  
compiles the text of the expressions contained in the argument aSourceCode and places the compiled method obtained in the method dictionary of the receiver under category aCategory. An error is indicated to the programmer if the argument aSourceCode cannot be compiled:

compile: aSourceCode classified: aCategory requestor:  
aRequestor

compiles the text of the expressions contained in the argument aSourceCode and places the compiled method obtained in the methods dictionary of the receiver under category aCategory. An error is sent to the argument

aRequestor if the argument aSourceCode cannot be compiled.

#### 8.2.5 Access to variable names

Three methods allow variable names to be returned in the form of strings. These three methods are:

    classVariablesString

returns a string of characters containing the names of all the class variables of the receiver,

    instanceVariablesString

returns a string of characters containing the names of all the instance variables of the receiver,

    sharedPoolsString

returns a string of characters containing the names of all the dictionaries that contain the shared variables of the receiver.

#### 8.2.6 Saving a class on a file

Three methods allow the description of a class to be saved on a file. These methods are:

    fileOutOn: aFile

saves the description of the receiver on the file aFile which is an instance of the class **FileStream**,

    fileOutCategory: aCategoryName

creates a file whose name is the name of the receiver concatenated to the string 'st' and places in it all the methods whose category is aCategoryName,

    fileOutChangedMessages: aListOfChanges on: aFile

the argument aListOfChanges is made up of class/message couplets; it is an instance of the class **Set**. The description of the couplets class/message is saved on the file aFile, which is an instance of the class **FileStream**.

The class **ClassDescription** possesses another subclass called **Metaclass**.

### 8.3 The class Metaclass

The metaclasses are instances of the class **Metaclass**. We have seen that the role of a metaclass is to allow the initialisation of class variables and the creation of instances of the class that is assigned to it. Each metaclass contains an instance variable called **thisClass** which allows its unique instance to be referenced.

### 8.4 Multiple inheritance

We have seen that a class inherits variables and methods from its superclass. In the general case a class may inherit from a single other class (this is so for all the classes of the Smalltalk system), but it is also possible for a class to inherit from several other classes; this is called multiple inheritance.

Multiple inheritance allows the properties of two or more existing classes in the system to be fused into one class. Multiple inheritance imposes certain additional constraints that arise from the conflicts that may appear between the different superclasses:

- A class cannot inherit from two classes that have a same instance variable (that is, a variable having the same name).

- When creating a class with multiple superclasses, there may be a conflict between different inherited methods with the same name. When this occurs, a new method with the same name is automatically created in the new class which signals an error if activated. If the user wishes to resolve the conflict, he should edit the new method to call one of the inherited methods, designating the method by its name, prefixed with the name of its class and a full stop. For example:

Point.+

designates the method of addition of two points.

To implement multiple inheritance, the Smalltalk system possesses a subclass of Metaclass called **MetaclassForMultipleInheritance**, whose instances have an instance variable called otherSuperclasses which contains the list of additional superclasses. The metaclasses of multiple inheritance classes are instances of MetaclassForMultipleInheritance.

To create a class that inherits from several other classes the following message must be sent to the class Class

```
named: aSymbol  
superclass: namesOfSuperclasses  
instVariableNames: namesOfInstanceVariables  
classVariableNames: namesOfClassVariables  
category: nameOfCategory
```

For example, consider the case of simulating a transport network in which the following two classes are defined:

MotorVehicle

which has two instance variables, namely:

tankCapacity and operatingHours

WheeledVehicles

which has one instance variable:

mileage

The class Automobile embraces the two aspects of classes MotorVehicle and WheeledVehicles; thus, to define it we need to construct it as the subclass of these two classes:

Class named: #Automobile

```
superclasses: 'MotorVehicle WheeledVehicles'  
instVariableNames: 'numberOfSeats'  
classVariableNames: ''
```

```
category: 'simulation-network'
```

Those methods that conflict will have to be redefined; for example, the method `printOn:` which generates a textual representation of an object.

For multiple inheritance to be attractive to use, the number of conflicting methods needs to be small; that is, the two superclasses should have few aspects in common.

# 9 The Magnitude Classes

The magnitude classes contain the objects on which an ordering relation can be established. They are objects for which the concept of size has a meaning.

## 9.1 The class **Magnitude**

The class **Magnitude** is a subclass of the class **Object**. All the ordered classes are subclasses of **Magnitude**; this is why **Magnitude** describes comparison methods that may be used by all the instances of its subclasses.

These methods are:

< aSize

returns true if the receiver is less than the argument aSize,

<= aSize

returns true if the receiver is less than or equal to the argument aSize,

> aSize

returns true if the receiver is greater than the argument aSize,

>= aSize

returns true if the receiver is greater than or equal to the argument aSize,

between: min and: max

returns true if the receiver is greater than or equal to the argument min and less than or equal to the argument max,

min: aSize

returns the minimum of the receiver and the argument aSize,

max: aSize

returns the maximum of the receiver and the argument aSize,

= aSize

this message must be redefined in subclasses of **Magnitude**. In class **Magnitude** the method = is described by the unique expression self subclassResponsibility.

For example:

2 > 1	returns true
1 > 2	returns false
2 between: 1 and: 3	returns true
2 min: 5	returns 2
2 max: 5	returns 5

## 9.2 The class Date

The class **Date** is a subclass of **Magnitude**. Its instances are dates that represent a day, a month and a year of the calendar. It knows the days of the week and the months of the year by storing index/name pairs.

There are a number of methods that allow these dates to be manipulated.

### 9.2.1 Creation of instances

When a new date is created, it is initialised with special values. The different creation methods are:

**today**

returns a new instance of the class **Date** representing today's date,

**fromDay: numberOfWorkDays**

returns a new instance of **Date** representing the date obtained by adding **numberOfWorkDays** to the first of January 1901,

**newDay: aDay month: aMonth year: aYear**

returns a new instance of **Date** whose day is **aDay**, the month **aMonth** and the year **aYear**. The arguments **aDay** and **aYear** are integers, whereas the argument **aMonth** is a symbol.

**newDay: numberOfWorkDays year: aYear**

returns a new instance of **Date** representing the date obtained by adding **numberOfWorkDays** to the first of January of the year **aYear**.

For example:

<b>Date today</b>	returns 22 October 1986
-------------------	-------------------------

<b>Date fromDays: 3</b>	returns 4 January 1901
-------------------------	------------------------

<b>Date newDay: 8 month: #Oct year: 86</b>	returns 8 October 1986
--------------------------------------------	------------------------

<b>Date newDay: 3 year: 86</b>	returns 4 January 1986
--------------------------------	------------------------

### 9.2.2 Information about the calendar

It is possible to obtain general information about the calendar through the following methods:

**dayOfWeek: nameOfDay**

returns the number of the day **nameOfDay** in the week (1 for **#Monday**, etc),

**nameOfDay: aNumber**

returns a symbol representing the day of the week with number **aNumber**,

**indexOfMonth: nameOfMonth**

returns the number of the month **nameOfMonth** in the year (1 for **#January** or **#Jan**, etc). The name of a month can be abbreviated,

**nameOfMonth: aNumber**

returns a symbol representing the month of the year of number **aNumber**,

```

daysInMonth: nameOfMonth forYear: aYear
returns the number of days of the month nameOfMonth for
the year aYear,
daysInYear: aYear
returns the number of days of the year aYear,
leapYear: aYear
returns 1 if the year is a leap year, otherwise 0,
dateAndTimeNow
returns an instance of Array whose first element is the
actual date and second element the actual time.

```

For example:

Date dayOfWeek: #Tuesday	returns 2
Date nameOfDay: 3	returns Wednesday
Date indexOfMonth: #Mar	returns 3
Date nameOfMonth: 4	returns April
Date daysInMonth: #February forYear: 1984	returns 29
Date daysInYear: 1984	returns 366
Date leapYear: 1984	returns 1
Date dateAndTimeNow	returns (a Date a Time)

### 9.2.3 Arithmetic

There are three methods for carrying out arithmetic operations on dates. These methods are:

addDays: numberOfDays	
returns a date obtained by adding numberOfDays to the receiver,	
subtractDays: numberOfDays	
returns a date obtained by subtracting numberOfDays from the receiver,	
subtractDate: aDate	
returns the number of days separating the receiver from the argument aDate.	

For example:

Date today addDays: 10	returns 1 November 1986
Date today subtractDays: 10	returns 12 October 1986
Date today subtractDate:	
(Date newDay: 5 month: #Oct year: 1986)	
	returns 17

### 9.2.4 Conversion

In order to be able to carry out calculations on dates and hours both are converted into seconds. The method that allows a date to be converted into seconds is:

asSeconds

which returns the number of seconds that have elapsed since 1 January 1901 up to the beginning of the current day.

### 9.3 The class Time

The class **Time** is a subclass of the class **Magnitude**. The instances of **Time** represent an instant of a day.

#### 9.3.1 Creation of instances

Three methods allow an instance of **Time** to be created. These three methods are:

now

returns an instance of **Time** representing the current time,

fromSeconds: numberOfSeconds

returns an instance of **Time** representing the time calculated by adding **numberOfSeconds** to midnight:

dateAndTimeNow

returns an instance of **Array** whose first element is the actual date and whose second element is the actual time.

#### 9.3.2 Information about the time

The class **Time** can respond to a certain number of general messages about the time:

millisecondsClockValue

returns the number of milliseconds since the clock passed zero,

millisecondsToRun: aBlock

returns the number of milliseconds it takes to execute the block **aBlock**,

timeWords

returns the number of seconds that have elapsed since 1 January 1901 at the Greenwich meridian.

totalSeconds

returns the number of seconds that have elapsed since 1 January 1901, allowing for corrections for time zones and summertime.

#### 9.3.3 Conversions

In order to be able to carry out calculations on dates and times, it is useful to be able to convert these into seconds. The method that allows this to be done is:

asSeconds

returns the number of seconds that have elapsed since midnight.

#### 9.3.4 Arithmetic

Two methods allow arithmetic operations to be carried out on times. These are:

addTime: aPeriod

returns an instance of **Time** representing the time calculated by adding **aPeriod** to the receiver. The argument **aPeriod** must be an instance of **Date** or of **Time**,

subtractTime: aPeriod

returns an instance of Time representing the time calculated by subtracting aPeriod from the receiver. The argument aPeriod must be an instance of Date or of Time.

#### 9.4 The class Character

The class **Character** is also a subclass of the class **Magnitude**. It is an ordered class because the instances of this class represent ASCII codes. There are therefore 256 instances in this class; they belong to the Smalltalk system and cannot be altered. It is not possible to make new Characters.

##### 9.4.1 Access to the instances

The methods that allow access to the instances of the class **Character** are:

```
    value: anInteger
returns the Character with ASCII code anInteger,
    digitValue: anInteger
returns the character representing the digit corresponding to anInteger (when anInteger is greater than 9, it returns the first letters of the alphabet),
    asciiValue
returns the ASCII code of the receiver,
    digitValue
returns the number corresponding to the digit represented by the receiver.
```

For example:

Character value: 67	returns \$C
Character digitValue: 1	returns \$1
Character digitValue: 10	returns \$A
\$A asciiValue	returns 65
\$B digitValue	returns 11

##### 9.4.2 Testing the instances

The following methods test for the nature of the instances of the class **Character**:

```
    isAlphaNumeric
returns true if the receiver is a letter or a digit,
    isDigit
returns true if the receiver is a digit (0-9),
    isLetter
returns true if the receiver is a letter,
    isLowerCase
returns true if the receiver is a lower case letter,
    isUpperCase
returns true if the receiver is an upper case letter,
    isSeparator
returns true if the receiver is one of the separator characters (cr, tab, line feed, form feed),
    isVowel
returns true if the receiver is an upper or lower case vowel.
```

For example:

\$+	isAlphaNumeric	returns false
\$2	isDigit	returns true
\$t	isLetter	returns true
\$G	isLowerCase	returns false
\$G	isUpperCase	returns true
\$;	isSeparator	returns false
\$a	isVowel	returns true

The following chapter describes another subclass of Magnitude, namely the class Number.

# 10 The Numeric Classes

The instances of the numeric classes are the objects that we use to carry out arithmetic operations or mathematical calculations. Each type of numeric value is represented by a class. All the numeric classes are subclasses of class Number which has three direct subclasses.

- the class **Float**
- the class **Fraction**
- the class **Integer**.

The class Integer has three subclasses

- the class **SmallInteger**
- the class **LargePositiveInteger**
- the class **LargeNegativeInteger**.

The classes Number and Integer are abstract subclasses, that is, they have no direct instance.

## 10.1 The class Number

A certain number of messages are understood by all the numeric objects. These messages allow arithmetic operations and comparisons to be made. The majority of these messages are redefined in the subclasses of class Number. The messages common to all the subclasses of class Number are described below.

### 10.1.1 Arithmetic

The messages allowing arithmetic operations on numbers to be carried out are:

+ aNumber

returns the sum of the receiver and the argument aNumber,

- aNumber

returns the difference between the receiver and the argument aNumber,

\* aNumber

returns the product of the receiver and the argument aNumber,

/ aNumber

returns the quotient between the receiver and the argument aNumber. If the division is not whole the result is an instance of the class Fraction,

```

// aNumber
returns the whole part of the division of the receiver
by the argument aNumber,
\\ aNumber
returns the whole part of the remainder of the division
of the receiver by the argument aNumber,
abs
returns the absolute value of the receiver,
negated
returns the opposite of the receiver,
quo: aNumber
returns the whole part of the division of the receiver
by the argument aNumber,
rem: aNumber
returns the whole part of the remainder of the division
of the receiver by the argument aNumber,
reciprocal
returns the inverse of the receiver, producing an error
if the receiver is equal to 0.

```

For example:

Expressions	Results
2 + 3	5
2 - 3	-1
2 * 3	6
4/2	2
4/3	(4/3) instance of the class Fraction with denominator 3 and numerator 4
4//3	1
-4//3	-2
4 \\ 3	1
-4 \\ 3	1
-5 abs	5
5 negated	-5
-4 quo: 3	-1
-4 rem: 3	-1
7 reciprocal	(1/7)

### 10.1.2 Mathematics

The messages that allow mathematical operations to be carried out on numbers are:

```

exp
returns the exponent of the receiver,
ln
returns the naperian logarithm of the receiver,
log: aNumber

```

```

returns the logarithm to base aNumber of the receiver,
    floorLog: aNumber
returns the whole part of the logarithm to base aNumber
of the receiver,
    raisedTo: aNumber
returns the receiver raised to the power aNumber,
    raisedToInteger: anInteger
returns the receiver raised to the power anInteger. The
argument anInteger must be an instance of class Integer.
    sqrt
returns the square root of the receiver,
    squared
returns the square of the receiver,
    sin
returns the sine of the receiver expressed in radians,
    cos
returns the cosine of the receiver expressed in radians,
    tan
returns the tangent of the receiver expressed in
radians,
    arcSin
returns the arcsine of the receiver in radians,
    arcCos
returns the arccosine of the receiver in radians,
    arctan
returns the arctangent of the receiver in radians.
For example:

```

Expressions	Results
1 exp	2.71828
1 ln	0.0
2 log: 2	1.0
3 floorLog: 2	1.0
4 raisedTo: 2.4	27.8576
4 raisedToInteger: 3	64
2 sqrt	1.41421
9 squared	81

### 10.1.3 Tests

There are the following messages that allow the characteristics of a number to be tested.

even

returns true if the receiver is an even number,

odd

returns true if the receiver is an odd number,

negative

returns true if the receiver is negative,

positive

returns true if the receiver is positive or zero,

strictlyPositive

```

returns true if the receiver is strictly positive,
      sign
returns 1 if the receiver is strictly positive; -1 if it
is strictly negative; 0 if it is zero.

```

#### 10.1.4 Truncation and rounding

Several messages allow a number to be truncated or rounded. These messages are:

```

ceiling
returns the smallest integer greater than or equal to
the receiver,
floor
returns the largest integer less than or equal to the
receiver (that is, its whole part),
truncated
returns the receiver without its decimal part,
truncateTo: aNumber
returns the largest multiple of the argument aNumber
whose absolute value is less than or equal to the absolute
value of the receiver,
rounded
returns the integer closest to the receiver,
roundTo: aNumber,
returns the multiple of the argument aNumber closest to
the receiver.

```

For example:

Expressions	Results
2.3 ceiling	3
-2.3 ceiling	-2
2.3 floor	2
-2.3 floor	-3
2.3 truncated	2
-2.3 truncated	-2
2.3 truncateTo: 0.7	2.1
-2.3 truncateTo: 0.7	-2.1
2.3 rounded	2
2.6 rounded	3
-2.3 rounded	-2
-2.6 rounded	-3
2.3 roundTo: 0.7	2.1
2.6 roundTo: 0.7	2.8

#### 10.1.5 Conversions

A number can represent an angle expressed either in degrees or in radians. Two messages allow conversions of angles to be made:

```

degreesToRadians
the receiver being expressed in degrees, it returns a

```

number representing the receiver expressed in radians,  
radiansToDegrees  
the receiver being expressed in radians, it returns a  
number representing the receiver expressed in degrees.

#### 10.1.6 Nature of the object returned by arithmetic operations

Arithmetic operations can be broken down into two categories.

- Those in which the argument and the receiver are instances of the same class, for example:

2 + 3

- Those in which the argument and the receiver are not instances of the same class, for example:

2.2 + 3

When the receiver and the argument are instances of the same class, one is able to determine the class of the result. Depending on the values represented by the argument and the receiver, the result returned may belong to different classes.

For example:

8/4 returns 2, an instance of class Integer

8/3 returns (8/3), an instance of class Fraction.

When the receiver and the argument are not instances of the same class, one of the operands must be transformed to get back to the preceding case. The transformation must be made with the minimum loss of data. The answer is therefore to transform the item that contains the minimum of data.

Thus a hierarchy of numbers can be defined depending on their generality. This is as follows:

Float

Fraction

LargePositiveInteger LargeNegativeInteger

SmallInteger

All numbers can be represented by instances of the class Float (see below), but the precision is not infinite.

There are three messages for making conversions:

coerce: aNumber

returns a number representing the number aNumber that is an instance of the same class as the receiver. This method must be defined in each subclass of Number.

generality

returns the degree of generality of the receiver,

retry: aSelector coercing: aNumber

the receiver and the argument not being instances of the same class, the least general must be converted, then the selector message aSelector must be re-evaluated. For example, in the expression

3.7\*2

the argument 2 and the receiver 3.7 are not instances of the same class. Evaluation of the expression will then look like

```
3.7 retry #* coercing: 2  
which will transform 2 into 2.0 and will return the  
result 7.4.
```

### 10.2 The class **Float**

The instances of class **Float** represent real numbers with a finite precision and lying between two limiting values. For example, the precision may be to 6 figures, with the numbers lying between  $10^{-32}$  and  $10^{32}$ .

Some examples of floating number are:

```
1.2 -1.2 2.1e4 -2.1e4
```

Whenever one of the operands in an arithmetic expression is a **Float**, the result is a **Float**.

### 10.3 The class **Fraction**

The instances of the class **Fraction** represent rational numbers. This representation is exact; a fraction has an integer numerator and denominator. The instances of the class **Fraction** can be obtained in an arithmetic expression when one of the operands is a fraction and the other is not a floating number (cf 10.1.6 and 10.2). Instances of class **Fraction** can also be obtained by dividing two integers, when the numerator is not an exact multiple of the denominator.

### 10.4 The class **Integer**

The instances of class **Integer** represent integer numbers. This class has three subclasses which are:

**SmallInteger**

which represents integers whose absolute value is less than a limit which allows for compact storage (typically in one word less 1 bit),

**LargePositiveInteger**

which represents arbitrarily large positive integers,

**LargeNegativeInteger**

which represents arbitrarily large negative integers,

Calculations are much more costly when they involve instances of the classes **LargePositiveInteger** and **LargeNegativeInteger**.

If an operation on two instances of the class **SmallInteger** gives a result outside the limits of the class, the result is transformed into a **LargePositiveInteger** or a **LargeNegativeInteger**.

#### 10.4.1 Enumeration

There is a method for evaluating a block several times. This is

```
timesRepeat: aBlock
```

which evaluates the block aBlock a number of times equal to the receiver. For example:

a <- 1.

5 timesRepeat: [a <- a/20 + a reciprocal]

assigns to a the value 1 and then evaluates the block 5 times. The final value of a is 1.41421 (close to the square root of 2).

#### 10.4.2 Some additional arithmetic functions

The class Integer contains some arithmetic functions that belong to integers. The messages that relate to these functions are:

factorial

returns the factorial of the receiver, which must not be negative,

gcd: anInteger

returns the greatest common divisor of the receiver and the argument anInteger,

lcm: anInteger

returns the largest common multiple of the receiver and the argument anInteger.

#### 10.4.3 Bit manipulation

An integer can be interpreted as a sequence of bits (the representation of the integer expressed to the base 2). Some messages allow access to these bits. These are:

allMask: anInteger

returns true if all the bits set in the argument anInteger are also set in the receiver,

anyMask: anInteger

returns true if at least one of the bits set in the argument anInteger is also set in the receiver,

noMask: anInteger

returns true if none of the bits set in the argument anInteger is set in the receiver,

bitAnd: anInteger

returns an integer whose bits are obtained by making a logical "and" between the corresponding bits of the argument anInteger and the receiver,

bitOr: anInteger

returns an integer whose bits are obtained by a logical "or" between the corresponding bits of the argument anInteger and the receiver,

bitXor: anInteger

returns an integer whose bits are obtained by making an "exclusive or" between the corresponding bits of the argument anInteger and the receiver,

bitAt: anIndex

returns the receiver bit positioned at anIndex,

bitInvert

returns an integer whose bits are complements of the

```

receiver bits,
    highBit
returns the position of the most significant bit of the
receiver or the number of bits necessary to represent
the receiver to base 2,
    bitShift: anInteger
returns an integer (in two's complement) obtained by
shifting the receiver bits (in two's complement) of
anInteger to the left.
For example:

```

Expressions	Results
2r1010	10
2r1000	8
10 allMask: 8	true
8 allMask: 10	false
8 anyMask: 10	true
10 noMask: 8	false
10 bitAnd: 8	8
10 bitOr: 8	10
10 bitXor: 8	2
10 bitAt: 3	0
10 bitInvert	-11
10 highBit	4
10 bitShift: 2	40

#### 10.4.4 Changing base

We have seen how a number could be expressed to another base. There is a message to do this:

```

radix: anInteger
returns a representation to base anInteger of the receiver.
For example:
2 radix: 2           returns 2r10
9 radix: 3           returns 3r100
4r10 radix: 2        returns 2r100

```

# 11 The Class Collection

A collection represents a set of objects that are not necessarily of the same kind.

Most of the subclasses of **Collection** use indexed instance variables. Indexed instance variables are accessed by an index.

The class **Collection** is an abstract class that describes the behaviour of all collections. It supports four categories of message described below.

## 11.1 Creation of instances

Some collections can be expressed in literal form (see chapter 5). For example

#(1 5 3 4)

represents an instance of **Array** with four elements,

'this is a string'

represents an instance of **String** with 16 elements.

The usual messages for creating an instance of a Smalltalk-80 class are **new** and **new:**. These messages are also usable for instances of the class **Collection**. However, there are additional messages allowing the creation of new collections in which the elements are initialised on creation. These messages are:

with: anObject

returns a new instance of the receiver containing the single element **anObject**,

with: aFirstObject with: aSecondObject

returns a new instance of the receiver containing the two elements **aFirstObject** and **aSecondObject**,

with: aFirstObject with: aSecondObject with: aThirdObject

returns a new instance of the receiver containing the three elements **aFirstObject**, **aSecondObject** and **aThirdObject**,

with: aFirstObject with: aSecondObject with: aThirdObject with: aFourthObject

returns a new instance of the receiver containing the four elements **aFirstObject**, **aSecondObject**, **aThirdObject** and **aFourthObject**.

For example:

Array with: 1 with: 5 with: 3 with: 4  
is equivalent to #(1 5 3 4).

String with: \$a with: \$b with: \$c  
is equivalent to 'abc'.

## 11.2 Manipulating the elements of a collection

### 11.2.1 Adding elements to a collection

There are two messages that allow elements to be added to a collection. They are:

add: anObject

adds to the receiver the element anObject,

addAll: aCollection

adds to the receiver all the elements of the collection aCollection.

For example:

#### Expressions

#### Contents of a

---

a <- Bag new.

Bag ()

a add: \$a.

Bag (\$a)

a addAll: #(2 3 \$b)

Bag (\$a 2 3 \$b)

### 11.2.2 Removing elements from a collection

Just as you can add elements to a collection, so you can remove them. There are three messages that allow elements to be removed from a collection. They are:

remove: anObject

removes one of the occurrences of the element anObject from the receiver and returns the object anObject. If no occurrence is found, an error is produced.

remove: anObject ifAbsent: aBlock

removes one of the occurrences of the element anObject from the receiver and returns the object anObject. If no occurrence is found, the block aBlock is evaluated.

removeAll: aCollection

removes an occurrence of each element of the collection aCollection. If, for an element of the collection aCollection, there is no occurrence in the receiver, an error is generated. If not, the collection aCollection is returned.

For example:

#### Expressions

#### Contents of a

---

a

Bag (\$a 2 3 \$b)

a remove: \$a

Bag (2 3 \$b)

a remove: \$a ifAbsent: [a add: \$b]

Bag (\$b 2 3 \$b)

a removeAll: #(2 \$b)

Bag (3 \$b)

### 11.3 Tests on a collection

Three messages allow collections to be tested. These are:

includes: anObject

returns true if the object anObject is equal to one of the elements of the receiver,

isEmpty

returns true if the receiver is an empty collection,

occurrencesOf: anObject

returns the number of elements of the receiver equal to the argument anObject.

For example:

Expressions	Result
a	Bag (\$a 2 3 \$b \$b)
a includes: 2	true
a isEmpty	false
a occurrencesOf: \$b	2

### 11.4 Enumeration of a collection

A collection generally includes several elements. There are several messages described in the class Collection that allow each element of a collection to be accessed and allow an operation to be performed on this element. These messages are:

do: aBlockWithOneArgument

evaluates the block aBlock for each element of the receiver, the element being the argument of the block,

select: aBlockWithOneArgument

evaluates the block aBlock for each element of the receiver and returns a collection of the same class as the receiver, containing those elements for which evaluation of the block has returned true,

reject:aBlockWithOneArgument

evaluates the block aBlock for each element of the receiver and returns a collection of the same class as the receiver, containing those elements for which evaluation of the block has returned false,

detect: aBlockWithOneArgument

evaluates the block aBlock for each element of the receiver and returns the first element of the receiver that evaluation of the block has returned as true. If no element is found, an error is produced.

detect: aBlock ifNone: anOtherBlock

evaluates the block aBlock for each element of the receiver and returns the first element of the receiver that evaluation of the block aBlock has returned as true. If no element is found, the block anOtherBlock is evaluated.

inject: anInitialValue into: aBlockWithTwoArguments evaluates the block aBlockWithTwoArguments for each element of the receiver. The second argument of the block is the element, the first argument of the block is the result of its preceding evaluation. At first evaluation, the first argument is anInitialValue.

For example:

```

sum <- 0
a <- #(1 2 3 4 5).
a do: [:anElement | sum <- sum +anElement].
places in the variable sum the sum of all the elements
in the array a;
a <- #(1 2 3 4 5).
a select: [:anElement | anElement even].
returns Array (2 4);
a reject: [:anElement | anElement even]
returns Array (1 3 5);
a detect: [:anElement | anElement > 2.5]
returns 3;
a detect: [:anElement | anElement > 5] ifNone:
[ 'Element not found']
returns 'Element not found';
a inject: 10 into: [:initialValue: anElement |
initialValue + anElement]
returns 25. On the first evaluation of the block,
initialValue is 10, anElement is 1; the result of the
evaluation of the block, which is also the result of the
evaluation of its last expression, is 11. On the second
evaluation, initialValue is 11, anElement is 2, and so
on.

```

The message inject: into: allows a local variable to be initialised in the message that will be passed between the successive evaluations of a block. To write the previous example without inject: into:, one could write, for example:

```

counter <- 10.
a do: [:anElement | counter <- counter + anElement]

```

### **11.5 Conversion of a collection**

Five messages understood by the instances of class Collection transform instances of a subclass of class Collection to instances of another subclass of class Collection.

```

asBag
returns an instance of class Bag whose elements are
those of the receiver,
asSet
returns an instance of class Set whose elements are
those of the receiver (an instance of class Set does not
possess two equal elements),
asOrderedCollection

```

returns an instance of class OrderedCollection whose elements are those of the receiver,  
asSortedCollection  
returns an instance of class SortedCollection whose elements are those of the receiver (the elements are stored in ascending order),  
asSortedCollection: aBlock  
returns an instance of class SortedCollection whose elements are those of the receiver (the criterion for the order of the elements being given by the evaluation of the block aBlock).

## 11.6 The subclasses of the class Collection

Class Collection has no instances; that is, it is an abstract class. It possesses several subclasses that allow the different structures of useful data to be represented. These subclasses are arranged in the following hierarchy.

### 11.6.1 The class Bag

Class Bag is a subclass of class Collection. The elements of an instance of Bag represent an arbitrary set whose elements are not accessible by keys. Because of this, the messages at: and at: put: are not understood by the instances of class Bag. In fact, an instance of the class Bag just represents a group of elements. There is only one additional message that is understandable by the instances of this class. It is as follows:

add: anObject withOccurrences: anInteger  
adds to the receiver a number of occurrences of anObject equal to anInteger.

### 11.6.2 The class Set

Class Set is also a subclass of class Collection. In contrast to the instances of class Bag, those of class Set cannot contain two equal elements.

### 11.6.3 The class Dictionary

Class Dictionary is a subclass of class Set. The elements of an instance of class Dictionary are instances of class Association. An instance of class Association is represented by a pair of objects; the first element is called the key, the second element is called the value assigned to the key.

The messages at: and at: put: are still defined for the instances of this class; however, their methods have been redefined. In contrast to the foregoing, the argument of the keyword at: is no longer a numeric index but any object that represents a key.

The additional messages that can be understood by the instances of class Dictionary are as follows.

**Access**

at: aKey ifAbsent: aBlock  
 returns the value assigned to the key aKey. If the key is not found, the block aBlock is evaluated.

associationAt: aKey  
 returns the association whose key is aKey. If the key is not found, an error is produced.

associationAt: aKey ifAbsent: aBlock  
 returns the association whose key is aKey. If the key is not found, the block aBlock is evaluated.

keyAtValue: aValue  
 returns a key that refers to the value aValue, and if several keys reference the same value, returns the first one found. If no key is found, nil is returned.

keys  
 returns an instance of class Set containing all the keys of the receiver,

values  
 returns an instance of class Bag containing all the values of the receiver.

For example:  
 dictionary <- Dictionary new.  
 dictionary at: #un put: #one.  
 dictionary at: #deux put: #two.  
 dictionary at: #trois put: #three.  
 dictionary add: (Association key: #quatre value: #four).

allows a dictionary to be created. We have used two messages to add an element to the dictionary, the message at:put: and the message add: whose argument must be an Association.

**Expressions****Result**


---

dictionary at: #deux	two
dictionary associationAt: #deux	deux -> two
dictionary keyAtValue: #three	trois
dictionary keys	Set (un deux trois quatre)
dictionary values	Bag (one two three four)

---

**Test**

Being a subclass of Collection, the class Dictionary inherits the test messages of this class, in particular, includes:, although the method assigned to this message is redefined in the class Dictionary. This message allows for testing the inclusion or otherwise of a value of an association. The same goes for the message occurrencesOf:. The new messages are:

includesAssociation: anAssociation

returns true if the receiver contains the element anAssociation,

includesKey: aKey

returns true if one of the elements of the receiver has aKey as its key.

For example:

Expressions	Result
dictionary includes: #two	true
dictionary includes: #deux	false
dictionary includesAssociation: (Association key: #deux value: #two)	true
dictionary includesKey: #two	false
dictionary includesKey: #deux	true

### Deletion

Three new messages allow an element to be deleted from a dictionary. The message remove: defined in Collection is unusable for a dictionary. So as to be able not to use it, the method assigned to it is redefined in Dictionary; it is written as follows:

remove: anObject

self shouldNotImplement

The new messages are:

removeAssociation: anAssociation

which deletes the element anAssociation from the receiver. If the element anAssociation is not found in the receiver, an error is produced.

removeKey: aKey

deletes the element of the receiver whose key is aKey. If the key is not found, an error is produced; if successful, the value assigned to the key is returned.

For example:

Expressions	Dictionary
dictionary removeAssociation: (Association key: #un value: #one)	Dictionary (deux -> two quatre -> four trois -> three)
dictionary removeKey: #deux	Dictionary (trois -> three quatre -> four)

### Enumeration

Being a subclass of class Collection, class Dictionary inherits the messages of this class, in particular, do:, even though the method assigned to this class is redefined in class Dictionary. This message allows evaluation of a block for all values of the receiver and not the associations. The new messages are:

associationsDo: aBlock

evaluates the block aBlock for all associations of the receiver.

keysDo: aBlock

evaluates the block aBlock for all keys of the receiver.

For example:

```
dictionary associationDo: [:anAssociation | anAssociation
    value:anAssociation value asString]
```

means that the values of the dictionary are no longer symbols but strings. The result for the dictionary is:  
 Dictionary (one -> 'un' two -> 'deux' three -> 'trois'  
 four -> 'quatre').

#### **11.6.4 The class IdentityDictionary**

The class **IdentityDictionary** is a subclass of class **Dictionary**. The difference between these two classes is that for class **Dictionary** the search for a key uses the equality (=), whereas for class **IdentityDictionary** the search for a key uses the equivalence (==).

#### **11.6.5 The class SequenceableCollection**

The class **SequenceableCollection** is a subclass of class **Collection**. The instances of **SequenceableCollection** have ordered elements that are accessible via a numeric index. An instance of **SequenceableCollection** possesses a first element and a last element.

##### **Access**

Being a subclass of **Collection**, the class **SequenceableCollection** inherits from the access messages of this class, especially **at:** and **at:put:**. The other messages are:

**atAll: aCollection put: anObject**

assigns the argument **anObject** to each index of the receiver that belongs to the argument **aCollection**,

**atAllPut: anObject**

assigns the argument **anObject** to all indices of the receiver,

**first**

returns the first element of the receiver, producing an error if the receiver is empty,

**last**

returns the last element of the receiver, producing an error if the receiver is empty,

**indexOf: anElement**

returns the first index of the element **anElement**, returning 0 if the receiver does not contain the element **anElement**,

**indexOf: anElement ifAbsent: aBlock**

returns the first index of the element **anElement**, and

evaluates the block aBlock if the receiver does not contain the element anElement,

```
indexOfSubCollection: aSubCollection
startingAt: anIndex
```

if the argument aSubCollection is a subcollection of the receiver, returns the index in the receiver of the first element of the subcollection, the search starting at the index anIndex, and returns 0 if the subcollection is not found,

```
indexOfSubCollection: aSubCollection startingAt:
anIndex ifAbsent: aBlock
```

if the argument aSubCollection is a subcollection of the receiver, returns the index in the receiver of the first element of the subcollection, the search starting at the index anIndex, and evaluates the block aBlock if the subcollection is not found,

replaceFrom: beginning to: end with: aCollection replaces the elements whose index lies between beginning and end with the elements of the argument aCollection; if the size of the argument aCollection does not match with beginning and end, an error is produced,

```
replaceFrom: beginning to: end with: aCollection
startingAt: indexBeginning
```

replaces the elements whose index lies between beginning and end with the elements of the argument aCollection, starting at the index indexBeginning; the size of aCollection must not be less than (end - beginning + indexBeginning). For example:

### Expressions

### Result

#(\$a \$b \$c \$d) size	4
#(\$a \$b \$c \$d) at: 3	\$c
#(\$a \$b \$c \$d) at: 3 put: \$e	#(\$a \$b \$e \$d)
#(\$a \$b \$c \$d) atAll: #(1 2) put: \$f	#(\$f \$f \$c \$d)
#(\$a \$b \$c \$d) atAllPut: \$a	#(\$a \$a \$a \$a)
#(\$a \$b \$c \$d) first	\$a
#(\$a \$b \$c \$d) last	\$d
#(\$a \$b \$c \$d) indexOf: \$b	2
#(\$a \$b \$c \$d) indexOf: \$t ifAbsent: [5]	5
'abcdefghijkl' indexOfSubCollection:	
'cd' startingAt: 1	3
'abcdefghijkl' indexOfSubCollection:	
'bc' startingAt: 3	5
'abcdefghijkl' replaceFrom: 5 to: 7 with 'FGH'	'abcdFGHL'
'abcdefghijkl' replaceFrom: 5 to: 8 with:	
'ABCDEFGHIJK' startingAt: 3	'abcdCDEF'

### Copy

There are messages that allow all or part of an instance of SequenceableCollection to be copied. They are:

```
, aCollection
```

returns a copy of the receiver concatenated with the argument aCollection which must be an instance of SequenceableCollection,

copyFrom: beginning to: end

returns a copy of the subcollection of the receiver starting at the index beginning and finishing at the index end,

copyReplaceAll: aSubCollection with: anOtherSubCollection

returns a copy of the receiver in which all the occurrences of the subcollection aSubCollection are replaced with the collection anOtherSubCollection,

copyReplaceFrom: beginning to: end with: aCollection if end is greater than beginning, replaces the subcollection starting at the index beginning and finishing at the index end with the collection aCollection; if end is less than beginning, inserts the collection aCollection in front of the index beginning and returns the copy of the new collection,

copyWith: anObject

returns a copy of the receiver concatenated with anObject,

copyWithout: anObject

returns a copy of the receiver from which all occurrences of the object anObject have been removed.

For example:

### Expressions

### Result

'abc', 'def'	'abcdef'
'abcdef' copyFrom: 2 to: 5	'bcde'
'abcdef' copyReplaceAll: 'cd' with: 'CDE'	'abCDef'
'abcdef' copyReplaceFrom: 2 to: 3 with: 'AAA'	'aAAAdef'
'abcdef' copyReplaceFrom: 3 to: 2 with: 'AAA'	'abAAACdef'
'abcdef' copyWith: \$g	'abcdefg'
'abcdef' copyWithout: \$f	'abcde'

### Enumeration

Because the elements of an instance of SequenceableCollection are ordered, enumeration of these elements is also made in an ordered manner. The new messages accessible by the instances of SequenceableCollection are:

findFirst: aBlock

which evaluates the block aBlock for each element of the receiver, and returns the index of the first element for which evaluation of the block returns true,

findLast: aBlock

evaluates the block aBlock for each element of the receiver, and returns the index of the last element for

which evaluation of the block returns true,  
 reverseDo: aBlock  
 evaluates the block aBlock for each element of the receiver starting with the last and going back to the first,  
 with: aCollection do: aBlock  
 evaluates the two argument block aBlock, taking as first argument an element of the receiver and as second argument an element of the SequenceableCollection aCollection. The SequenceableCollection aCollection must be the same size as the receiver.

For example:

Expressions	Result
'abcDEFG' findFirst: [:i   i isUpperCase]	4
'abcDEFG' findLast: [:i   i isUpperCase]	7
mult <- Bag new.	
#(1 2 3 4) with: #(5 6 7 8)	
do: [:i :j   mult add: i*j].	
mult	Bag(5 12 21 32)

SequenceableCollection possesses several subclasses; these are OrderedCollection, LinkedList, Interval and ArrayedCollection.

#### 11.6.5.1 The class OrderedCollection

The class **OrderedCollection** is a subclass of SequenceableCollection. The order of the elements of an instance of OrderedCollection is fixed by the order in which they are added. The new messages accessible by the instances of OrderedCollection are:

addLast: anObject

which adds the object anObject to the end of the receiver,

addFirst: anObject

adds the object anObject to the beginning of the receiver,

add: anObject after: anOtherObject

adds the object anObject to the receiver after the first occurrence of anOtherObject. If the object anOtherObject is not an element of the receiver, an error message is produced,

add: anObject before: anOtherObject

adds the object anObject to the receiver in front of the first occurrence of anOtherObject. If the object anOtherObject is not an element of the receiver, an error message is produced,

addAllFirst: aCollection

adds to the beginning of the receiver all the elements of the OrderedCollection aCollection,

addAllLast: aCollection

adds to the end of the receiver all the elements of the OrderedCollection aCollection,  
 removeFirst  
 removes the first element of the receiver and returns this element,  
 removeLast  
 removes the last element of the receiver and returns this element. An error is produced if the receiver is empty,  
 after: anObject  
 returns the element that appears after the first occurrence of the object anObject. If the receiver contains no object anObject or there is no element after anObject, an error is produced,  
 before: anObject  
 returns the element that appears in front of the first occurrence of the object anObject. If the receiver contains no object aObject or there is no element in front of anObject, an error is produced.

For example:

Expressions	List
list <- OrderedCollection new.	OrderedCollection ()
list add: 1	OrderedCollection (1)
list addFirst: 2	OrderedCollection (2 1)
list addLast: 3	OrderedCollection (2 1 3)
list add: 4 after: 1	OrderedCollection (2 1 4 3)
list add: 5 before: 1	OrderedCollection (2 5 1 4 3)
list removeFirst	OrderedCollection (5 1 4 3)
list removeLast	OrderedCollection (5 1 4)
list before: 1	returns 5
list after: 1	returns 4
list addAllFirst: list	OrderedCollection (5 1 4 5 1 4)
list addAllLast: list	OrderedCollection (5 1 4 5 1 4 5 1 4 5 1 4)

### 11.6.5.2 The class **SortedCollection**

The class **SortedCollection** is a subclass of **OrderedCollection** whose elements are ordered. To order these elements the class uses a function represented by a two argument block. To create an instance of **SortedCollection**, it is sufficient to send the message **sortBlock:** to the class. The argument of **sortedBlock:** is a two argu-

ment block, for example:

```
SortedCollection sortBlock: [:arg1 :arg2 | arg1 > arg2]
```

This instance of SortedCollection will contain a collection of objects stored from largest to smallest.

The message new allows creation of a SortedCollection with the following as default sort block

```
[:arg1 :arg2 | arg1 <= arg2]
```

There are two other messages that allow any collection to be converted into a SortedCollection. These are:

```
asSortedCollection
```

which transforms the receiver into a SortedCollection with the block as default sort function,

```
asSortedCollection: aBlock
```

which transforms the receiver into a SortedCollection with the block given as argument being the sort function.

There are two messages sortBlock and sortBlock: that are understood by instances of SortedCollection; the first returns the sort block of the receiver, the second allows the sort block of the receiver to be altered.

For example:

Expressions	Results
collection <- SortedCollection new	SortedCollection ()
collection add: 'abcdefg'	'abcdefg'
collection add: 'abdef'	'abdef'
collection add: 'aa'	'aa'
collection	SortedCollection ( 'aa' ''abcdefg' ''abdef' )
collection sortBlock: [:a :b   a size > b size]	SortedCollection ( 'abcdefg' ''abdef' ''aa' )
collection add: 'abc'	'abc'
collection	SortedCollection ( 'abcdefg' ''abdef' ''abc' ''aa' )

The only constraint imposed on the sort block is that it returns true or false. In particular, it is important to ensure that the objects of the SortedCollection can be compared by the sort block.

#### 11.6.5.3 The class LinkedList

The class **LinkedList** is a subclass of SequenceableCollection whose instances contain a list of instances of the class **Link** ordered explicitly at the time of addition or removal. The class **Link** is a subclass of **Object** whose instances make reference to another instance of **Link**.

To create an instance of the class `Link`, the message `nextLink: aLink` needs to be sent to the class. The argument of the message `nextLink:` must be an instance of `Link`.

There are two messages that are understood by instances of `Link`:

`nextLink`

which returns the instance of `Link` referenced by the receiver; and

`nextLink: aLink`

which allows the reference of the receiver to be altered.

The class `Link` has no direct instances because they would not be usable (that is, the class `Link` is an abstract class). To use the class `LinkedList`, subclasses of `Link` have to be defined whose instances will have other usable variables.

The instances of `LinkedList` understand the same messages as the instances of `Link`; they also understand all the messages of the class `SequenceableCollection` (`addFirst`, `removeFirst`, etc.). The following example, which allows a circular list of `n` objects to be created, shows how the class `LinkedList` may be used:

name of class	Bond
superclass	Link
name of instance variables	object
instance methods	
<b>object</b>	
↑object	
<b>object: anObject</b>	
object <- anObject	
name of class	CircularList
superclass	LinkedList
class method	
<b>new: ASize</b>	
aList	
aList <- super new.	
(1 to: aSize) inject: nil into:	
[:i :j  aList addFirst: (Bond nextLink: i)].	
aList last nextLink: aList first	

#### 11.6.5.4 The class `Interval`

The class `Interval` is a subclass of `SequenceableCollection` whose instances represent arithmetic sequences. An arithmetic sequence is characterised by its first element and the increment which steps from one element to the next up to a finite limit. There are two messages that create `Interval` instances.

`from: beginning to: end`

returns an instance of the class Interval starting at the beginning number and finishing at the end number, with increments of 1.

from: beginning to: end by: anIncrement  
returns an instance of the class Interval starting at the beginning number and finishing at the end number and incrementing by anIncrement.

The class Number also supports two messages that create Interval instances:

to: aNumber

which returns an instance of the class Interval starting at the receiver and finishing at the number aNumber, with an increment of 1.

to: aNumber by: anIncrement

which returns an instance of the class Interval starting at the receiver and finishing at the number aNumber, with an increment of anIncrement.

For example:

1 to: 5 by: 2

returns an instance of Interval whose elements are 1, 3 and 5.

1 to: 5 by: 2.1

returns an instance of Interval whose elements are 1 and 3.1.

#### 11.6.5.5 The class **ArrayedCollection**

The class **ArrayedCollection** is a subclass of SequenceableCollection whose instances are arrays of objects. Each object is accessible by a numeric key representing its position in the array. The class **ArrayedCollection**, which is an abstract class, has five subclasses, as described below.

##### - The class **Array**

An instance of the class **Array** can store any type of object. The messages understood by instances of **Array** are the messages supported by the class **Collection**.

##### - The class **String**

In contrast to class **Array** the instances of class **String** only contain characters. The instances of **String** are therefore strings of characters. The additional messages supported by the class **String** are:

#### **Creation**

Two messages allow instances of **String** to be created:

fromString: aString

returns a copy of the string aString;

readFrom: aStreamOfCharacters

returns an instance of **String** created from the characters contained in the stream of characters aStreamOfCharacters (an instance of **Stream**, see chapter 12).

**Comparison**

Six messages allow strings of characters to be compared with one another. The distinction between upper and lower case is ignored. These messages are:

```

<aString
returns true if the receiver precedes the argument
aString alphabetically,
<=aString
returns true if the receiver precedes the argument
aString alphabetically or if two strings are equal,
>aString
returns true if the receiver follows the argument
aString alphabetically,
>=aString
returns true if the receiver follows the argument
aString alphabetically or if the two strings are equal,
match: aString
returns true if the argument aString appears in the receiver. The characters # and * have a special meaning in the receiver:
- # can represent any character
- * can represent any string of characters (including the empty string).
sameAs: aString
returns true if the receiver and the argument aString are equal.
```

**Conversion**

Three messages allow character strings to be converted:

```

asLowerCase
returns a string of characters equivalent to the receiver containing only lower case characters,
asUpperCase
returns a string of characters equivalent to the receiver containing only upper case characters,
asSymbol
returns the symbol made up of characters cf the receiver.
```

For example:

Expressions	Result
'String' = 'String'	true
'String' = 'string'	false
'String' sameAs: 'String'	true
'String' sameAs: 'string'	true
'#tring' match: 'string'	true
'*ng' match: 'string'	true
'string' asUpperCase	'STRING'
'string' asSymbol	string

**The class Symbol**

The class **Symbol** is a subclass of the class **String** whose instances are used to name objects in the system. The objects that have to be named are for example classes, messages, variables, etc. In the interests of saving space, symbols are only represented once in the system. Two messages allow new symbols to be created:

intern: aString  
returns the unique symbol made up of the characters of the string aString;

internCharacter: aCharacter  
returns the unique symbol made up of the character aCharacter.

Remember that the syntax for writing a symbol as a literal form is #nameOfSymbol.

**The class Text**

The class **Text** is a subclass of **ArrayedCollection**. The class **Text** allows strings of characters to be stored as **String**, with the additional possibility of assigning to each character a further characteristic that can for example represent the style of the character when displayed.

**The class ByteArray**

The class **ByteArray** is a subclass of the class **ArrayedCollection** whose instances contain integers coded into a byte (between 0 and 255). This class allows arrays of such integers to be stored efficiently.

**The class RunArray**

The class **RunArray** is a subclass of the class **ArrayedCollection**. This class allows the convenient storage of arrays containing sequences of similar objects. The instances of **RunArray** have two named instance variables that are runs and values. The variable runs contains a list of integers and the variable values contains a list of objects. Each object of values is assigned to an integer of runs that represents the number of consecutive occurrences of this object.

## 12 The Class Stream

Collections allow objects to be grouped. Each object of a collection can be accessed individually; also, an operation can be carried out on all the objects of a collection, provided that this is done in a loop. However, it is not easy to mix the two types of access and operate on the enumeration of objects.

The class **Stream** is a subclass of the class **Object** that allows streams of objects to be controlled. The instances of **Stream** have a variable representing the current object in the objects stream. The instances of **Stream** do not directly store objects, but they do control access to these objects.

Several types of access are possible; they will be defined in the subclasses of **Stream**, which accounts for why **Stream** is an abstract class.

To summarise, the instances of **Stream** allow access to the collections to be controlled.

The messages of class **Stream** can be classified into several categories, as follows.

### **Reading**

Four access messages are understood by the instances of **Stream**:

```
next  
returns the next object accessed by the receiver,  
next: anInteger  
returns the next anInteger objects accessed by the receiver, in the collection form of the same class as the collection accessible by the receiver,  
nextMatchFor: anObject  
accesses the next object and returns true if this object is equal to the argument anObject,  
contents  
returns all the objects accessed by the receiver.
```

### **Writing**

Three write messages are understood by the instances of **Stream**:

```
nextPut: anObject  
stores the argument anObject at the next position accessed by the receiver and returns anObject,  
nextPutAll: aCollection
```

stores all the objects of the argument aCollection at the next positions accessed by the receiver and returns aCollection,

next: anInteger put: anObject

stores the argument anObject at the next anInteger positions accessed by the receiver and returns anObject.

### Test

The message atEnd establishes whether the receiver can access any more objects.

### Enumeration

The message do: aBlock evaluates the argument aBlock for all the objects accessible by the receiver.

## 12.1 The class PositionableStream

The class **PositionableStream** is a subclass of **Stream**. It is again an abstract class that has three subclasses:

- the class **ReadStream**
- the class **WriteStream**
- the class **ReadWriteStream**

In addition to the protocol defined within class **Stream**, instances of the class **PositionableStream** allow positioning within the object stream; these streams of objects therefore must themselves be ordered, that is, they must be instances of the class **SequenceableCollection**.

To create an instance of **PositionableStream**, two messages are possible:

on: aCollection

returns an instance of **PositionableStream** allowing control of the stream made up of the objects of the argument aCollection.

on: aCollection from: indexBeginning to: indexEnd  
returns an instance of **PositionableStream** allowing control of the stream made up of the the objects of the argument aCollection that lie between the arguments indexBeginning and indexEnd.

The additional messages understood by the instances of **PositionableStream** fall into several categories, as follows.

### Test

The message isEmpty returns true if the collection controlled by the receiver is empty.

### Access

Four messages allow access to the objects of the collection controlled by the receiver:

peek

returns the next object of the collection, but does not

alter the current position in the stream; it returns nil if the receiver is at the end,

```
    peekFor: anObject
```

returns true and advances one position in the stream if the next object is equal to the argument anObject; if the object is not found, it returns the remainder of the collection,

```
    reverseContents
```

returns a collection in inverse order of that controlled by the receiver.

### **Positioning**

The messages that alter or give the position in the objects stream are:

```
    position
```

returns the integer representing the current position in the stream,

```
    position: anInteger
```

places the receiver at the position given by the argument anInteger; if the position given as argument is not within the limits of the collection controlled by the receiver, an error is produced,

```
    reset
```

places the receiver at the beginning of the collection that it controls,

```
    setToEnd
```

places the receiver at the end of the collection that it controls,

```
    skip: anInteger
```

shifts the current position of the argument anInteger, ensuring that this position never goes beyond the permitted bounds,

```
    skipTo: anObject
```

shifts the current position to place it immediately after the next occurrence of the argument anObject in the collection controlled by the receiver; it returns true if such an object exists and false otherwise.

For example:

### **Expression**

### **Result**

---

pS <- ReadStream on: # (peter paul john)	
pS reset	
pS next	peter
pS next	paul
pS next	john
pS next	nil
pS position	3
pS position: 2	2
pS peek	paul
pS next	paul
pS contents	(peter paul john)
pS reverseContents	(john paul peter)

### 12.1.1 The class ReadStream

The class **ReadStream** is a subclass of the class **PositionableStream**. The instances of this class can read the stream of objects of a collection but cannot alter this stream. None of the messages **nextPut:**, **next:** and **nextPutAll:** is recognised.

### 12.1.2 The class WriteStream

The class **WriteStream** is a subclass of the class **PositionableStream**. The instances can alter the objects stream of a collection but cannot read the objects of this stream. In the system this class is used on numerous occasions to implement the methods **printString** and **printOn** that allow objects to be displayed in a textual form.

Additional messages are supported by the class **WriteStream**. They fall into the following categories.

#### **Creation**

The messages that allow instances to be created are:

**with: aCollection**

creates an instance of **WriteStream** allowing the collection given as argument to be controlled, and places it at the end of the stream (the next objects will be added at the end),

**with: aCollection from: indexBeginning to: indexEnd**  
creates an instance of **WriteStream** allowing control of the collection of objects lying between the arguments **indexBeginning** and **indexEnd** in the collection given as argument, and places it at the end of the stream (the next objects will be added at the end).

#### **Writing characters**

The following methods allow a character to be added to an objects stream:

**cr**

adds a carriage-return after the current position of the receiver,

**tab**

adds a tabulation after the current position of the receiver,

**space**

adds a space after the current position of the receiver,

**crTab**

adds a carriage-return and a tabulation after the current position of the receiver,

**crTab: n**

adds a carriage-return and tabulations after the current position of the receiver.

An interesting case is where the collection is a string of characters. In such a case the control of the character stream amounts to carrying out editing funct-

ions on a string of characters.

### 12.1.3 The class `ReadWriteStream`

The class `ReadWriteStream` is a subclass of `WriteStream` which allows control of a stream of objects when reading or writing. It supports all the messages of `ReadStream` and `WriteStream`.

### 12.2 The class `ExternalStream`

The class `ExternalStream` is a subclass of the class `ReadWriteStream`. In the Smalltalk system, communication with the external files is also carried out by means of data streams. In contrast to the subclasses of `Stream` that we have already met, the input/output streams do not contain any objects but only bytes. The messages supported by the class `ExternalStream` fall into two categories.

#### Access

The access messages to byte streams are:

`nextNumber: n`

returns the next `n` bytes of the collection accessed by the receiver and interprets them as a positive instance of `SmallInteger` or of `LargePositiveInteger`, depending on the value read,

`nextNumber: n put: anInteger`

stores in `n` bytes the positive integer `anInteger` (an instance of `SmallInteger` or of `LargePositiveInteger`) in the collection accessed by the receiver,

`nextString`

returns in the form of a character string (instance of `String`) the next bytes of the collection accessed by the receiver,

`nextStringPut: aString`

stores the string of characters `aString` in the form of bytes in the collection accessed by the receiver,

`nextWord`

returns the next two bytes interpreted as an integer,

`nextWordPut: anInteger`

stores in two bytes the integer `anInteger` in the collection accessed by the receiver.

#### Positioning

The messages for positioning on the byte streams are:

`padTo: aSize`

places itself at the next multiple position of `aSize`, and returns the number of bytes skipped,

`padTo: aSize put: aCharacter`

places itself at the next multiple position of `aSize` writing the character `aCharacter` the required number of

```
times, and returns the number of bytes skipped,
    padToNextWord
places itself on the next word (first even byte position), and returns the character skipped if there is one,
    padToNextWordPut: aCharacter
places itself on the next word writing the character aCharacter as many times as necessary,
    skipWords: n
skips n words,
    wordPosition
returns the current position in words,
    wordPosition: aPosition
places itself at the position aPosition in words.
```

### 12.2.1 The class FileStream

The class **FileStream** is a subclass of the abstract class **ExternalStream**. All accesses to files are carried out by means of instances of **FileStream**. In contrast to its superclasses, the class **FileStream** does not work on collections present in central memory, but on streams of bytes present in auxiliary memories. In addition to the messages supported by the class **ExternalStream**, the class **FileStream** supports messages relating to the status of a file (open, closed, etc). Three class methods allow instances of **FileStream** to be created:

```
oldFileNamed: aString
creates an instance of FileStream accessing the file named aString and verifies that it exists,
newFileNamed: aString
creates an instance of FileStream accessing the file named aString, creating it if it does not exist and deleting it if it does,
fileNamed: aString
creates an instance of FileStream accessing the file named aString and creating it if it does not exist.
```

The two important new messages supported by **FileStream** are:

```
open
which allows the file accessed by the receiver to be opened; read and write operations cannot take place until this message has been sent.
close
which allows the file accessed by the receiver to be closed.
```

### 12.2.2 The class FileDirectory

The class **FileDirectory** is a subclass of the class **FileStream**. The class **FileStream** allows files to be named, while the class **FileDirectory** allows them to be located in the logical organisation of the secondary memory. The

methods of the class `FileDirectory` depend on the system support, that is, the file organisation of the operating system on which the Smalltalk interpreter has been written. The instances of the class `FileDirectory` represent directories or volumes.

# 13 Processes

A process represents a sequence of instructions that has to execute a particular task. A conventional machine, for example, contains processes for reading the keyboard, writing to disk, printing out text via a printer, etc. When a machine is what is known as a uniprocessor, it can only execute one process at a time; it is therefore necessary to be able to select a process to execute and synchronise the various processes required. Smalltalk-80 contains several classes that allow such concepts to be defined.

## 13.1 The class Process

A process is a sequence of expressions that are interpreted by Smalltalk. There are already in the system processes that inspect the input interfaces (keyboard and mouse) and a process that is activated when the amount of available memory is running low, notifying the user of a potential problem.

Processes are instances of the class **Process** which is a subclass of **Link**. To create a new process, four messages are supported by the class **BlockContext**:

**fork**

creates a new process from the receiver which is a block containing the expressions to be interpreted, and places it on the list of processes to be executed; the priority of the process created is equal to the priority of the process that creates it,

**forkAt: aPriority**

creates a new process from the receiver which is a block containing the expressions to be interpreted, and places it on the list of processes to be executed with the priority **aPriority**,

**newProcess**

creates a new process from the receiver which is a block containing the expressions to be interpreted, but does not place it on the list of processes to be executed,

**newProcessWith: anArray**

creates a new process from the receiver which is a block containing the expressions to be interpreted, without

putting it on the list of processes to be executed. The arguments of the block are initialised by the objects contained in the argument `anArray`.

The messages that allow action on processes are:

- `resume`
- which authorises execution of the receiver,
- `suspend`
- which halts executions of the receiver, but keeps it on the list of processes to be executed; to authorise continuation, the message `resume` must be sent,
- `terminate`
- halts execution of the receiver and removes it from the list of processes to be executed,
- `priority`
- returns the priority of the receiver,
- `priority: aPriority`
- alters the priority of the receiver and gives the value of the argument `aPriority`.

### 13.2 The class ProcessorScheduler

The Smalltalk virtual machine has only one processor; since as a consequence only one process can be executed at a time, the activation of the various processes must be controlled.

It is the role of the class `ProcessorScheduler`, a subclass of the class `Object`, to control a list of processes. The class `ProcessorScheduler` has only one instance which is stored in a global variable called `Processor`.

The messages supported by the class `ProcessorScheduler` are:

- `activePriority`
- returns the priority of the process currently being executed,
- `activeProcess`
- returns the process (instance of `Process`) currently being executed,
- `terminateActive`
- terminates the process currently being executed,
- `yield`
- authorises execution of a process of the same priority as the active process,
- `highIOPriority`
- returns the priority of the fast input/output processes,
- `lowIOPriority`
- returns the priority of the slow input/output processes,
- `systemBackgroundPriority`
- returns the priority of the Smalltalk processes operating in the "background",
- `timingPriority`

returns the priority of the processes that consult the real-time clock,

    userBackgroundPriority

returns the priority of the user processes operating in the "background",

    userInterruptPriority

returns the priority of the user processes considered to be of high priority,

    userSchedulingPriority

returns the priority of the non-priority user processes.

The list below gives the rank order of priorities, from the highest to the lowest:

timingPriority	-> Priority 8
----------------	---------------

highIOPriority	-> Priority 7
----------------	---------------

lowIOPriority	-> Priority 6
---------------	---------------

userInterruptPriority	-> Priority 5
-----------------------	---------------

userSchedulingPriority	-> Priority 4
------------------------	---------------

userBackgroundPriority	-> Priority 3
------------------------	---------------

systemBackgroundPriority	-> Priority 2
--------------------------	---------------

### 13.3 The class Delay

The class **Delay** is a subclass of the class **Object**. The instances of the class **Delay** allow the active process to be suspended for a determined period. There are three messages to create instances of **Delay**, as follow.

    forMilliseconds: aNumberOfMilliseconds

returns a new instance of the class **Delay** that will suspend execution of the active process for the number of milliseconds equal to the argument,

    forSeconds: aNumberOfSeconds

returns a new instance of the class **Delay** that will suspend execution of the active process for the number of seconds equal to the argument,

    untilMilliseconds: aNumberOfMilliseconds

returns a new instance of the class **Delay** that will suspend execution of the active process until the real-time clock reaches the value of the argument.

To suspend execution of the active process, the message **wait** must be sent to an instance of **Delay** that is initialised for the required delay.

The message **resumptionTime** sent to an instance of **Delay** returns the value in milliseconds of the real-time clock at which the interrupted process can resume.

For example:

```
[[true] whileTrue:
```

```
    [Time now printString asDisplayText displayAt:
```

```
        Sensor cursorPoint. (Delay forSeconds: 1) wait]] fork
```

will display the current time every second at the cursor position.

### 13.4 Semaphores

We have seen that processes are independent of one another. In some circumstances it is useful to be able to synchronise processes and communicate data between them.

The instances of the class **Semaphore**, which is a subclass of the class **LinkedList**, allow processes to be synchronised.

To allow a process to wait for an event from another process, it places itself in a wait state by sending a **wait** message to a semaphore. When the second process sends a signal to the first, the latter is able to resume.

The order in which these messages are sent has no effect on the execution of the processes, except that it may or may not cause the process placed in signal await state to wait.

The two messages that are involved are:

**wait**

which suspends the active process until the semaphore receives a signal,

**signal**

which allows a process to signal an event to the semaphore.

Several processes can wait for a semaphore; they resume execution in the same order that they were suspended as and when the semaphore receives the messages signalling an event.

One of the useful applications of semaphores is for achieving mutual exclusion between two processes. Two processes are in mutual exclusion when one of them cannot interrupt the other when it is in a critical section.

To define a mutual exclusion semaphore, a new semaphore must be created and the signal message sent to it. This is done by:

**forMutualExclusion**

A critical section is a part of a process that cannot be interrupted by another process. To define it, there must be a mutual exclusion semaphore to which the message **wait** is sent; that is, from that moment any process that sends the message **wait** to that same semaphore is blocked. When a critical section is terminated, it is then sufficient to signal to the semaphore to free access to the protected resource. The following message defines a critical section in a method:

**critical: aBlock**

which evaluates the argument **aBlock**, while preventing any other process from being activated.

**The class SharedQueue**

The instances of this class are queues that are used in Smalltalk to transmit data between processes while protecting their integrity. The instances of this class have two instance variables that refer to two semaphores. The first counts the elements in the queue and blocks the process that reads this queue when it is empty. The second is used for read and write operations to protect the queue (two processes cannot alter the queue at the same time, because of the danger of losing data).

The messages understood by a queue are:

isEmpty

returns true if the queue is empty, false if otherwise,  
next

returns the next object in the queue and positions itself at the following object; if there is no object in the queue, the process is suspended until an object is added,

peek

returns the next object in the queue without positioning itself at the following object; if there is no object in the queue, the process is suspended until an object is added,

nextPut: anObject

places the argument anObject at the next position in the queue.

# 14 The Classes Point and Rectangle

Just as we have numbers to represent scalar sizes, we also need objects to represent sizes in two dimensions. These sizes are essential in order to represent screen positions. The class **Point**, a subclass of the class **Object**, was introduced for this purpose.

## 14.1 The class Point

The instances of the class **Point** have two instance variables called **x** and **y**, which represent respectively the abscissa and ordinate of a point.

### 14.1.1 Creating a point

A class method creates new points:

**x: anAbscissa y: anOrdinate**

returns a new instance of the class **Point** whose abscissa is **anAbscissa** and ordinate **anOrdinate**.

There is another instance method of the class **Number** that allows new points to be created:

**@ aNumber**

returns a new instance of the class **Point** whose abscissa is the receiver and whose ordinate is **aNumber**.

### 14.1.2 Messages supported by the class Point

Like numbers (instances of the class **Number**), points are basic objects and therefore they understand a large number of messages. These can be divided into the following categories.

#### Access

The messages that allow access to the characteristics of a point are:

**x**

returns the abscissa of the receiver,

**y**

returns the ordinate of the receiver,

**x: aNumber**

replaces the abscissa of the receiver by the argument **aNumber**,

**y: aNumber**

replaces the ordinate of the receiver by the argument **aNumber**.

**Arithmetic**

The arithmetic operations are:

+ anOffset

returns a new point whose x-coordinate is the sum of the x-coordinates of the receiver and the argument anOffset; similarly for the y-coordinate. If the offset is not a point but a number, the two coordinates are incremented by the same value,

- anOffset

returns a new point obtained by subtracting from each coordinate of the receiver the corresponding coordinate of the argument anOffset; if the offset is not a point but a number, the two coordinates are decremented by the same value,

\* aCoefficient

returns a new point obtained by multiplying each coordinate of the receiver by the argument aCoefficient; if the coefficient is not a point but a number, the two coordinates are multiplied by the same value,

/ aCoefficient

returns a new point obtained by dividing the receiver by the argument aCoefficient; if the coefficient is not a point but a number, the two coordinates are divided by the same value,

// aCoefficient

returns a new point obtained by carrying out the integer division for each coordinate between the receiver and the argument aCoefficient; if the coefficient is not a point but a number, the two coordinates are divided by the same value,

abs

returns a new point obtained by taking the absolute value of each coordinate of the receiver,

rounded

returns a new point whose coordinates are the integers closest to each coordinate of the receiver,

truncateTo: aCoefficient

returns a new point whose coordinates are the largest multiples of the argument aCoefficient smaller than the coordinates of the receiver,

dist: aPoint

returns the Euclidean distance between the receiver and the argument aPoint,

grid: aPoint

returns the multiple of the argument aPoint closest to the receiver,

truncateGrid: aPoint

returns the largest multiple of the argument aPoint smaller than the receiver,

transpose

returns a new point whose x-coordinate is the receiver's y-coordinate, and whose y-coordinate is the receiver's x-coordinate.

### Comparison

The messages that allow two points to be compared are:

< aPoint

returns true if the two coordinates of the receiver are less than the two coordinates of the argument,

<= aPoint

returns true if the two coordinates of the receiver are less than or equal to the two coordinates of the argument,

> aPoint

returns true if the two coordinates of the receiver are greater than the two coordinates of the argument,

>= aPoint

returns true if the two coordinates of the receiver are greater than or equal to the two coordinates of the argument.

For example:

### Expressions


### Results

Point x: 10 y: 20	10@20
10@20	10@20
(10@20)x	10
(19@20)y	20
(10@20)x: 30	30@20
(10@20)y: 30	10@30
(10@20)+(5@8)	15@28
(10@20)+5	15@25
(10@20)-(5@8)	5@12
(19@20)-5	14@15
(10@20)*(2@3)	20@60
(10@20)*2	20@40
(10@20)/(2@4)	5@5
(10@20)/3	(10/3)@(20/3)
(10@20)//(3@7)	3@2
(10@20)//3	3@6
(-10@20) abs	10@20
(10.7@20.2) rounded	11@20
(10@20) truncateTo: 3	9@18
(4@4) dist: (1@0)	5
(10@20) grid: (3.3@7)	9.9@21
(10@20) truncatedGrid: (3.3@7)	9.9@14
(10@20) transpose	20@10

### 14.2 Coordinate system within Smalltalk environment

Each point on the screen can be defined by two coordinates, an abscissa and an ordinate. The diagram


below shows the position and orientation of the two coordinate axes in the Smalltalk system.



#### 14.3 The class Rectangle

A rectangle parallel to the axes can be defined by two points. In the Smalltalk system many objects use this concept of a rectangle (for example, windows, menus, etc). The class Rectangle, which is a subclass of the class Object, allows rectangles to be defined by two points, an origin and a corner.

The origin and the corner of a rectangle are defined as shown in the following diagram:



#### 14.3.1 Creating a rectangle

Three messages create new rectangles:

origin: anOriginPoint corner: aCornerPoint

returns a new instance of the class Rectangle whose origin is anOriginPoint and corner aCornerPoint,

origin: anOrigin extent: aVector

returns a new instance of the class Rectangle whose origin is anOriginPoint and whose corner is obtained by adding to the argument anOrigin the argument aVector (which is an instance of Point),

left: limitLeft right: limitRight top: limitTop

bottom: limitBottom

returns a new instance of the class Rectangle whose coordinates are delimited by the limits given as arguments.

There are two messages understood by instances of the class Point that also allow new rectangles to be created:

extent: aPoint

returns a new instance of the class Rectangle whose origin is given by the receiver and whose corner is obtained by adding the origin to the argument aPoint,

corner: aPoint

returns a new instance of the class Rectangle whose origin is given by the receiver and whose corner is the argument aPoint.

#### 14.3.2 Instance methods

The instance methods of the class Rectangle may be divided into several categories:

##### Access

There are nine methods for accessing particular points of a rectangle:

topLeft

returns the point located at the top left of the receiver,

topCenter

returns the point located at the top centre of the receiver,

topRight

returns the point located at the top right of the receiver,

leftCenter

returns the point located at the left centre of the receiver,


center

returns the point located at the centre of the receiver,

rightCenter

returns the point located at the right centre of the receiver,  
bottomLeft  
returns the point located at the bottom left of the receiver,  
bottomCenter  
returns the point located at the bottom centre of the receiver,  
bottomRight  
returns the point located at the bottom right of the receiver,  
origin  
returns the same point as the method topLeft,  
corner  
returns the same point as the method bottomRight.

These points are located as shown in the following diagram:



Other messages give access to the characteristics of the receiver:

origin: anOrigin corner: aCorner  
replaces the origin and the corner of the receiver by the arguments anOrigin and aCorner,  
origin: anOrigin extent: aSize  
replaces the origin by the argument anOrigin and the corner by the point obtained by adding the origin with the argument aSize,  
width  
returns the width of the receiver,  
height  
returns the height of the receiver,  
extent  
returns aPoint whose abscissa is the width and whose ordinate is the height,  
top

```
returns the ordinate of the origin of the receiver,  
    bottom  
returns the ordinate of the corner of the receiver,  
    left  
returns the abscissa of the origin,  
    right  
returns the abscissa of the corner,  
    area  
returns the area of the rectangle (the product of its  
height and its width).
```

#### **Operations on rectangles**

These operations return new objects, in contrast to those that we shall examine later, which transform the receiver. These operations are:

```
    amountToTranslateWithin: aRectangle  
returns a point corresponding to the minimal translation  
that is required to be made to the receiver for it to be  
placed within the argument aRectangle,  
    areasOutside: aRectangle  
returns a collection of rectangles contained in the rec-  
eiver and not intersecting the argument aRectangle,  
    expandBy: anIncrement  
returns a new rectangle obtained by increasing the rec-  
eiver by the value given as argument; the argument can  
be a number, a point or a rectangle,  
    insetBy: aDecrement  
returns a new rectangle obtained by decreasing the rec-  
eiver by the value given as argument; the argument can  
be a number, a point or a rectangle,  
    insetOriginBy: anOriginIncrement  
    cornerBy: aCornerDecrement  
returns a new rectangle whose origin is the origin of  
the receiver increased by the argument anOriginIncre-  
ment and whose corner is the corner of the receiver de-  
creased by the argument aCornerDecrement; the arguments  
can be points or numbers,  
    intersect: aRectangle  
returns the rectangle obtained by taking the intersec-  
tion of the receiver and the argument,  
    merge: aRectangle  
returns the smallest rectangle containing the receiver  
and the argument,  
    rounded  
returns a new rectangle whose coordinates are the  
rounded coordinates of the receiver,  
    translateBy: aTranslation  
returns a new rectangle whose coordinates are the co-  
ordinates of the receiver translated from the argument  
aTranslation; aTranslation can be a number or a point,  
    scaleBy: aCoefficient
```

returns a new rectangle whose coordinates are the coordinates of the receiver multiplied by the argument aCoefficient; aCoefficient can be a number or a point.

### Test

Three methods allow the characteristics of a rectangle to be tested:

contains: aRectangle

returns true if the argument aRectangle is within the receiver,

containsPoint: aPoint

returns true if the argument aPoint is within the receiver,

intersects: aRectangle

returns true if the intersection of the receiver and the argument aRectangle is non-empty.

### Transformations

The methods that alter rectangles are:


moveBy: aTranslation

translates the origin and the corner of the receiver by the argument aTranslation; aTranslation can be a number or a point,

moveTo: aPoint

translates the receiver in such a way that the origin of the receiver is located at the argument aPoint.

We shall take as an example three rectangles placed as in the following diagram:



Expression	Result
r1 <- Rectangle origin: (0@0) corner: (70@50)	
r2 <- Rectangle origin: (40@30) extent: (90@50)	
r3 <- 30@90 corner: (130@150)	
r1 centre	35@25
r2 width	90
r3 area	6000
r1 amountToTranslateWithin: r3	30@90
r1 areasOutside: r2	OrderedCollection ((0@0 corner: 70@30) (0@0 corner: 40@50))
r3 expandBy: 20	10@70 corner: 150@170
r3 expandBy: 20@10	10@80 corner: 150@160
r3 expandBy: (20@10 corner: 10@5)	10@80 corner: 140@155
r1 intersect: r2	40@30 corner: 70@50
r1 merge: r3	0@0 corner: 130@150
r1 contains: r2	false
r1 intersects: r2	true
r3 containsPoint: 50@100	true
r1 moveBy: 10@20	10@20 corner: 80@70
r1 moveTo: 100@100	100@100 corner: 170@150
r1 scaleBy: 2	200@200 corner: 340@300
r1 scaleBy: 2@4	200@400 corner: 340@600
r1 translateBy: -50	50@50 corner: 120@100
r1 translateBy: -50@50	50@150 corner: 120@200

# 15 The Graphics Classes

An object can be represented in different ways. For example, a collection of numbers may be seen as

- a sequence of numbers
- a histogram
- a piechart
- a curve passing through points

a rectangle may be seen as

- two points giving the extremities of the rectangle
- four right-angled segments
- a system of four mathematical equations.

We have seen that all the objects of the Smalltalk system respond to the message `printString` by returning a string of characters. This string allows the characteristics of the object to be explained, that is, its class and its instance variables. For certain objects, this representation is not helpful; they require graphic representation in addition.

In the Smalltalk system, the graphics interface consists of a bit-mapped screen. A bit-map screen is a screen in which each point is individually accessible, and where the characteristics of each point are stored in memory. A point on the screen is called a pixel, and in the case of Smalltalk each pixel can have one of two possible states, on or off. Thus, a single bit is sufficient to store this state and there are as many bits in memory as there are pixels on the screen.

## 15.1 The class `Bitmap`

The class `Bitmap` is a subclass of the abstract class `ArrayedCollection`. It therefore has indexed instance variables, each of which allows a machine word to be stored. The instances of the class `Bitmap` are able to store sequences of pixels. Since the size of one of these instance variables is the same as that of the machine word, pixel manipulation is carried out efficiently.

## 15.2 The class `DisplayObject`

The class `DisplayObject` is an abstract subclass of the class `Object`. An instance of the class `DisplayObject` represents a picture described by its width, its height,

and a point that provides the offset applied when it is displayed. The class `DisplayObject` contains five subclasses, each of which adopts a different way of storing a graphics object. The methods supported by the class `DisplayObject` are:

width  
returns the width of the minimum rectangle containing the receiver,  
height  
returns the height of the minimum rectangle containing the receiver,  
extent  
returns a point whose coordinates are the width and height of the minimum rectangle containing the receiver,  
offset  
returns the point giving the offset of the minimum bounding rectangle when the receiver is displayed,  
offset: anOffset  
replaces the offset of the receiver by the argument anOffset,  
scaleBy: aPoint  
multiplies the offset of the receiver by the argument aPoint, coordinate by coordinate,  
translateBy: aPoint  
adds the argument aPoint to the offset of the receiver,  
align: aPointOfAlignment with: anOtherPoint  
translates the offset of the receiver such that the argument aPointOfAlignment is placed at the position of the argument anOtherPoint,  
boundingBox  
returns the minimum rectangle containing the receiver, taking the offset into account,  
display  
displays the receiver at the point 0@0,  
displayAt: aPoint  
displays the receiver at the point aPoint.

### 15.2.1 The class `DisplayText`

The class `DisplayText` is a subclass of the abstract class `DisplayObject`. The objects that we want to represent in the class `DisplayText` are instances of the class `Text`. Recall that a text is made up of a string of characters and of an instance of the class `RunArray` representing the style of the characters. To display a text, additional data is required, like for example the graphical representation of the character fonts or what action is to occur in the case of tabulation. Such information is collected together in the class `TextStyle`, each instance of which represents a style of display.

### 15.2.1.1 Creating an instance of the class `DisplayText`

The message `asDisplayText`, when sent to an instance of the class `String` or to an instance of the class `Text`, returns a new instance of the class `DisplayText` initialised with the character string or the text. For example:

```
Time now printString asDisplayText displayAt: 100@100  
displays the current time at position 100@100.
```

Two other messages sent to the class `DisplayText` allow instances to be created:

```
text: aText
```

returns an instance of the class `DisplayText` consisting of the text given as argument; the style of the instance is given by a shared variable (pool variable) placed in the dictionary `TextConstants` and called `DefaultTextStyle` (recall that the dictionaries containing the pool variables are referenced by global variables).

```
text: aText textStyle: aStyle
```

returns an instance of the class `DisplayText` consisting of the text given as argument with as style the argument `aStyle`.

### 15.2.1.2 Instance methods

The instance methods of the class `DisplayText` are:

```
string
```

returns the character string that makes up the text of the receiver,

```
text
```

returns the text of the receiver,

```
text: aText
```

replaces the text of the receiver by the argument `aText`,

```
setText: aText textStyle: aStyle offset: anOffset
```

replaces the text, style and offset of the receiver by the values given as argument,

```
textStyle
```

returns the style of the receiver,

```
textStyle: aStyle
```

replaces the style of the receiver by the argument `aStyle`.

### 15.2.2 The class `DisplayMedium`

The abstract class `DisplayMedium` is another subclass of the class `DisplayObject`. The graphics object is stored in pixel form. The class `DisplayMedium` implements additional messages that allow a drawing to be filled with a motif and a border added to the rectangle containing the drawing. A motif is a 16 x 16 pixel drawing that is repeated in every direction. The messages can be divided into categories, as follows.

### 15.2.2.1 Background colour

The messages that fill a drawing with a motif are:

```
black  
fills the receiver with black,  
black: aRectangle  
fills the intersection between the argument aRectangle  
and the rectangle containing the receiver with black,  
darkGrey  
fills the receiver with dark grey,  
darkGrey: aRectangle  
fills the intersection between the argument aRectangle  
and the rectangle containing the receiver with dark  
grey,  
lightGrey  
fills the receiver with light grey,  
lightGrey: aRectangle  
fills the intersection between the argument aRectangle  
and the rectangle containing the receiver with light  
grey,  
veryLightGrey  
fills the receiver with very light grey,  
veryLightGrey: aRectangle  
fills the intersection between the argument aRectangle  
and the rectangle containing the receiver with very  
light grey,  
white  
fills the receiver with white,  
white: aRectangle  
fills the intersection between the argument aRectangle  
and the rectangle containing the receiver with white.
```

### 15.2.2.2 Drawing borders

The messages that add a border to an instance of the class DisplayMedium are:

```
border: aRectangle width: aWidth  
draws a black border inside the rectangle aRectangle of  
width aWidth,  
border: aRectangle width: aWidth mask: aMotif  
draws a border with the motif aMotif inside the rect-  
angle aRectangle of width aWidth,  
border: aRectangle widthRectangle: aRectangleInterior  
mask: aMotif  
constructs a rectangle from the rectangle aRectangle and  
from the rectangle aRectangleInterior, in which the new  
rectangle has as origin the origin of the first, plus  
the origin of the second, and as corner the corner of  
the first less the corner of the second. It fills the  
difference between the rectangle aRectangle and the new  
rectangle with the motif aMotif.
```

**15.2.2.3 Filling with a motif**

The aim of these messages is to compose two pictures, the first having the same size as the receiver, with background colour a motif propagated in all directions; the second being the picture of the receiver.

There are several ways of composing the two pictures. In the Smalltalk system an integer is assigned to each composition rule. There are 16 of them.

<b>Composition rule</b>	<b>Final image</b>
-------------------------	--------------------

---

0	all the pixels are white
1	logical AND between the two pictures
2	logical AND between the first picture and the second inverted picture
3	the final picture is the first picture
4	logical AND between the first inverted picture and the second picture
5	the final picture is the second picture
6	exclusive OR between the two pictures
7	logical OR between the two pictures
8	logical AND between the two inverted pictures
9	exclusive OR between the first inverted picture and the second picture
10	the final picture is the second inverted picture
11	logical OR between the first picture and the second inverted picture
12	the final picture is the first inverted picture
13	logical OR between the first inverted picture and the second picture
14	logical OR between the two inverted pictures
15	all the pixels are black

The messages that allow a picture to be composed with a motif are:

fill: aRectangle mask: aMotif  
fills the intersection between the argument aRectangle and the rectangle containing the receiver with the motif aMotif, applying composition rule 3,

fill: aRectangle rule: aCompositionRule mask: aMotif  
fills the intersection between the argument aRectangle and the rectangle containing the receiver with the motif aMotif, applying composition rule aCompositionRule,

reverse  
inverts the picture of the receiver, that is, fills the

picture of the receiver with a black motif, applying composition rule 6 (exclusive OR),  
reverse: aRectangle  
inverts the intersection of the receiver and the argument aRectangle,  
reverse: aRectangle mask: aMotif  
fills the intersection of the receiver and the argument aRectangle with the motif given as argument, applying composition rule 6 (exclusive OR).

### 15.2.3 The class Form

The class **Form** is a subclass of the class **DisplayMedium**. An instance of the class **Form** allows a rectangular picture to be stored. This picture is described by its height, its width and the value of the pixels that make it up. The instance variable called **bits** which stores pixels is an instance of the class **Bitmap**.

#### 15.2.3.1 Creating a form

Six messages create and initialise a **Form**:

extent: aSize  
returns a new instance of **Form** whose size is equal to the argument **aSize**, and whose **bitmap** is not yet initialised,  
extent: aSize fromArray: anArray  
returns a new instance of **Form** whose size is equal to the argument **aSize**, and whose **bitmap** is initialised to the values contained in the argument **anArray**,  
extent: aSize fromArray: anArray offset: anOffset  
returns a new instance of **Form** whose size is equal to the argument **aSize**, and whose **bitmap** is initialised to the values contained in the argument **anArray** and **offset** is initialised by the argument **anOffset** (the **offset** is a shift that is added when the **form** is displayed),  
fromDisplay: aRectangle  
returns a new instance of **Form**, created from the screen by taking the **bitmap** contained in the **rectangle** given as argument,  
fromUser  
returns a new instance of **Form**, created from the screen by taking the **bitmap** contained in the **rectangle** designated by the user,  
fromUser: aPoint  
returns a new instance of **Form**, created from the screen by taking the **bitmap** contained in the **rectangle** designated by the user; the position and size of this **rectangle** are multiples of the argument **aPoint**,  
readFrom: aFileName  
returns a new instance of **Form** read from the file called **aFileName**.

### 15.2.3.2 Other class methods

The Smalltalk system contains a number of predefined forms used to construct motifs. The size of these forms is 16 x 16 pixels. The forms available are:

black	returns a form in which all the pixels are black
darkGrey	returns a form in which almost all the pixels are black
grey	returns a form in which one pixel is black
lightGrey	returns a form in which the number of white pixels is greater than the number of black
veryLightGrey	returns a form in which almost all the pixels are white
white	returns a form in which all the pixels are white

Some class methods return an integer corresponding to a composition rule. These methods are:

and	
returns 1	which corresponds to the logical AND between the two pictures,
over	
returns 3	which corresponds to the replacement of the first picture by the second,
erase	
returns 4	which corresponds to the erasure in the first picture of all the pixels that are black in the second,
reverse	
returns 6	which corresponds to the exclusive OR between the two pictures,
under	
returns 7	which corresponds to the logical OR between the two pictures, giving the impression that the second picture is located under the first.

### 15.2.3.3 Instance methods of the class Form

Some useful messages supported by the class Form are:

bits	
returns the instance of the class Bitmap that stores the pixels,	
bits: aBitmap	
replaces the bitmap of the receiver by the argument aBitmap	
borderWidth: aWidth	
inserts a black border of width aWidth around the receiver,	

```

displayOn: aDisplayMedium
at: aPoint
clippingBox: aRectangle
rule: aCompositionRule
mask: aMotif
displays the receiver on the argument a DisplayMedium at
the position aPoint, without exceeding the limits of
aRectangle and applying the mask aMotif with the compos-
ition rule aCompositionRule,
follow: aBlock while: aBlockCondition
so long as evaluation of the block aBlockCondition
returns true, the block aBlock is evaluated to give the
position at which the receiver must be displayed and the
zone hidden by the receiver at the previous evaluation
is redisplayed,
valueAt: aPoint
returns 0 or 1 depending on whether the pixel located at
the position aPoint in the receiver is white or black,
valueAt: aPoint put: aBit
sets the pixel located at the position aPoint in the
receiver to white or black depending on whether the
argument aBit is 0 or 1.

```

#### 15.2.4 The class Cursor

The class Cursor is a subclass of the class Form. Depending on the system status, different cursors appear. These cursors indicate, by their position, the place on the screen pointed to by the mouse, and by their form the status of the system. The cursor forms occupy 16 x 16 pixels.

The message cursorLink: sent to the class Cursor allows movement of the cursor to be linked or not to the movement of the mouse, depending on whether its argument is true or false.

The messages supported by the class cursor are:

```

show
where the receiver becomes the current cursor,
showGridded: aPoint
the receiver becomes the current cursor, with its
position being a multiple of the argument aPoint,
showWhile: aBlock
the receiver becomes the current cursor while the block
given as argument is evaluated, after which the current
cursor becomes the previous cursor.

```

#### 15.2.5 The class DisplayScreen

The class DisplayScreen is a subclass of the class DisplayMedium. The instances of the class DisplayScreen serve to represent the complete screen (the screen is a rectangular form comprising a certain number of pixels; it may therefore be likened to an instance of the class

Form). There is in the Smalltalk system a single instance of the class `DisplayScreen` referenced by the global variable `Display`. Note the following message:

flash: aRectangle

which causes to flash that part of the screen corresponding to the rectangle given as argument.

#### 15.2.6 The class Path

The class `Path` is a subclass of the class `DisplayObject` which is used to draw a form along the length of a "trajectory". A trajectory is an `OrderedCollection` of points. The instances of the class `Path` have two instance variables that are a form and a trajectory. The messages of the class `Path` may be divided into the following categories.

##### Access

The messages that allow access to the characteristics of the instances are:

form

returns the form referenced by the receiver,

form: aForm

replaces the form referenced by the receiver by the form `aForm`,

at: anIndex

returns the point of the trajectory referenced by the receiver of index `anIndex`,

at: anIndex put: aPoint

replaces the trajectory point of index `anIndex` by the point `aPoint`,

size

returns the number of points of the trajectory referenced by the receiver.

##### Test

The message `isEmpty` returns true if the trajectory referenced by the receiver contains no point.

##### Add

The message `add:` adds the point given as argument at the end of the trajectory referenced by the receiver.

##### Remove

The message `removeAllSuchThat:` removes from the trajectory referenced by the receiver all the points for which evaluation of the block given as argument returns true.

##### Enumeration

The three messages that allow the points of a trajectory to be handled successively are:

do: aBlock

evaluates the block given as argument for each point of the trajectory referenced by the receiver,

collect: aBlock

evaluates the block given as argument for each point of the trajectory referenced by the receiver and returns an instance of OrderedCollection containing the results of each evaluation,

select: aBlock

evaluates the block given as argument for each point of the trajectory referenced by the receiver and returns an instance of OrderedCollection containing the points for which the evaluation has returned true.

The subclasses of the class Path correspond to particular trajectories. These are:

- straight lines
- arcs of circles
- circles

#### 15.2.6.1 The class Line

The class Line is a subclass of the class Path that allows rectilinear trajectories to be represented by giving the two extremities of the trajectory. The additional messages implemented in this class are:

beginPoint

returns the first point of the receiver,

endPoint

returns the second point of the receiver,

beginPoint: aPoint

replaces the first point of the receiver by the point aPoint,

endPoint: aPoint

replaces the second point of the receiver by the point aPoint.

#### 15.2.6.2 The classes Arc and Circle

The classes Arc and Circle are subclasses of the class Path which allow sections of circles and circles to be represented.

### 15.3 The class BitBlt

The class BitBlt is a subclass of the class Object. (The name BitBlt stands for Bit Block Transfer, a technique for altering and copying rectangular areas of a picture.) The instances of the class are objects that are capable of transferring pixels of one form to another form, while respecting certain constraints given by their instance variables. The instance variables that control pixel transfer are:

destForm

which is a form in which the transferred pixels will be stored,

sourceForm

which is a form from which the pixels to be transferred are taken,

halftoneForm

which is the motif with which the starting form will be composed. There are four possible cases:

- sourceForm and halftoneForm are equal to nil; the form transferred will be a black form
- sourceForm is equal to nil; the form transferred will be the motif propagated in all directions
- halftoneForm is equal to nil; the form transferred will be the sourceForm
- halftoneForm and sourceForm are not nil; the form transferred will be obtained by making a logical AND between the sourceForm and the motif propagated in all directions.

combinationRule

which is an integer representing the composition rule of the starting form and arrival form (cf composition rules defined in 15.2.2).

sourceX, sourceY, width, height

which are four integers designating the region of the starting form from which the pixels will be copied,

destX, destY

which are two integers designating the region of the destination form where the pixels will be placed,

clipX, clipY, clipWidth, clipHeight

which are four integers that limit the region that can be altered by the copy.

The message that makes the copy, while respecting the constraints imposed by the instance variables, is copy-Bits. Generally, the call to BitBlt is made by other classes like Form, Line, etc., and the user does not have to invoke it explicitly.

The class BitBlt implements another important message:

drawFrom: aStartPoint to: anEndPoint

which allows a form to be propagated between the two points given as argument.

#### 15.4 The class Pen

The class Pen is a subclass of the class BitBlt. The instances of the class Pen simulate a stylus that can be moved on the screen and which leaves a trace, or does not, depending upon whether it is lowered or raised.

Since the class Pen is a subclass of the class BitBlt, drawing operations are carried out more quickly using it rather than instances of the class Line.

The instances of the class Pen have four instance variables:

penDown

which is a boolean indicating whether the stylus is lowered or raised,

location

which is a point indicating the current stylus position,  
direction

which is an angle expressed in degrees indicating the direction of the stylus trace,

frame

which is a rectangle indicating the limits within which the stylus can trace.

The messages of the class Pen may be divided into the following categories.

#### **Access**

The messages that allow access to the characteristics of the stylus are:

down

lowers the receiver,

up

raises the receiver,

location

returns the current position of the receiver,

direction

returns the direction in degrees of the trace of the receiver,

turn: anAngleInDegrees

adds to the current direction of the trace of the receiver the argument anAngleInDegrees, which is equivalent to turning the angle anAngleInDegrees in relation to its current direction,

north

orients the trace of the receiver towards the top of the screen,

frame

returns the rectangle in which the receiver may trace,

frame: aRectangle

replaces the rectangle in which the receiver may trace by the argument aRectangle,

defaultNib: aWidth

replaces the width of the receiver by the argument aWidth,

colour: aMotif

replaces the motif of the receiver by the motif given as argument.

#### **Movement**

The messages that allow the stylus to be moved are:

place: aPoint

replaces the current position of the receiver by the point aPoint,

home

places the receiver at the centre of the rectangle in which it may trace,  
go: aDistance  
moves in the current direction by a distance aDistance,  
tracing if the receiver is lowered,  
goto: aPoint  
moves in a straight line to the point given as argument,  
tracing if the receiver is lowered.

# 16 Basic Elements of the Interface

Every information system has an interface to allow communication with a user. This interface usually consists of a keyboard and a screen.

There are two basic types of screen, a character type and a graphics type. Smalltalk uses a graphics screen.

Two technologies are used for graphics screens, depending on their function - vector scan screens and raster scan screens. The former allow lines to be drawn very quickly, and they are used for CAD and similar applications. The latter manipulate points; if these points are individually accessible, such screens are said to be bit-mapped. It is this last type of screen that we find in the Smalltalk system.

The number of points on the screen depends on the machine and the Smalltalk system can adapt itself to this number of points.

The Smalltalk system uses a standard keyboard as input device, but another device, commonly known as a "mouse", is also required.

## 16.1 Mouse

A mouse is a device that allows the user to designate a point on the screen by moving the mouse around on a flat surface. The mouse incorporates buttons that cause different operations to occur when they are pressed. There are usually three buttons in the Smalltalk system.

These buttons are identified by colour:

the left button is the red button (redButton),

the centre button is the yellow button (yellow-Button),

the right button is the blue button (blueButton).

In the Smalltalk system each button has a particular function:

the red button is used to select or designate a visible area,

the yellow button is used to call up the menu that allows the inside of the viewed area to be changed,

the blue button is used to call up the menu that allows the viewed area itself to be changed.

## 16.2 Windows

A window is a rectangular portion of the screen in which

a program specific to this window runs, independently of the content of other windows. For each application there are specialised windows. For example, there are windows in which text can be edited and others in which drawings can be edited.

In the Smalltalk system, activities are not assigned to the screen but to the individual windows that go to make up the screen.

A window is said to be active when the user can interact with the contents of that window. There can be only one window active at a time.

When windows are superimposed, the active window is the one that is visible.

If there is no active window, as soon as the cursor moves inside a window, that window becomes active. To change windows, the cursor must be placed in an inactive window and a button on the mouse pressed. If one of the mouse buttons is pressed with the cursor outside any window, there is no longer any window active.

### 16.3 Pop-up menus


A pop-up menu is a window that appears at the cursor position when the yellow or blue button on the mouse is pressed. This menu remains displayed for as long as the user holds the button down.

These menus offer a list of commands. To select a command the user must move the cursor to the command that he wishes to execute and then release the button. Positioning the cursor on the command has the effect of inverting the display of this command line. If the user releases the button without selecting any command, the pop-up menu disappears and no command is executed.

Pop-up menus are instances of the class `PopupMenu` which is a subclass of the class `Object`.

### 16.4 Standard windows

One of the important window classes is the class `StandardSystemView`. The instances of this class are windows that have a title indicating the functions of the view. These windows can be manipulated; they can be moved, altered in size, etc. These actions are carried out by means of a pop-up menu which appears on the screen when the blue button is pressed while within the window. The diagram below shows a menu example that appears when the blue button is pressed in a standard window:



The actions possible are:

under

calls for a window below the cursor to be activated. If there are several inactive windows under the cursor, one of these will be activated. If there is no inactive window under the cursor, this command has no effect,

move

moves the active window without altering its size. The window follows the movement of the cursor until the red button is pressed,

frame

changes the dimensions of the active window. An origin cursor follows the movement of the mouse until the red button is pressed. The position thus obtained corresponds to the top left corner of the window. Subsequently, a corner cursor follows the movement of the mouse for as long as the red button remains pressed. The new position obtained corresponds to the bottom right corner bottom of the window.

collapse


removes all of the window except its title. To make this window reappear, the command frame must be used to change its dimensions,

close

causes the window to disappear completely from the screen.

### **16.5 The system menu**

The system menu is the pop-up menu that appears when the yellow button is pressed outside any window. This is the menu that, for example, is used to create windows or to exit the Smalltalk system. The diagram below shows an example of a system menu.



The commands offered by this menu are:

restore display

redisplays the complete screen, beginning with the background in a grey motif, then displaying all the windows,

exit project

allows you to exit from the current project. A project is a list of windows that relate to a particular application. When you return to a project, the screen is cleared and the windows of this project are displayed. When you have entered a project, the system menu is the same,

- project
- opens a window that will allow entry to a project,
- file list
- opens an access window to the external files,
- browser
- opens a window that allows classes and methods to be edited,
- workspace
- opens a window that allows text to be edited and expressions evaluated,
- system transcript
- opens a window used by the Smalltalk system to display its messages,
- system workspace
- opens a text window whose contents are initialised by the Smalltalk system with a number of expressions that are useful to know,
- save
- allows the Smalltalk virtual image to be saved to a file,
- quit
- allows you to quit the Smalltalk system.

# 17 Text Editing Windows

Numerous types of window use a text editor. Whatever the window type, the editor used is always the same.

To edit text, the user can use the keyboard and the mouse. The keyboard provides the means of entering new characters, while the mouse allows a section of text to be selected with the red button or commands to be executed using the yellow button.

If a section of text is too long to be displayed all at once in the window, there needs to be a means of displaying only part while retaining easy access to the remainder. This is what happens on a conventional word processor page, where at most 24 lines are displayed, but the rest of the text can be accessed by means of arrow keys. In the Smalltalk system, the windows have another means of accessing a particular section of text.

## 17.1 Scrollbars

A "scrollbar" is a rectangular area that appears on the left of an editing window of active text when the cursor is located inside the window.

Whenever the cursor moves outside of this window, the scrollbar disappears, only to reappear again when the cursor returns to the window.


This scrollbar represents the totality of the text. It contains a grey rectangle whose position in the scrollbar represents the position of the text shown on the screen relative to the text in its entirety. Its height, in relation to the height of the scrollbar, represents the proportion of text shown on the screen in relation to the text in its entirety.

To access a particular section of text, the rectangle within the scrollbar must be moved. When the mouse points to the interior of the scrollbar, there are three different cursors, depending on the relative position of the cursor and of the interior rectangle.

If the mouse points to the left of the rectangle, the down cursor appears. If the user presses the red button, the inside rectangle moves up and the text displayed in the text window moves down, as shown in the following diagram.



If the mouse points to the right of the interior rectangle, the up cursor appears. If the user presses the red button, the interior rectangle moves down and the text displayed in the text window moves up, as shown below.




If the mouse points along the same level as the interior rectangle, the marker cursor appears. If the user presses the red button, the interior rectangle centres itself on the cursor and causes the text to move, as shown below.



### 17.2 How to select text


Selecting part of a text means choosing a continuous section of text with a view to manipulating it with an edit command.

In a text window there is a special cursor that indicates the position of the next text input. This cursor is shaped like an inverted "v", as can be seen on the previous three diagrams. This cursor can be moved around in the window by pressing and releasing the red button on the mouse until the correct position for the mouse cursor is found. There are several ways of selecting a text; these are described in the following paragraphs.





The mouse cursor is placed at a position on the text, the red button is pressed without being released and the mouse cursor is moved to another position on the text, then the button is released. The text located between these two positions will be displayed in reverse video, as shown in the figure on page 124 (bottom).

The mouse cursor is placed on a word, the red button is pressed and released twice without moving the mouse. This will select a complete word, as shown below.



The cursor is placed at the beginning or at the end of the text, the red button is pressed and released twice without moving the mouse. This will select all of the text, as shown on page 125.

The cursor is placed after or in front of a carriage return, the red button is pressed and released twice without moving the mouse. This will select that section of text that extends backwards or forwards to the next carriage return, as shown below.




The cursor is placed after or in front of a delimiter, the red button is pressed and released twice without moving the mouse. This will select that portion of text located between delimiters, as shown in the example at the foot of page 126.

If the Escape key on the keyboard is pressed, all text entered since the last use of the mouse buttons will be selected.

If the text contains a selected section and the user enters characters via the keyboard, the entered characters will replace those in the selected section.

### 17.3 Edit commands

The edit commands can be accessed by a menu that appears when the yellow button is pressed. Each command is executed as soon as the button is released. This menu is illustrated below.



The commands are:

cut

removes the selected section from the text,

copy

places a copy of the selected text in a buffer,

paste

replaces the selected text with the contents of the buffer,

again

re-executes the last command,

undo

cancels the effect of the last command.

### 17.4 Workspaces

A workspace is a text-editing window which makes it possible to evaluate expressions contained in a selected section of text. Smalltalk expressions can be evaluated

by two commands that can be accessed via the yellow button menu. They are:

doIt

which evaluates the selected text, and

printIt

which evaluates the selected text and displays the text that describes the object returned by the last evaluation. The returned text becomes the current selection.

When a text is evaluated, the text expressions are compiled. If there is a syntax error, an error message is displayed immediately in front of the position corresponding to the detected error. This message becomes the selected section of text, which allows it to be amended immediately or replaced with other text.

In the basic version of Smalltalk-80, when there is an unknown variable in an expression, a menu appears offering a choice of possible declarations for this variable. This menu also offers the possibility of correcting the variable name or of abandoning the evaluation.

More recent versions of Smalltalk-80 consider that variables can be local to a workspace. These variables will be defined when first assigned and will be stored in a dictionary local to the workspace.

During evaluation of expressions, the scrollbar disappears from the window. It reappears when the evaluation is finished.

### 17.5 System workspace

The System workspace is a workspace whose contents consists of expressions that it may be useful to know in order to change certain operating parameters in Smalltalk. The parameters that can be changed are, for example:

the size of the screen

the "sources" and "changes" files

The System workspace also contains expressions for reconstructing an image after an accidental break.

# 18 Browsers

The browsers are windows that allow access to the classes and methods. A browser has two possible uses; it allows the user to:

- consult or change the existing classes and methods. Since the Smalltalk environment is itself written in Smalltalk, the user has access to all of its code. From this point of view Smalltalk can be regarded as a self-documented language. The advantage of having access to the existing classes encourages the reuse of code. It is not advisable to change the existing classes or methods, as any such changes could have consequences well beyond any that could be foreseen.

- create new classes or new methods. In the Smalltalk environment, the building of an application amounts to enriching the environment with new classes and new methods.

## 18.1 Description


A browser is made up of five main windows. Four of these give access to the following lists:

- list of categories of classes,
- list of classes belonging to the selected category,
- list of categories of methods or protocols of the selected class,
- list of methods of the selected protocol.

The fifth window, located beneath the other four, is a Smalltalk code editing window, the contents of which will be interpreted according to the options selected from the lists.

There are two additional windows located under the list of classes; the one on the left contains the 'instance' text, that on the right the 'class' text. These two windows allow access to a class or to its metaclass, depending on whether the left or right window is selected (indicated by reverse video). If the left window is selected, the methods displayed are the instance methods; if the right window is selected, the class methods are displayed. These two windows cannot be selected simultaneously.

Each window has a particular menu linked to the central button on the mouse. The diagram below shows these different menus.



## 18.2 The categories of class

The categories of class group together the classes by function. Once again, be reminded that the categories have no influence on the behaviour of the system. The list of these categories appears in the left window. The menu in this window, linked to the central button, contains eight functions, namely:

*file out*

which, for all classes of the selected category, outputs the text of the class and its methods to a file. This file can then be reread by the message *fileIn*. The name of the file produced is the name of the category with the suffix '*.st*'.

*print out*

which causes the same action as the function '*file out*', except that in this case the file created cannot be reread. Special formatting is inserted into the text so that when printed it can be more easily read by humans. In particular, the names of methods appear in bold characters. The name of the file produced is the name of the category with the suffix '*.pp*'. (Note that the exact nature of this option varies from machine type to type, so your machine may do something different.)

*spawn*

which opens a browser dedicated to the selected category,

*rename*

which allows the name of the selected category to be changed,

*remove*

which removes the selected category, along with all the classes that belong to it,

update

which updates the list of categories. This list may have been changed in another browser or by reading a file,

edit all

which presents the list of categories and their classes in the Smalltalk code editing window of the browser. This list can then be changed using the menu available in that window.

### 18.3 The classes

The list of classes of the selected category appears in the second window. If no class category is selected, this window is empty. The functions offered by the menu in this window are:

file out


as before, this outputs the definition of the selected class and the text of its methods to an external file. The name of the file produced is the name of the class concatenated with the suffix '.st',

print out


as for the class categories, this outputs the definition of the selected class and the text of its methods to an external file. The name of the file produced is the name of the class with the suffix '.pp',

spawn


opens a browser dedicated to the selected class, as shown in the diagram below for the class Point:



spawns hierarchy  
 opens a browser on the hierarchy of the selected class  
 (that is, its subclasses and superclasses), as shown in  
 the following diagram:



hierarchy  
 displays in the code window the hierarchy of the selected class; the depth of the classes in their hierarchy is represented by indentation of the text,  
 definition  
 displays in the text window the definition of the selected class,  
 comment  
 displays in the code window the commentary relating to





the selected class. This can then be changed, protocols

displays in the code window the list of protocols and methods belonging to these protocols, as shown on page 132 (foot). The name of a protocol is a string of characters, whereas the name of a method is a symbol.

```
inst var refs
```

causes a pop-up menu to appear containing the list of named instance variables of the selected class and of its superclasses. It is then possible to select one of the variables from the offered menu, which has the effect of opening a browser on the list of methods that make reference to this variable. The following two diagrams show the selection of the instance variable `labelText` from the class `StandardSystemView`:




class var refs

causes a pop-up menu to appear containing the list of class variables of the selected class and its super-classes. It is then possible to select one of the class variables from the menu, which has the effect of opening a browser on the list of methods that make reference to this variable.

class refs

opens a browser on the list of methods that make reference to the selected class, as shown in the diagram below:



rename

allows the name of the selected class to be changed. If any methods make reference to this class, a browser is opened on the list of these methods,

remove

removes the selected class and its subclasses if they exist.

#### 18.4 Protocols


The third window contains the list of protocols of the selected class. If no class is selected, this window is empty. The reader is reminded that the sole function of protocols in the system is to clarify the presentation of the methods in browsers. The functions offered by this menu are:

file out

which outputs the text of all the methods belonging to the selected protocol. The name of the file produced is the name of the class concatenated with the name of the protocol, concatenated with the string '.st',

print out  
 causes the same action as file out, but produces a formatted text,

spawn  
 opens a browser on the list of methods of the selected protocol. The figure below shows an opened browser on the protocol comparing of the class Point:



add protocol  
 allows a protocol to be added to the list of existing protocols. The name of this protocol is requested in an edit window,  
 rename  
 allows the name of the selected protocol to be changed,  
 remove  
 removes the selected protocol, together with the list of all its methods.

### 18.5 Methods

The list of methods belonging to the selected protocol is given in the right window. If no protocol is selected, this window is empty. The functions offered by the menu are:

file out  
 which allows the text of the selected method to be output. The name of the file produced is the name of the class, concatenated with the name of the method, concatenated with the string '.st',

print out  
 causes the same action as file out, but produces a formatted text,

spawn  
 opens a browser on the selected method,  
 senders  
 opens a browser on the list of methods that send the message with the same name as the selected method,

implementors  
 opens a browser on the list of methods with the same name as the selected method,  
 messages  
 causes a menu to appear giving the list of all the messages sent in the selected method. Using this menu, it is possible to access all the implementations of one of the messages given in the menu,  
 move  
 allows the selected method to be moved. If it is moved to another class, it is copied; if its protocol is changed, it is simply moved,  
 remove  
 removes the selected method.


### 18.6 The code window

The Smalltalk code window, which appears in the browser, is like the text windows described in the previous chapter. In addition to its editing functions, this window offers three more functions in its menu:

format  
 this function relates exclusively to methods and will thus only be valid when a method is displayed. Its purpose is to present the code of the method in the clearest possible form. To achieve this, the name of the method is shown in bold and some conventions in its presentation are observed; the most important of these are:


- there is only one statement per line
- loops are shown with additional indentation.

The figure below shows the effect of the formatting on a method.



spawn  
which opens a new browser on the selected method,  
explain

this function is very useful; it allows the nature of the selected word to be explained. Only one word can be selected at a time. The following example shows the explanation provided by the system for the word 'x'. This word can be explained in two ways: x is an instance variable of the class Point, and x is a message defined in the class Point.



Depending upon the options chosen from the list at the top, the code window has two main uses, described in the following sections.

#### 18.6.1 Editing a class

When a category is selected, without a class being selected, a text giving the general form of a class is displayed; this can be changed. To capture this newly defined class, the accept function accessible via the middle button must be used.

When a class is selected, without a protocol being selected, the definition of this class appears in the text window. It is then possible to change and accept it, as before.

#### 18.6.2 Editing a method

When a protocol is selected, without a method being selected, the general form of a method is displayed in the text window. The user may change this general form in order to produce a new method. If the name of the method that is displayed is the same as that of a method already in existence in this class, the newly accepted

method will replace the old. To compile the method, the accept option (via the centre button) must be used.

When a method is selected, the text of this method is displayed in the text window. This text may then be changed. If the method name is changed, a new method will be created.

# 19 Other Windows

## 19.1 Inspectors


Just as a browser allows classes to be accessed and changed, so an inspector allows access to any object in the system. It provides access to all the instance variables of the object inspected and all the Smalltalk objects can be inspected. Inspectors are used chiefly during the development phase of an application.

To open an inspector on an object, the message inspect has to be sent to it. For example:

```
#(1 2 3 4) inspect  
opens an inspector on the receiver of the message,  
Set someInstance inspect  
opens an inspector on an instance of the class Set.
```

### 19.1.1 Description

An inspector type window consists of two subwindows, as shown in the figure below.



The left window contains a list made up of the symbol `self`, representing the object inspected, and the list of names of its instance variables. In the case where the inspected object has indexed instance variables, the latter are represented by their index.

In the left window, the menu linked to the middle button offers the choice of inspecting the object that is selected. If the selected object is `self`, a second inspector will be opened on the same object.

The right window is a text window which contains the textual description of the object that is selected from the list on the left. The reader is reminded that this textual description is obtained by sending the message `printString` to the selected object.

The menu in this window, linked to the middle button, offers the usual functions suggested in a text window, namely:

- three text-editing functions: cut, copy and paste
- two evaluation functions: do it and print it
- two functions for manipulating the window contents: accept and cancel. The contents are the text that appears when it is open.

In contrast to the text windows that we have met before (for example, workspaces), evaluations in an inspector are carried out in the context of the inspected object, that is, at the time of evaluation references can be made to an object contained in the list on the left. In particular, an expression can be built by sending a message to `self` or to one of its instance variables, which is not the case in a workspace.

The function `accept` assigns the result of the evaluation of the text present in the text window to the variable selected from the list on the left.

### 19.1.2 Other inspectors

We have seen how in order to inspect an object the message `inspect` has to be sent to it. The method for responding to this message is defined in the class `Object`; thus all objects can be inspected. In some classes, this message is redefined, which allows the opening of windows adapted to certain types of objects.

Dictionaries have special inspectors that allow keys to be manipulated. The following figure shows an inspector on a well known dictionary, namely the dictionary called `Smalltalk`, which contains all the global variables in the system, it will be recalled. In the left window the keys of the dictionary appear, while in the right window the value assigned to the selected key appears.

The left window offers the following options:  
inspect  
which opens an inspector on the selected variable,

references


which opens a browser on all the methods that refer to the selected variable,

add field

which allows a key to be added to the dictionary,


remove

which removes the selected key from the dictionary.



Views (instances of the class View) also have special inspectors which allow access to the model and to the controller assigned to the view to form a MVC (model-view-controller) system. The MVC system is described in detail in chapter 21.

The figure below shows an inspector on an instance of



BrowserView, which is the view assigned to a browser. This inspector is in fact made up of three inspectors on the three objects of the MVC system.



## 19.2 Debuggers

Development tools are very important tools in programming environments and in Smalltalk they have been very carefully designed.

Debuggers are windows that serve to visualise the methods and variables implicated in the evaluation of an expression during the actual evaluation. A debugger can be opened from a window called a Notifier, in which the menu, controlled by the middle button, offers two actions:

```
proceed
which resumes the interrupted activity,
debug
which opens a debugger on the interrupted context.
```

The Notifier contains the list of the last methods that were used before the interruption, as shown in the following figure.



Several circumstances can cause the opening of a Notifier. These are usually:

- an object does not understand a message that it receives,
- the user causes interruption of an activity by pressing control C,
- an object receives a halt message.

A debugger type window is made up of several subwindows, as shown in the bottom figure on page 142. The six windows involved are described below.

#### **19.2.1 List of methods**

This list contains the name of the last methods that were active before the interruption. When the user selects one of the methods contained in the list, the text of this method appears in the code window located below, with the expression that was being evaluated when the interruption occurred highlighted. This subwindow contains a menu controlled by the middle button which offers the following options:

full stack

which completes the list of methods by adding to it all the methods that were active within the process,

proceed

which resumes execution of the interrupted process,

restart

which restarts evaluation of the selected method from the beginning,

senders

which opens a browser on the list of methods that send a message with the same name as the selected method,

implementors

which opens a browser on the list of methods with the same name as the selected method,

messages

which cause a pop-up menu to appear giving the list of all messages sent from the selected method. From this menu it is possible to access all the implementations of a message present in the menu,

step

which causes the next expression to be evaluated. The selected part in the code window is updated,

send

which causes the next message in the selected expression to be sent.

#### **19.2.2 The code window**

This window is located under the window that provides the list of methods that have been used. When a method is selected from the list of methods, its text appears in the code window, as in a browser. The expression that

was being evaluated at the time of the interruption is shown in inverse video. The functions offered in the menu, which is controlled by the centre button, are the same as those offered in the code window of a browser.

#### **19.2.3 Inspector of the receiver**

This inspector is an inspector which describes the receiver of the message selected from the list at the top. In this window it is possible to access all the instance variables of the receiver and change them.

#### **19.2.4 Inspector of local variables**

The inspector enables access to the arguments of the method selected from the list of methods, as well as access to all its local variables. It is possible to consult or change these variables.

### **19.3 Views of external files**

These windows, called 'fileLists', allow access to the file system underlying the Smalltalk virtual machine. The windows allow access to both directories and files. They are used in particular for integrating new classes into the system; rather than transfer complete images, it is more economical to transfer definitions of classes and methods that have been obtained by means of a fileOut.

The fileList type windows are made up of three sub-windows, as shown in the following figure.

The top window is a text window in which the user enters the name of the file that is to be accessed. The character '\*' can be used to replace any character string when designating the file; for example:

\*.st

designates all the files suffixed by '.st', that is, all the files obtained by a fileOut in a browser.

Once the name of the file(s) has been entered, the user can obtain a list of all the filenames matching the name entered by using the accept function controlled by the middle button.

The central window contains the list of the files whose name is compatible with the filename specification of the top window.


If there are files whose name corresponds to the name in the top window, the menu controlled by the middle button offers the following options:

get contents

allows the text from the selected file to be loaded into the bottom window,

fileIn

this operation carries out the opposite function to the browser fileOut; it allows the text of the selected file



to be compiled. The latter must contain the definitions of classes and methods; that is, it must have been obtained by a fileOut,

copy name

which allows the name of the selected file to be copied into a text buffer (this is the equivalent of the copy function in a text window),

rename

which allows the name of the selected file to be changed. The name of a new file must not be the name of a file that already exists,

remove

which allows the selected file to be destroyed.

When the name of the top window does not correspond to any existing file, the central window contains a list, whose only element is the name entered in the top window, together with a menu whose options are:

copy name

which allows, as before, the name of the selected file to be copied into the text buffer,

rename

which allows the selected file to be renamed,

new file

which allows a new file to be created whose name will be

the selected name  
new directory

which allows a new directory to be created whose name will be the selected name.

The bottom window is a text window; it contains the text of the selected file in the central window if the option 'get contents' has been used in this window. The menu offered is the one generally offered in this type of window, but it contains three additional options, namely:

file it in


which allows the selected text to be compiled by interpreting it as text defining classes and methods, that is, this text must be in the format of the external files,

put

which writes the text in the bottom window to the selected external file,

get

which replaces the text in the bottom window with the text from the selected file.



#### 19.4 Graphics editors

The Smalltalk system has two graphics editors that allow a form to be drawn. Such forms are instances of the class `Form` (see chapter 15) or of one of its subclasses.

All the entities displayed on the screen are forms (drawing, cursor, character font) and can therefore be edited graphically.

##### 19.4.1 The Form Editor

This editor consists of two windows, one containing the form to be edited and the other a graphics menu from which the user selects the functions required to draw on the form. To edit a form with this graphics editor necessitates sending it the message `edit`. The figure on page 146 illustrates the graphics editor.

To select an option from the graphics menu the user places the cursor on the option and presses the red button on the mouse. The options in this menu are as follows:



allows selection of the pattern that will serve as the paintbrush. This pattern is obtained by designating a rectangle on the screen.



indicates that the paintbrush is to be copied at the position indicated by the mouse as soon as the red button is pressed.



indicates that the paintbrush is to be copied at the position indicated by the mouse as long as the red button is pressed.



indicates that the paintbrush is to be copied along the line defined by two presses on the red button.



allows a curve to be drawn with the paintbrush.



allows a rectangle to be filled with the selected pattern.



indicates that paintbrush overwrites the background when it is applied.



indicates that the paintbrush is ORed on to the background.



indicates that paintbrush is XORed on to the background.



indicates that paintbrush is ANDed on to the background.



allows part of the drawing to be edited with the bit editor (see next section).



initialises the fill colour with a white pattern.



initialises the fill colour with a light grey pattern.



initialises the fill colour with a grey pattern.



initialises the fill colour with a dark grey pattern.



initialises the fill colour with a black pattern.



allows the horizontal grid to be used.



allows the vertical grid to be used.



allows the minimum number of points between two lines and two columns of the grid to be set.



allows the edited form to be output to an external file.



allows a form to be read in from an external file.

The menu linked to the yellow button has two options, namely:

accept


which allows the existing drawing form to be replaced by the one currently shown on the screen,

cancel

which replaces the drawing on the screen with the original edited drawing.

#### 19.4.2 The Bit Editor

There is a second graphics editor called by the message `bitEdit` which allows access to each pixel of the drawing. In this case, the drawing is enlarged and each pixel is represented by a square of 8 x 8 points. The figure below shows an example of the use of this editor.



This view is made up of two subviews, one giving access to the enlarged drawing and the other which lets you choose to draw in either black or white. The menu, controlled by the yellow button, is the same as that offered by the editor edit. This editor is generally used for small size forms; for example, it can be used to create a new cursor:

Cursor new bitEdit

# 20 Management of User Events

In order to interact with the Smalltalk system, the user can use two different devices - the keyboard and the mouse. Any alteration to the state of the keyboard or the mouse is called a user event. The user events that are possible are:

- a movement of the mouse,
- the pressing of one of the mouse buttons,
- the striking of a key on the keyboard.

Within the Smalltalk system there are several classes that define methods for detecting and managing user events. Because only one user may interact with a Smalltalk system at once, these classes have only one instance.

## 20.1 The class `InputState`

The class `InputState`, which is a subclass of the class `Object`, interfaces Smalltalk with the hardware that detects user events. There is in the system one instance of this class which at any time represents the status of the keyboard and the mouse. This class has a class variable called `InputProcess`, which references the process reading the status of the keyboard and the mouse; this process is an instance of `Process` (cf chapter 13).

## 20.2 The class `InputSensor`

The class `InputSensor`, which is a subclass of the class `Object`, interfaces between the user and the instance of `InputState` that represents the status of the keyboard and the mouse. It has a class variable called `CurrentInputState`, which references the instance of the class `InputState`. This class supports the methods that allow the keyboard and the mouse to be tested.

Because there is only one keyboard and one mouse, this class only has one instance that is referenced by the global variable called `Sensor`. To test the status of the keyboard or the mouse, messages have to be sent to the global variable `Sensor`. The messages understood by `Sensor` are:

```
anyButtonPressed  
returns true if one of the buttons on the mouse is  
pressed,
```

```
noButtonPressed  
returns true if none of the buttons on the mouse is  
pressed,  
    redButtonPressed  
returns true if the red button on the mouse is pressed,  
    yellowButtonPressed  
returns true if the yellow button on the mouse is  
pressed,  
    blueButtonPressed  
returns true if the blue button on the mouse is pressed,  
    waitButton  
waits for one of the buttons on the mouse to be pressed  
and returns the cursor position,  
    waitNoButton  
waits until none of the buttons on the mouse is pressed  
and returns the cursor position,  
    keyboardPressed  
returns true if one of the keys on the keyboard has been  
pressed,  
    keyboard  
waits until a character is struck on the keyboard and  
returns this character,  
    keyboardPeek  
reads a character from the keyboard queue without  
removing it from the queue,  
    keyboardNext  
reads a character from the keyboard queue and withdraws  
it from this queue,  
    cursorPoint  
returns the current cursor position,  
    cursorPoint: aPoint  
places the cursor at the position given as argument,  
    currentCursor  
returns the current cursor,  
    currentCursor: aCursor  
replaces the current cursor by the argument aCursor.  
For example:
```

```
Sensor waitButton.  
Display reverse.  
Sensor waitNoButton.  
Display reverse  
reverses the screen as soon as one of the mouse buttons  
is pressed and keeps it in this state until the button  
is released.
```

# 21 Model-View-Controller System

In the preceding chapters we have met several different window types in the Smalltalk system, each designed to solve a particular problem. The programmer can also build special windows for his own individual applications. We recall that a window specifies a part of the screen in which interaction takes place between the user and the application.

In this chapter we shall see how interactive applications in the Smalltalk system are constructed and how to build new ones.

Within Smalltalk all interactive applications are built to the same plan, called the MVC system. The approach is based on the following three entities:

- the object on which the work is to be done,
- an output interface that serves to present this object to the user,
- an input interface that allows interaction with the object or with its representation.

In the Smalltalk system, the object on which the work is to be done is called the **model**, the output interface is called the **view**, and the input interface is called the **controller**.


Whenever the model is changed, its representation must be updated. For this updating to be automatic, a dependency must be established between the view and the model that it represents. The concepts of dependence were examined in chapter 7 on the class Object.

In addition to the dependency of the view on the model, the view and the controller must be aware of one another and aware of the model on which they are to work.

The basic classes that allow views and controllers to be defined are the classes **View** and **Controller**. These two classes are subclasses of the class **Object**. The relationships between a model, its view and its controller are shown in the following figure.


## 21.1 The class View

The instances of the class **View** are windows. Two of the instance variables of this class are **model** and **con-**



troller, which refer respectively to the model and the controller assigned to the view.

When several objects are to be represented in the same view, the view can be divided into several subviews. Thus, a hierarchy of views is defined, since a subview is a completely separate view. To implement this hierarchy each view has a list of subviews and one superview. The following figure shows the relationships between models, views and controllers in the case where a view is made up of several subviews.



#### **21.1.1 The class WindowingTransformation**

The coordinates of a view may be defined in several coordinate systems:

- in its own coordinate system,
- in the coordinate system of its superview,
- in the coordinate system of the screen.

In order to move from one coordinate system to another, simple geometric transformations are used.

The class WindowingTransformation is a subclass of the class Object whose instances serve to represent the

geometric transformations consisting of a translation and a scaling factor. The two instance variables of this class are scale and translation. These geometric transformations will be used to change coordinate system in a view.

The messages that allow new geometric transformations to be created are:

identity

which returns a transformation without translation and with an enlarging factor of 1,

scale: anEnlargingFactor translation: aTranslation  
which returns an initialised geometric translation with the arguments of the method,

window: aWindow viewport: aViewport

which returns a transformation to move from the rectangle aWindow to the rectangle aViewport. This method is the one that constructs the transformation for passing the coordinates of a view to the coordinates of its superview.

The messages that can be sent to the instances of WindowingTransformation are:

applyTo: aPoint

which returns the point obtained by applying the geometric transformation to the point aPoint,

applyInverseTo: aPoint

which returns the point obtained by applying the inverse geometric transformation to the point aPoint,

compose: aGeometricTransformation

which returns a new instance of WindowingTransformation obtained by composing the receiver and the argument aGeometricTransformation.

### 21.1.2 Instance variables of the class View

The instance variables of the class View are:

model

which is the object represented by the view,

controller

which is the controller that allows interaction with the view or with the model,

superView

which is the superview of the view,

subViews

which is an instance of OrderedCollection containing the subviews of the view,


transformation

which is an instance of WindowTransformation allowing transfer from the coordinate system belonging to the view (given by window) to the coordinate system belonging to the superview,

viewport

which is a rectangle giving the coordinates of the view

expressed in the coordinate system of the superview,  
 window  
 which is a rectangle giving the coordinate system of the view,  
 displayTransformation  
 which is an instance of WindowTransformation allowing transfer from the coordinate system belonging to the view (given by window) to the coordinate system of the screen,  
 borderWidth  
 which may be a number, a point or a rectangle, giving the width of the border that encloses the view,  
 borderColour  
 which is a motif giving the colour of the border,  
 insetDisplayBox  
 which is a rectangle expressed in the screen coordinates representing the interior rectangle of the view (the view without its border),  
 insideColour  
 which is a motif giving the background colour of the interior of the view,  
 boundingBox  
 which is the minimum rectangle expressed in the view coordinates containing the view and its subviews.  
 The following figure recapitulates the different coordinate systems in which the position of a point can be expressed and the transformations that allow transfer from one system to another.



The coordinates belonging to a view are indicated inside the view and the coordinates of that view within its superview are shown outside.

**21.1.3 Messages understood by instances of View**

The class View, which is the basis of the entire windowing system, supports a large number of messages. These messages are divided into several categories which are examined in the following sections.

**Manipulation of views and subviews**

```
    addSubView: aView
adds the argument aView to the list of subviews of the receiver,
    addSubView: aView align: aPoint with: anOtherPoint
adds the argument aView to the list of subviews of the receiver, by fixing the viewport such that the point aPoint, expressed in the coordinate system of aView, is located at the point anOtherPoint, expressed in the receiver's coordinate system,
    addSubView: aView below: aPoint
adds the argument aView to the list of subviews of the receiver, by fixing the viewport such that the origin of the view aView is located at the point aPoint, expressed in the receiver coordinate system,
    firstSubView
returns the first subview of the receiver,
    lastSubView
returns the last subview of the receiver,
    removeFromSuperView
removes the receiver from the list of subviews of its superview, if it exists,
    removeSubView: aView
removes the argument aView from the list of subviews of the receiver,
    resetSubviews
removes all the subviews from the list of subviews of the receiver,
    superView
returns the superview of the receiver,
    superView: aView
replaces the superview of the receiver by the view given as argument,
    topView
returns the main view of the hierarchy of views containing the receiver,
    isTopView
returns true if the receiver has no superview,
    release
before removing a view, one must also remove the dependence of this view in relation to the model, together with all the dependencies of its subviews; this is the function of the message release.
```

**Access to coordinate systems**

    displayBox  
    returns the rectangle containing the receiver,  
        insetDisplayBox  
    returns the rectangle interior to the receiver,  
        window  
    returns the coordinate system of the receiver,  
        window: aRectangle  
    replaces the coordinate system of the receiver by that  
    given as argument,  
        window: aRectangle viewport: anOtherRectangle  
    replaces the coordinate system of the receiver and the  
    coordinates of the receiver within its superview by the  
    coordinates given as argument,  
        viewport  
    returns the coordinates of the receiver expressed in the  
    coordinate system of its superview,  
        boundingBox  
    returns the minimum rectangle containing the receiver  
    and its subviews,  
        containsPoint: aPoint  
    returns true if the point given as argument is within  
    the receiver.

**Access to the controller**

    controller  
    returns the controller assigned to the receiver,  
        controller: aController  
    replaces the controller of the receiver by the controller  
    given as argument,  
        defaultController  
    returns a new controller which is a default instance of  
    the class.  
        defaultControllerClass  
    returns the class of the controller.

**Access to the model**

    model  
    returns the object represented by the receiver,  
        model: anObject  
    replaces the object represented by the receiver by the  
    object given as argument,  
        model: anObject controller: aController  
    replaces the object represented by the receiver and the  
    assigned controller by the objects given as argument.

**Display**

    borderColour: aMotif  
    replaces the motif that provides the border colour of  
    the receiver by the motif given as argument,  
        borderWidth: aWidth

replaces the width of the receiver border by the width given as argument,  
    clear: aMotif  
fills the area of screen occupied by the receiver with the motif given as argument,  
    clearInside  
fills the interior of the receiver with the motif given by the instance variable insideColour,  
    clearInside: aMotif  
fills the interior of the receiver with the motif given as argument,  
    display  
displays the receiver border, the receiver itself, and its subviews,  
    displayBorder  
displays the receiver border,  
    flash  
makes that part of the screen occupied by the receiver flash,  
    highlight  
inverts that part of the screen occupied by the receiver,  
    insideColour  
returns the motif that defines the background colour of the receiver,  
    insideColour: aMotif  
replaces the background colour of the receiver by the motif given as argument.

The class View has several subclasses which allow different types of object to be represented (text, forms, list, etc). Interaction varies, according to the object represented; it is therefore necessary to assign a specific controller to each view. All its controllers are defined in subclasses of the class Controller.

## 21.2 The class Controller

The class Controller is a subclass of the class Object. Defined in this class are the basic operations for controlling interactions and locating the active controller.

### 21.2.1 Instance variables of Controller

This class has three instance variables:

    model  
which is the object to be interacted with,  
    view  
which is the view representing the model,  
    sensor  
which is a reference to the global variable Sensor, the only instance of the class InputSensor; use of such a reference saves time in locating this object.

**21.2.2 Messages understood by instances of Controller**

The messages supported by the class Controller are:

controlActivity

finds the subview of the view assigned to the receiver which requires control and activates the controller assigned to this subview,

controlInitialize

this message is sent to the receiver when the latter has just been activated, thus allowing certain parameters to be set,

controlTerminate

this message is sent to the receiver just before the latter is deactivated, thus allowing certain parameters to be restored,

controlLoop

loops for as long as the view assigned to the receiver contains the cursor, in order to discover which of the subviews wishes to take control (perhaps none),

startUp

this message is sent to a controller to activate it; the method sends to the controller the messages controlInitialize, controlLoop and controlTerminate.

isControlActive

returns true if the view is to retain control,

model

returns the object with which the receiver is interacting,

model: anObject

replaces the object with which the receiver interacts by the object given as argument,

view

returns the view assigned to the receiver,

view: aView

replaces the view assigned to the receiver with the view given as argument,

release

removes the references of the receiver to the view and the model.

To activate a controller, the message startUp must be sent to it. The startUp method is written as follows:

startUp

    self controlInitialize.

    self controlLoop.

    self controlTerminate.

The methods controlInitialize and controlTerminate are redefined in the subclasses of Controller. For example, for the class ScrollController, which controls scrolling in a text window, these methods cause the scrollBar to appear on the left of the view when the controller is first activated and disappear at the end.

The controlLoop method is written as follows:

```
controlLoop
  [self isControlActive] whileTrue:
    [Processor yield.
     self controlActivity]
```

When the cursor moves from one view to another, it is necessary to be able to change the active controller in order to interact with the view on which the cursor is located. There is within the Smalltalk system a class called **ControlManager** whose task is to control the activation of the different controllers.

### 21.3 The class ControlManager

The class **ControlManager** is a subclass of the class **Object**. The instance of the class **ControlManager** which controls the controllers at any given moment is referenced by the global variable called **ScheduledControllers**.

Its instance variables are:

  scheduledControllers

which references the list of controllers that may be activated; this list further contains an instance of **ScreenController** which controls the screen apart from any window,

  activeController

which references the active controller,

  activeControllerProcess

which references the process connected to the active controller; a new process is executed every time that the active controller changes,

  screenController

which references an instance of the class **ScreenController** that controls those parts of the screen not occupied by a window.

When there is no longer an active controller, the instance of **ControlManager** referenced by **ScheduledControllers** decides to which window it should hand over control.

If there is no window that wishes to take control, it hands control to the screen as soon as the user presses the yellow button on the mouse.

Each project contains a collection of windows and is directed by an instance of the class **ControlManager** that belongs to it. When a project is changed, the **ControlManager** of this project is still referenced by the global variable **ScheduledControllers**.

When the controller of a window is activated, it will decide itself whether it is to retain control or hand it to one of the subviews of the window.

### 21.4 The class MouseMenuController

The class **MouseMenuController** is a subclass of the class

Controller. The instances of this class handle mouse events. The instance variables of this class are:

redButtonMenu

which is a PopUpMenu that appears whenever the user presses the red button,

redButtonMessages

which is an array of symbols corresponding to the messages that will be sent to the controller if one of the menu options assigned to the red button is selected,

yellowButtonMenu

which is a PopUpMenu that appears whenever the user presses the yellow button,

yellowButtonMessages

which is an array of symbols corresponding to the messages that will be sent to the controller if one of the yellow button menu options is selected,

blueButtonMenu

which is a PopUpMenu that appears whenever the user presses the blue button,

blueButtonMessages

which is array of symbols corresponding to the messages that will be sent to the controller if one of the blue button menu options is selected.

If no menu is to be assigned to a mouse button, the instance variable containing the menu corresponding to this button will be set to nil.

The messages understood by the instances of the class MouseMenuController are:

redButtonMenu: aMenu

redButtonMessages: anArrayOfMessages

allows initialisation of the red button menu and list of messages,

yellowButtonMenu: aMenu

yellowButtonMessages: anArrayOfMessages

allows initialisation of the yellow button menu and list of messages,

blueButtonMenu: aMenu

blueButtonMessages: anArrayOfMessages

allows initialisation of the blue button menu and list of messages,

redButtonActivity

this message is sent to the receiver when the red button is pressed,

yellowButtonActivity

this message is sent to the receiver when the yellow button is pressed,

blueButtonActivity

this message is sent to the receiver when the blue button is pressed,

For the preceding three messages, if there is a menu assigned to the button, then the menu is displayed for

as long as the button is held down. When it is released, if one of the options was selected, the corresponding message is sent to the receiver.

Almost all the controllers are instances of Mouse-MenuController or of its subclasses, because almost all of them use the mouse.

The two important subclasses of this class are:

  ScrollController

whose instances allow control of the scrollbar that appears to the left of windows that require scrolling,

  StandardSystemController

whose instances allow screen windows to be manipulated (movement, changing size, closing, etc.).

## 21.5 The classes StandardSystemView and StandardSystem-Controller

These two classes implement the basic mechanisms that allow windows to be defined and controlled in the most general sense.

A window in the Smalltalk system is made up of a rectangle defining a part of the screen, surmounted by another rectangle containing the name of the window.

We have already seen how a window may contain subwindows which may themselves contain other windows. The window located at the top of the hierarchy, again called topView, is generally an instance of the class StandardSystemView. Only these windows may be deactivated or reactivated by the ControlManager.

The class StandardSystemView is a subclass of the class View. The additional messages supported by this class are:

  label

returns the name of the receiver,

  label: aString

replaces the name of the receiver by the string given as argument,

  labelDisplayBox

returns the rectangle containing the name,

  reverseLabel

inverts the name of the receiver,

  expand

changes the size of the receiver interactively,

  minimumSize

returns the minimum size of the receiver that is used when beginning to interact with it,

  minimumSize: aSize

fixes the minimum size of the receiver,

  maximumSize

returns the maximum size of the receiver that is used when beginning to interact with it,

  maximumSize: aSize

fixes the maximum size of the receiver.

The class **StandardSystemController** is a subclass of the class **MouseMenuItemController** whose instances generally control the windows that are instances of the class **StandardSystemView**. It has two class variables, namely:

ScheduledBlueButtonMenu

which is the PopUpMenu that is assigned by default to the blue button for all instances of this class,

ScheduledBlueButtonMessages

which is an array of messages sent to the receiver according to the command selected from the menu assigned to the blue button.

It has an instance variable called **status** which provides the status of the controller. The different possible states of a window are:

open

which is the status of the controller when interaction with the designated window is begun,

closed

which is the status of the controller when the designated window is closed (it is no longer visible on the screen),

active

which is the status of the controller when the designated window is active (that is, when the user can interact with this window),

inactive

which is the status of the controller when the designated window is no longer active.

To open a window, the message **open** must be sent to its controller.

## 21.6 Text-scrolling windows

When a section of text is too long to appear in a window in its entirety, it is only possible to display part of it, while allowing access to the remainder using a scrollbar.

The control mechanisms for the scrollbar are implemented in the class **ScrollController** which is a subclass of the class **MouseMenuItemController**.

When a window has a scrollbar, the latter is only visible when the cursor is located within the window. Before making the scrollbar appear, the **ScrollController** saves the area of the screen that will be covered by the scrollbar in an instance variable called **savedArea**. When the **ScrollController** makes the scrollbar disappear, it redisplays the form contained in **savedArea**.

In the Smalltalk system the scrollbars are used in two types of window:

- text-editing windows

- windows for selecting an element from a list (see figure below).



The two subclasses of the class `ScrollController` that allow management of these two types of view are the classes `ParagraphEditor` and `ListController`.

#### **21.6.1 Text-editing windows**

We have seen that, in order to represent sections of text on the screen, we may use instances of the class `DisplayText`. To present a text in a window, it also has to be formatted; that is, there has to be automatic control for moving to a new line when a line is filled or a word will not fit on the current line. The objects that allow formatted text to be produced are instances of the class `Paragraph`.

##### **21.6.1.1 Paragraphs**

The class `Paragraph` is a subclass of `DisplayText`, whose instances responds to certain additional messages.

A paragraph may be created from a section of text or from a string of characters by sending them the message `asParagraph`, for example:

```
'this is a string of characters' asParagraph
returns a paragraph that represents the receiver.
```

As the text of a paragraph is formatted, there are several ways of putting this text in shape, and the messages that allow choice of page layout are:

```
centred
which allows the text to be centred on a line,
justified
which allows the text to be aligned on the left and
right margins,
```

```
leftFlush
which allows the text to be aligned on the left margin;
this is the default option,
```

```
rightFlush
which allows the text to be aligned on the right margin.
```

When the size of a text window changes (for example, with the option frame via the blue button), the text formatting of the text in this window must also be changed. There is a message that allows a section of text to be reformatted, namely:

```
recomposeIn: aRectangle clippingBox: anOtherRectangle
which recalculates the page layout for the paragraph, so
```

that the lines will fit into the limits of the argument aRectangle and the display is restricted to the argument anOtherRectangle.

The controllers that allow a paragraph to be edited are instances of the class ParagraphEditor, which is a subclass of ScrollController. These controllers respond to a significant number of messages which the user has no need to access; these are therefore not described in this chapter.

In a text-editing window, the trio model-view-controller is made up of:

- a section of text or a string of characters for the model,
- an instance of View or of one of its subclasses, which displays the paragraph created from the model, for the view,
- an instance of ParagraphEditor or of one of its subclasses for the controller.

#### 21.6.1.2 **StringHolder** **StringHolderView**

##### **StringHolderController**

These three classes make up a trio of model-view-controller that allow the text of the model to be edited in a text-editing window.

The class **StringHolder** is a subclass of the class Object, whose instance variables are:

contents

which is the string of characters that is edited,

isLocked

which is a boolean that indicates whether the text of the view is the same as that of contents.

The class **StringHolderView** is a subclass of View. One of its instance variables, called displayContents, references the paragraph that represents the formatted text of the model. The instances of StringHolderView are not main views (main views are generally instances of StandardSystemView), but simply subviews.

The class **StringHolderController** is a subclass of the class ParagraphEditor which allows the paragraph assigned to the model to be edited.

#### 21.6.2 **Lists**

Lists are views which contain one or more lines of text and from which it is possible to select a line. For example, lists may be found in the four windows at the top of a browser; these allow access to the categories, classes, protocols and methods.

The MVC trio assigned to the control of a list consists of:

- an object that may in certain cases be considered as

a list,  
 - a view that is an instance of a subclass of **ListView**,  
 - a controller that is an instance of a subclass of **ListController**.

The subclasses of **ListView** and **ListController** that allow any lists to be manipulated are **SelectionInListView** and **SelectionInListController**.

The object which represents the model may be any object that responds to the messages that cause the system to:

- return the selected element
- change the selected element
- return the list of all elements
- return a menu assigned to the centre button which allows action on the elements in the list.

The name of the messages that will cause the above functions to occur will be given to the view at the time that it is created. Views of this type are rather special, because their model may belong to any class; they are said to be 'pluggable' views.

To create a view of this type, a message consisting of the following various selectors must be sent to the class **SelectionInListView**:

```

on: anObject
printItems: aBoolean
oneItem: anOtherBoolean
aspect: aspectMessage
change: changeMessage
list: listMessage
menu: menuMessage
initialSelection: selectionMessage

```

which returns an instance of **SelectionInListView** whose model is the argument **anObject**.

The argument **aBoolean** determines if the text of the view must be obtained by sending the message **printString** to all the elements in the list or if these elements are already strings of characters.

The argument **anOtherBoolean** determines if the list contains one and only one element.

The argument **aspectMessage** is the message that will be sent to the model to access the element selected from the list.

The argument **changeMessage** is a message to an argument that will be sent to the model to change the selected element.

The argument **listMessage** is the message that will be sent to the model to access the list of elements.


The argument **menuMessage** is the message that will be sent to the model to have a menu assigned to the centre button.

The argument **selectionMessage** is the message that


will be sent to the model to access the element selected at the start.

### 21.6 Example


The following review example shows how to construct an application that will allow a list of curves to be manipulated in one window, while the selected curve is represented in a second window, as shown below.



The equation of a curve may be represented by a block with one argument. The orientation of the axes in the view is that normally adopted in mathematics. The scale goes from 0 to 100 along the abscissa and the ordinate. To implement this application, only two classes are required, as shown below.




The class CurveList is a subclass of OrderedCollection whose elements will be the strings of characters that provide the equations of the curves.



The class CurveListView is a subclass of View which simply serves to redefine the method displayView that allows a curve to be displayed.

To open a curves window it is simply a question of evaluating the following expression.



The first method to be written is the method for creating a window. The text of this method appears in the window below.

Category Browser	
Example	instance creation examples
CurveList (CurveListView)	openOn:
Instance	class

```
openOn: aCurveList
    | topView listView subView |
    topView + StandardSystemView new.
    topView insideColor: Form white; borderWidth: 1; window: (0 @ 0 extent: 1
@ 1); label: 'Curve Plot'.
    listView + SelectionInListView
        on: aCurveList
        printItems: false
        oneItem: false
        aspect: #selectedCurve
        change: #selectedCurve:
        list: #curveList
        menu: #curveMenu
        initialSelection: #selectedCurve.
    listView window: (0 @ 0 extent: 1 @ 0.3); borderWidth: 1.
    subView + self new.
    subView insideColor: Form white; borderWidth: 1; window: (0 @ 100 extent:
100 @ -100)
        viewport: (0 @ 0.3 extent: 1 @ 0.7); model: aCurveList.
    topView addSubView: listView; addSubView: subView.
    topView controller open
```


The window is made up of three views:

- a main view which is an instance of StandardSystemView; the coordinates of this view (given by window:) range from 0@0 to 1@1,
- a subview (an instance of SelectionInListView) which allows the list of curves to be represented and manipulated; the coordinates of this view in its superview are 0@0 and 1@0.3,
- a subview (an instance of CurveListView) which is the view in which the selected curve is drawn; the coordinates of this view in its superview are 0@0.3 and 1@1. The coordinates belonging to this view range from 0@100 to 100@0, which represent a conventional (mathematical) coordinate system (the axis of the abscissas from left to right, and the axis of the ordinates from bottom to top).


To create the instance of SelectionInListView, it was necessary to give the messages that will be sent to the model in order to access the list of curves; now, the methods that correspond to these methods must be defined. These methods will be defined within the model, namely ListCurve.

The method curveList allows access to the list of elements of CurveList; here, this list directly consists


of the receiver, as shown in the figure below.




The method `curveMenu` returns a menu (instance of the class `ActionMenu`) which is a subclass of `PopUpMenu` including an additional instance variable which references the list of messages assigned to the menu options. If no line is selected, the menu returned only includes the option to add a curve to the list. If a curve is selected from the list, the menu returned includes three options, allowing a curve to be added, removed, or the selected curve to be altered.



The method `selectedCurve` returns the variable called `selectedCurve` (see below).



The method which allows a different curve to be selected is described below. Note that sending the message changed to the receiver has the effect of updating the view; in fact, it should be recalled that there is a dependence between the model and its view, and therefore any alteration of the model is signalled to the view.




There are two protocols in the class `CurveList`: the first, which contains the access methods, has already been described; the second contains the methods for manipulating the list.

To add a curve, its curve must be input in the form of a block; this block must have an argument which, during evaluation, will take the different values of the abscissa.


The class `FillInTheBlank` allows text in a window to

be captured; the window closes automatically when the capture is completed by using the accept option on the middle button or <return>. The expression

```
self changed: #selectedCurve  
redisplays the list.
```




The following method allows a curve to be removed from the list.




To alter a curve, a window needs to be opened on the text of the selected curve; then, once the text has been altered, the selected curve is replaced by the new text.

All the methods of CurveList are written; all that remains is to write the instance methods of CurveListView. The first of these methods is very simple, since it involves redefining the message update:. Remember that this message is sent to all the dependents of an object when the latter receives the message changed:. Here, any alteration of the model



(that is, the list of curves), must be result in a redisplay of the view.



The method that displays the interior of a view is called `displayView`. Here, this method must be redefined in order to display the curve within the view.

The method begins by erasing the interior of the view, since if there is a selected curve it draws it. The first task to be done is to transform the string of characters that represents the curve into a block; to do this, the text of the curve must be evaluated.

The pen can then be placed on the first point and be made to draw straight lines between all the following points. The transformation called `displayTransformation` must be applied to all the points in order to transfer the coordinates of the view to the coordinates of the screen.

This example illustrates how it is possible to construct graphics applications quickly, using but a few classes and methods. It will be noted that some methods

The screenshot shows the Smalltalk-80 Category Browser interface. The left pane displays the class hierarchy: 'Example' is the root category, containing 'CurveList' and 'CurveListView'. 'CurveListView' is selected. The right pane shows two methods: 'displaying updating' and 'displayView'. The 'displayView' method is expanded, showing its source code:

```

displayView
| textBlock block bic |
self clearInside.
(textBlock := model selectedCurve) notNil
ifTrue:
    [block := Compiler new
        evaluate: textBlock
        in: nil
        to: nil
        notifying: nil
        ifFail: []].
bic := Pen new frame: insetDisplayBox.
bic place: (displayTransformation applyTo: 0 @ (block value: 0)).
1 to: 100 do: [:i | bic goto: (displayTransformation applyTo: i @ (block
value: i)) rounded]]]

```

serve to redefine existing methods, which is a feature often found in Smalltalk, because an application is often defined by the differences it has in comparison with an existing application (differential programming).

This example is also typical of what might be an application prototype.

The constraints imposed on such prototypes are as follows:

- shorter realisation time
- ease of development
- minimum operation; that is, there is no need to wait for the complete realisation of the prototype before running those parts already defined.

The MVC system is fundamental to Smalltalk, because it provides a methodology for breaking down graphics and interactive applications. Given that there are more than 30 subclasses of the class View, as well as more than 30 subclasses of the class Controller, it has only been possible to discuss the main classes in this chapter.

# Index

@, 96  
[, 15  
^, 22

## A

accept, 30  
activeProcess, 92  
add:, 68  
addFirst, 77  
addLast, 77  
addSubView:, 156  
again, 28, 127  
allInstances, 44  
and, 111  
applyTo:, 154  
**Arc**, 88  
arithmetic, 59  
**Array**, 81  
**ArrayedCollection**, 81-82  
arrays, 14  
asciiValue, 57  
asDisplayText, 107  
**Association**, 71  
at:, 34, 72  
at:put:, 34

## B

**Bag**, 71  
**Behavior**, 41  
binary message, 16  
**BitBlt**, 114-115  
Bit editor, 148-149  
bit manipulation, 65  
bitmap, 26  
**Bitmap**, 105  
black, 108  
**BlockContext**, 91-92  
blocks, 18-20  
blueButtonActivity, 161  
border:, 108  
borderWidth:, 157  
boundingBox, 155  
browsers, 28, 129  
**ByteArray**, 83

## C

calendar, 54  
cascades, 18  
category, 48, 130  
changed, 37  
**Character**, 31, 57  
characters, 14  
character strings, 14  
**Circle**, 114  
class, 10-12, 31  
**Class**, 12, 40, 47  
**ClassDescription**, 12, 47  
class methods, 11  
classes, 7, 10-11, 131  
clearInside, 158  
close, 89, 120  
code window, 136, 143  
coerce, 63  
**Collection**, 67-70  
comment, 48  
compilation, 49  
contains, 103  
containsPoint:, 103  
contents, 84  
Control-C, 143  
controlActivity, 159  
controlInitialise, 159  
controller, 152, 157  
**Controller**, 152, 158  
controlLoop, 159, 160  
**ControlManager**, 160  
controlTerminate, 159  
conversions, 56, 62  
copy, 28, 32, 33, 127  
corner:, 100  
cornerBy:, 102  
**CurrentInputState**, 150  
**Cursor**, 112-113  
cursorPoint, 151  
cut, 28, 127

**D**

**Date**, 54  
**deepCopy**, 33  
**debugger**, 30, 142  
**defaultControllerClass**, 157  
**Delay**, 93  
**dependencies**, 36-37  
**destForm**, 114  
**detect:**, 69  
**Dictionary**, 71-74  
**display**, 6  
**displayAt:**, 21, 106  
**displayOn.at:clippingBox:**  
  **rule:mask:**, 112  
**Display**, 113  
**DisplayMedium**, 107-110  
**DisplayObjecte**, 105-106  
**DisplayScreen**, 112  
**DisplayText**, 196-107  
**displayTransformation**, 155  
**dist.**, 97  
**do:**, 69  
**doesNotUnderstand:**, 22  
**doIt**, 28, 128

**E**

**edit commands**, 127  
**enumeration**, 64  
**equality**, 33  
**equivalence**, 33  
**erase**, 111  
**error**, 35  
**expandBy:**, 102  
**explain**, 137  
**extent:**, 21, 100  
**ExternalStream**, 88-89

**F**

**false**, 19  
**file "changes"**, 128  
**file "sources"**, 128  
**fileIn**, 144  
**file list**, 27  
**fileOut**, 144  
**FileDirectory**, 89  
**fileNamed:**, 89  
**FileStream**, 50, 89  
**fill:rule:mask:**, 109  
**FillInTheBlank**, 171  
**findFirst**, 76

**findLast**, 76  
**first**, 74  
**firstSubView**, 156  
**flash:**, 113  
**Float**, 48, 59, 64  
**follow:while**, 112  
**fork**, 91  
**Form**, 21, 110-112, 147-148  
**format**, 136  
**Fraction**, 48, 59, 64  
**frame**, 116  
**fromDisplay**, 110  
**fromUser**, 110

**G**

**get contents**, 144  
**goto:**, 117  
**graphics editors**, 147  
**grid:**, 97

**H**

**halftoneForm**, 115  
**halt**, 143  
**height**, 101, 106  
**hierarchy**, 46, 132

**I**

**IdentityDictionary**, 74  
**ifTrue:ifFalse:**, 19  
**includes:**, 69  
**includesKey:**, 73  
**indexOf:**, 74  
**InputSensor**, 150  
**InputState**, 150  
**insetBy:**, 102  
**insetDisplayBox**, 157  
**inspect**, 139

**Inspectors**, 139-141  
**Instance**, 7, 9, 43  
**instances**, 43-45, 67  
**Integer**, 48, 59, 64  
**interface**, 118  
**intersect:**, 102  
**intersects:**, 103  
**Interval**, 80-81  
**isControlActive**, 159  
**isEmpty**, 95  
**isKindOf:**, 31  
**isMemberOf:**, 31  
**isNil**, 34

**K**

keyboard, 151  
 keyword messages, 16  
 keys, 72

**L**  
 label, 162  
**LargeNegativeInteger**, 48, 59  
**LargePositiveInteger**, 48, 59  
 last, 74  
 lastSubView, 156  
**Line**, 114  
**LinkedList**, 79-80  
 lists, 165  
**ListController**, 164, 166  
**ListView**, 166

**M**

**Magnitude**, 53  
 mathematics, 60  
 max:, 53  
 maximumSize:, 162  
 menus, 26, 119, 120-121  
 merge:, 102  
 messages, 6, 15, 21, 35  
**Metaclass**, 12, 40, 50  
 metaclasses, 11  
 methods, 8, 21, 25  
 millisecondsToRun:, 56  
 min:, 53  
 minimumSize:, 162  
 model, 152, 157  
 mouse, 26, 118  
**MouseMenuItemController**, 160-162  
 motif, 107  
 multiple inheritance, 50  
 MVC, 152

**N**

name, 48  
 new, 21, 43  
 newFileNamed:, 89  
 next, 84, 95  
 nextPutAll:, 84  
 nil, 23, 34  
 notNil, 34  
 now, 56

**Number**, 48, 59  
 numbers, 13

**O**  
**Object**, 31  
 objects, 6  
 oldFileNamed:, 89  
 on:, 85  
 open, 89  
**OrderedCollection**, 37, 77-78  
 origin, 100, 101  
 offset, 106  
 over, 111

**P**  
 padTo:, 88  
**Paragraph**, 164-165  
**ParagraphEditor**, 164  
 paste, 28, 127, 140  
**Path**, 113-114  
 peek, 85, 95  
**Pen**, 115-117  
 perform:, 36  
 pixel, 105  
**Point**, 96-99  
**PopUpMenu**, 119  
**PositionableStream**, 85-86  
 primitives, 25, 38  
 printIt, 28, 128  
 printString, 34  
 priority, 17, 92, 93  
 proceed, 142  
**Process**, 91  
 Processor, 92  
**Processor-Scheduler**, 92-93  
 protocol, 129, 134  
 pseudo-variables, 23

**Q**

quit, 27

**R**

raisedTo:, 16-17  
**Random**, 48  
 readFrom:, 110  
**ReadStream**, 87  
**ReadWriteStream**, 88  
 receiver, 16, 24

**Rectangle**, 99-104  
**redButtonActivity**, 161  
**reject:**, 69  
**release**, 159  
**remove:**, 68, 73  
**removeFirst**, 78  
**removeLast**, 78  
**rename**, 130  
**respondsTo:**, 31  
**restore display**, 27  
**resume**, 92  
**reverse**, 110, 111  
**reverseDo:**, 77  
**rounding**, 62  
**RunArray**, 83

**S**  
**save**, 27  
**savedArea**, 163  
**ScheduledControllers**, 160  
**ScreenController**, 160  
**scrollbar**, 122  
**ScrollController**, 162, 163  
**select:**, 69  
**selectors**, 16, 42  
**SelectionInListController**,  
  166  
**SelectionInListView**, 166  
**self**, 23  
**Semaphore**, 94  
**Sensor**, 150-151  
**SequenceableCollection**,  
  36, 74-77  
**Set**, 71  
**shallowCopy**, 32  
**SharedQueue**, 95  
**showWhile:**, 112  
**signal**, 94  
**size**, 34  
**skip**, 86  
**skipTo:**, 86  
**SmallInteger**, 31, 48, 59  
**Smalltalk**, 10  
**someInstance**, 44  
**SortedCollection**, 79-79  
**sourceForm**, 114  
**spawn**, 130  
**StandardSystemController**,  
  162-163  
**StandardSystemView**,  
  119-120, 162-163

**startUp**, 159  
**status**, 163  
**storeString**, 35  
**Stream**, 84-85  
**String**, 81  
**StringHolder**, 165  
**subclasses**, 10  
**subViews**, 154  
**super**, 24  
**superclass**, 10, 40, 46  
**superView**, 154  
**suspend**, 92  
**Symbol**, 31, 83  
**symbols**, 14

**T**  
**terminate**, 92  
**terminateActive**, 92  
**tests**, 42, 61  
**Text**, 42, 83  
**TextStyle**, 107  
**Time**, 54  
**timesRepeat:**, 64  
**to:**, 81  
**today**, 54  
**topView**, 156  
**transcript**, 30  
**true**, 19  
**truncateTo:**, 97  
**truncation**, 62

**U**  
**unary messages**, 16  
**under**, 111  
**undo**, 28  
**update:**, 37

**V**  
**value**, 16  
**values**, 72  
**variables**, 9, 15  
  class, 9  
  global, 10  
  indexed instance, 9  
  named instance, 9  
  pool, 10  
  temporary, 22  
**View**, 152-153, 154-158,  
  168  
**view**, 152  
**virtual machine**, 92

**W**

wait, 94  
waitFor, 151  
whileFalse, 19-20  
whileTrue:, 19-20  
white, 108  
width, 101, 106  
window, 26, 118, 157  
with:, 87  
**WindowingTransformation**,  
    153-154

**X**

x, 21, 96

**Y**

yield, 92  
y, 21, 96  
yellowButtonActivity, 161